



UVM Training Guide

Table of Contents

<i>Revision History</i>	6
<i>Introduction</i>	7
About This Tutorial.....	7
About utut.....	7
<i>Prerequisites</i>	9
<i>The ALUTB Testbench</i>	10
<i>Lesson 1: Your First Test</i>	12
Using the UVM Template Generator	12
flist Files.....	14
Extended from the Base Test.....	14
Changing the Clock Frequency	14
Overriding Constraints.....	15
Conclusion.....	17
<i>Lesson 2: Class Factories</i>	18
Extending the Driver.....	18
Using The Factory	19
Conclusion.....	22
<i>The ALU Protocol</i>	23
<i>Lesson 3: Creating a Transaction</i>	25
Creating a vcomponent	25
A Sequence Item is a Transaction	26
Typedefs Rule	27
Sequence Item U-Turns	28
Adding Constraints	28
Printing Your Classes	29
Packing and Unpacking	31
Printing, Comparing, Packing...Recording?	35
Adding Coverpoints.....	36
Conclusion.....	36
<i>Lesson 4: Creating an Interface</i>	37
The ALU Interface Explained	40
Connecting the Interface to the DUT.....	40
Storing the Interface in the Database	41

Conclusion.....	43
Lesson 5: Creating an Agent	44
Creating Multiple Files	44
Objects vs. Components	44
All About Configurations	49
All About The Class Factory	51
Factory Overrides on Non-Components	52
How To See Everything.....	53
Conclusion.....	53
Lesson 6: TLM Ports, Imps, and Exports	54
TLM Ports	54
TLM Implementations (Imps)	55
Connecting	55
TLM Exports	56
TLM Summary	56
Imp Declarations	57
Compared to TestBuilder.....	58
Many-to-One and One-to-Many	59
Broadcasting to a Listener.....	60
Conclusion.....	62
Lesson 7: Beginning Sequences	64
A Simple Generator.....	64
`uvm_do Macros	66
Getting a Response	67
Response Queue Handler.....	69
Complex Routines	71
Sequencing other Sequences	77
Sequence Hierarchy.....	80
The Sequencer.....	81
Locking and Grabbing	83
Prediction Encapsulation	84
Conclusion.....	86
Lesson 8: Drivers and Monitors	87
Fetching the Interface	87
Getting the Next Item	89
Driving a Transaction	90

Monitoring Activity	92
Conclusion	96
Lesson 9: Writing a Predictor	97
Monitor Prediction	97
Scoreboards	97
PW Scoreboard	98
Comparators	98
Prediction Encapsulation	99
Roll-Your-Own.....	99
Reference Modeling	101
Conclusion	106
Lesson 10: Configuration Registers	107
Register Organization	107
References for Everyone!.....	108
Configuring Your Block	109
The kval Test	109
How It All Works	111
Fixing Our Predictor	113
Register Callbacks	114
Built-In Register Sequences	115
Conclusion	116
Lesson 11: Writing Tests	117
CFG Classes	117
Test Overrides	120
Conclusion	123
Lesson 12: Advanced Sequences I	125
Default Sequences.....	125
Library Sequences	126
Library Sequence Configuration	127
Persistent Sequences	129
Sequencer References	130
Handling Cyclical Dependencies	132
Sequence Configurations	137
Conclusion	138
Lesson 13: A New Testbench	139
The ALU Framer	139

Creating the FRM Testbench	139
UVM New Testbench (untb)	140
<i>Lesson 14: Advanced Sequences II</i>	144
Virtual Sequences and Sequencers	144
Virtual Sequences	146
Adapter Sequences	150
Adding Time Delays	152
Conclusion	154
<i>Appendix A: Phases and Heartbeats</i>	155
Watchdog Timer	155
Phase Objections	155
Deadlock Checking	158
<i>Appendix B: Common Recipes</i>	160
Implementing a Watchdog	160
Forking Multiple Instances of a Method	160
Phase-Boundary Crossing Tasks	161
Handling a Reset	161

Revision History

Version	Author	Date	Notes
0.1	Brian Hunter	Oct. 5, 2011	Initial release to SPOC working group.
0.2	Brian Hunter	Oct. 27, 2011	Second release to SPOC working group.
0.3	Brian Hunter	Dec. 1, 2011	All major content complete.
0.4	Brian Hunter	Dec. 30, 2011	Release for testing.

Introduction

Welcome to Cavium's UVM Tutorial. In these lessons, you will learn the key concepts of UVM and how to use it with SystemVerilog for verification within Cavium's environment.

To get to this document, you should have already run the `uvm_tutorial.py` script. The script should also have unpacked a usable directory that exists on an svn branch that you will own. You may check your modified files in or out without affecting anybody else's work.

About This Tutorial

This tutorial walks you through the process of adding to an existing testbench. It takes a dive-right-in approach. Each lesson introduces one or more concepts and presents problems that you are expected to attempt to solve. *Some of these problems are intentionally difficult.* The objective of the problems is to challenge you to find the correct answers through trial-and-error, in order to better familiarize yourself with what works, what doesn't, and where to find the answers.

If you need help solving the problems, your best courses of action are to:

- read through the existing UVM code that comprises the testbench;
- refer to the UVM reference manual and user's guide, the SystemVerilog Language Reference Manual, or;
- visit <http://www.uvmworld.org>.

If you are unable to solve a problem, a solution is presented on the following page.

About utut

`utut` is a script that is located in `$CAD_DIR/common/bin`. Use this script to deploy a working area of the UVM tutorial codebase, and various other operations.

Create

To create a working area branch of the UVM tutorial codebase in which you can modify files and check in without affecting anyone else's branch, run the following command. Give it the name of the directory you would like to create:

```
> utut create mydir
```

Code

Many code snippets presented in this tutorial are preceded by a small number. Using the `utut` script as follows will print out the code that you can then cut-and-paste into your own files. This example prints out code snippet #17:

```
mydir/verif/alutb> utut code 17
```

Lesson

Each lesson throughout this tutorial has a checkpoint file available. You can run the following command to update your working copy to be identical to the solutions. This will, however, delete any modifications that currently exist in your working directory.

```
mydir/verif> utut lesson 3
```

Cleanup

Whenever you want to remove any working branches that you have created with `utut`, you can run the `cleanup` command. This command will allow you to remove one or all of the working branches associated with the user. It will not delete any working directories on your system.

```
mydir> utut cleanup
++ Looking for branches for user username
++ Found these branches:

    0: uvmtut_username_mytut/
    1: uvmtut_username_mytut2/

Select branches to be deleted [0-1 or all]: 0
Removing these branches:
['uvmtut_username_mytut/']
++ Running 'rs4 rm -m RemoveBranchesFor:username svn+ssh://masvn/svn/o68/branches/
uvmtut_username_mytut/' in directory mydir
```


Prerequisites

Before beginning the UVM Tutorial, you should already be quite familiar with SystemVerilog's verification components, including but not limited to these chapters of specification:

- Classes and Inheritance
- Constrained Random Value Generation
- Processes
- Functional Coverage Collection
- Data Types and Aggregate Data Types
- Interfaces, Modports, and Clocking Blocks
- Interprocess Synchronization and Communication

Also, you should have the **SystemVerilog 1800-2009** specification handy for your reference.

While this is a tutorial on Cavium's usage of UVM, it assumes at least a nascent knowledge of UVM. Therefore, it is helpful for you to read through the **UVM User's Guide 1.1** and familiarize yourself with the concepts it lays out.

It is also helpful to keep the **UVM 1.1 Class Reference** handy.

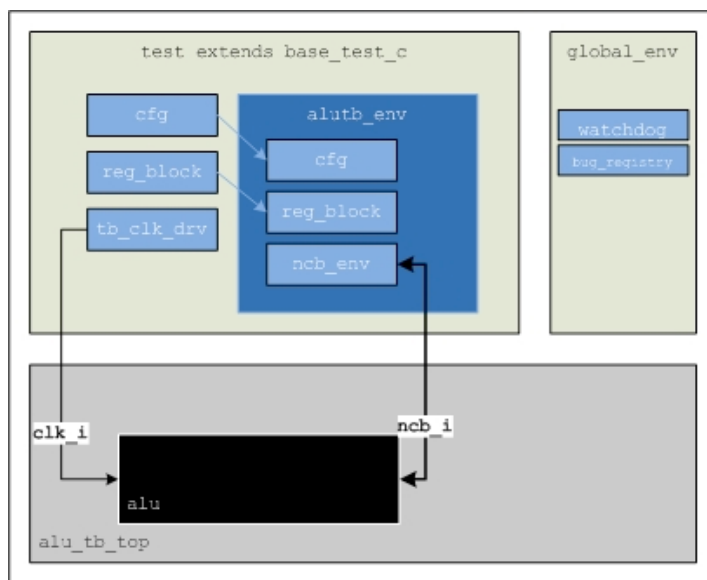
You should definitely familiarize yourself with our company's **SV Guidelines** on our internal wiki. This document describes coding guidelines, naming conventions, testbench organization, and generally preferred best practices. It also points you to the **NaturalDocs website**, which describes how code should be documented to work with this documentation generator.

All of these aforementioned documents are linked [here](#).

The ALUTB Testbench

The ALU testbench (`verif/alutb`) is a pure UVM testbench created for the purposes of this tutorial. The Device Under Test (DUT) is a block called `alu (rtl/alu)`. This block is not a real RTL block in the chip. It has CSRs on the NCB interface, and has a simple interface to perform math computation requests. In this tutorial, you will be responsible for creating UVM testbench code to test this device (and find bugs!).

The top-level hierarchy of the current ALUTB testbench approximates this diagram:



We'll discuss these particulars in more detail later. But first note that the top-level of the environment hierarchy is the test itself. This is a radical departure from TestBuilder. Also notice that each testbench instantiates a `global_env`, which is a container for functionality that is needed by every UVM testbench.

Feel free to read some of the SystemVerilog code that is already present in this testbench, and familiarize yourself with its layout and style. The code is located in these directories:

- `verif/alutb` - Contains this testbench and all of its tests.
- `verif/vcomponents/alutb` - Contains the `alutb` vcomponent (verification code that can be reused in a higher level testbench).
- `verif/vcomponents/ncb` - Contains the NCB vcomponent.
- `verif/vcomponents/global` - Contains the global vcomponent and things that live in the global environment.

- `verif/vcomponents/cn` - Contains miscellaneous Cavium-specific utility functions, macros, message customizations, and components.
- `verif/vcomponents/uvm/1_1` - Contains the UVM 1.1 library itself.

If you enter the `verif/alutb` directory and run the following, you should get a test to run and PASS.

```
verif/alutb> cmake sim SIMU=VCS
```

If your simulation fails, proceed no further and seek some help.

Lesson 1: Your First Test

In this lesson, you will learn to use the UVM Template Generator to write a test in the `alutb` testbench.

Using the UVM Template Generator

The first objective is to write a test that speeds up the clock frequency.

The template generator automates the creation of UVM files and classes. You should use it all the time because it saves you a considerable amount of typing. It also creates files and classes with a consistent look and feel, which encourages readability.

The first thing to do is go into the `alutb/tests` directory and create a new test. The script is called `utg` and it takes as its argument the name of a template. You also give it the name of the file/class to create:

```
verif/alutb/tests> utg test -n fast_clk
++ Creating fast_clk.sv
fast_clk test: Enter substitution for <description>: A test that makes the clock go super fast!
/*-- mode: Verilog; verilog-indent-level: 3; indent-tabs-mode: nil; tab-width: 1 --*/

// *****
// * CAVIUM CONFIDENTIAL
// *
// *
// * PROPRIETARY NOTE
// *
// * This software contains information confidential and proprietary to
// * Cavium, Inc. It shall not be reproduced in whole or in part, or
// * transferred to other documents, or disclosed to third parties, or
// * used for any purpose other than that for which it was obtained,
// * without the prior written consent of Cavium, Inc.
// * (c) 2011, Cavium, Inc. All rights reserved.
// * (utg v0.5.1)
// *****
// File: fast_clk.sv
// Author: username
/* About: A test that makes the clock go super fast!
*****/

`ifndef __FAST_CLK_SV__
`define __FAST_CLK_SV__

`include "base_test.sv"
// (`includes go here)

// class: fast_clk_test_c
// (Describe me)
class fast_clk_test_c extends base_test_c;
`uvm_component_utils_begin(fast_clk_test_c)
`uvm_component_utils_end

//-----
// Group: Configuration Fields

//-----
// Group: Fields

//-----
```

```

// Group: Methods
function new(string name="test",
             uvm_component parent=null);
    super.new(name, parent);
endfunction : new

////////////////////////////////////
// func: build_phase
virtual function void build_phase(uvm_phase phase);
    super.build_phase(phase);
endfunction : build_phase

////////////////////////////////////
// func: connect_phase
virtual function void connect_phase(uvm_phase phase);
    super.connect_phase(phase);
endfunction : connect_phase

////////////////////////////////////
// func: end_of_elaboration_phase
virtual function void end_of_elaboration_phase(uvm_phase phase);
    super.end_of_elaboration_phase(phase);
endfunction : end_of_elaboration_phase

////////////////////////////////////
// func: run_phase
virtual task run_phase(uvm_phase phase);
    super.run_phase(phase);
endtask : run_phase

endclass : fast_clk_test_c

`endif // FAST_CLK_SV

```

Because we did not add `-f` to the command-line, `utg` just printed its output to the screen.

As you can see, `utg` chose to name the class `fast_clk_test_c`, even though we asked it to be named `fast_clk`. This is because according to our coding guidelines, tests must have the word `test` as a suffix, and all classes must end in `_c`. Some of our internal tools expect these conventions to be followed, and it's much easier to tell what kind of class it is, isn't it? It also created lots of (nearly) empty tasks and functions for you, to save you some typing. If you don't need these, you can leave them there for future expansion, or you can delete them if you like smaller files.

`utg` also asked you to fill in a description. The templates have tagged identifiers in them, and if `utg` doesn't know what to put in there, it will ask you. In order to skip these requests and just have `utg` leave the tags in the document, add `-q` to the command-line. Note that some tags are not within comments, and will break the compile if they remain.

You can cut and paste this into a file, or you can use this technique to make sure that `utg` will do what you want it to do. Instead, let's just have `utg` pipe this directly into the file called `fast_clk.sv`, and we'll have it leave the tags in the resultant file:

```

verif/alutb/tests> utg test -n fast_clk -f -q
++ Creating fast_clk.sv

```

Now, you have a file called `fast_clk.sv` in the tests directory and it's time to actually do something with it.

flist Files

But before we go much further, there's one bit of housecleaning we need to do. The `cnmake` script relies upon files called `flists` to help it know what SystemVerilog files to compile and what order to compile them in. Edit the file `verif/alutb/alutb.flist` and add your new test file to it, so that it looks like this:

```
1.      verif/alutb/alutb.flist
+incdir+../../verif/alutb
+incdir+../../verif/common
+incdir+../../verif/uvm_common
+incdir+../../verif/hdl
+incdir+../../verif/alutb/tests
../../verif/alutb/alutb_tb_top.sv
../../verif/alutb/alu_wrapper.sv
../../verif/alutb/tests/basic.sv
../../verif/alutb/tests/base_test.sv
../../verif/alutb/tests/fast_clk.sv
```

We've marked the change you need to make in **bold**. Each directory has an `flist` file in it, but as you'll see later you won't need to do this all the time.

Extended from the Base Test

Now, take a look at your new test, `fast_clk.sv`. Pretty, isn't it? You'll notice that it extends from the class called `base_test_c`. In fact, all tests descend from this test, but it's not a test at all. Well, it is, but it's not one that does any testing.

Actually, `base_test_c` will be the tippy-top of the component hierarchy in every simulation you run. It's the one that instantiates all of the same components that (most) tests will share, so they don't have to do the same thing over and over.

Go ahead and take a gander at `base_test.sv`. It looks a lot like `fast_clk_test_c`, except it has *stuff* in it. As the coding guidelines say, it has all of its fields at the top, and all of its methods at the bottom. The phase methods are declared in the order that they will run. Each section has a `Group:` comment, each field has a `var:` comment, and each method has a `func:` comment. These are parsed by the documentation generator to make all that pretty documentation.

Changing the Clock Frequency

You're probably eager to actually write some code, so let's get to it. In `base_test_c`, look for the instantiation of the clock driver, `tb_clk_drv`, in the `build_phase` function. This driver is a component, and so can be configured directly through the `set_config`

functions. As you can see in the line just before it is instantiated, the base test has configured the period of the clock driver to be 2,000ps. Let's get it going faster in our new test.

Back in our `fast_clk_test_c` class, modify its `build_phase()` function to look like this:

```
2.      verif/alutb/tests/fast_clk.sv
      virtual function void build_phase(uvm_phase phase);
      super.build_phase(phase);
      set_config_int("tb_clk_drv", "period", 1800);
      endfunction : build_phase
```

This will re-configure the clock driver's period to be 1800ps. A detailed explanation of how this works will be given in Lesson 5. For now, know that it's important that these configuration calls happen during a component's build phase, and that they happen after you call `super.build_phase()`.

Now go ahead and simulate, and you should see a period of 1.8ns. Here's the command-line, which you'll execute from the `verif/alutb` directory:

```
verif/alutb> cmake sim TEST=fast_clk SIMU=VCS
```

See `cmake`'s help function for more information on how to use that tool.

Overriding Constraints

That's great, you say, but what if I want more randomness? What if I want a clock period between 1.8ns and 2ns?

Well, that's what constraints are great for. A component is a SystemVerilog class, and the `period` field is a `rand int`, so it can be randomized anytime after it is `new`'ed. So, let's change our `fast_clk_test_c::build_phase` to this instead:

```
3.      verif/alutb/tests/fast_clk.sv
      virtual function void build_phase(uvm_phase phase);
      super.build_phase(phase);
      tb_clk_drv.randomize() with {
          period inside {[1800:1999]};
      };
      `cn_info($psprintf("Selected a period of %dps", tb_clk_drv.period))
      endfunction : build_phase
```

Note our first use of the Cavium-specific messaging macro ``cn_info`. These work just like their corresponding macros in TestBuilder, but you need to call `$psprintf` to get formatted printing.

Go ahead and run that, but add `NOBLD=lib` to your command-line. This prevents `cmake` from re-compiling the C code, which hasn't changed at all:

```
verif/alutb> cmake sim TEST=fast_clk SIMU=VCS NOBLD=lib
```

Question

Didn't work, did it? Instead, you got a period of 2,000ps, which is what the base test specified.

Do you know why? Do you know how to fix it?

For the answer, see the next page.

Answer

Here's what happens:

1. The base test's build phase stuffs a period of 2,000ps into the resource database.
2. It then creates the `tb_clk_drv` component.
3. Then, we randomize the `tb_clk_drv` component and get a nice random value.
4. Later, the `tb_clk_drv` component's build phase runs, fetches the value of 2,000ps, and assigns it to the period.

So how do we fix it? Simple. We randomize it in a later phase. Either the `connect`, `end_of_elaboration`, or `start_of_simulation` phases will do, since they run after the `build_phase`.

You can move your randomize call from `build_phase` to `end_of_elaboration` instead. Re-simulate and check your results.

Conclusion

This was a dive-right-in lesson that familiarized you with the UVM template generator, flist files, writing a little bit of SystemVerilog, and running using `cnmake`.

Lesson 2: Class Factories

In this lesson, you will learn to create a derived driver class and use the factory to instantiate it instead of the base class.

Extending the Driver

The goal of this lesson is to create a clock driver that has a 75/25 duty cycle. The generic clock driver, located in `verif/vcomponents/cn/cn_clk_drv.sv`, has lots of optional but complex functionality such as jitter and drift that we will not need. Instead, we want to create a simple clock driver that has a '1' for 75% of its period and a '0' for the remaining 25%.

To start, head over to the `verif/alutb` directory and create a generic component using `utg`:

```
verif/alutb> utg component -n clk_duty_cycle -f -q  
++ Creating clk_duty_cycle.sv
```

This doesn't precisely do what you want. When you open the file, you'll see that there are still template variables all around because we specified `utg` to run quietly (`-q`). It also didn't know that we wanted to extend from an existing driver. But, using `utg` is still much faster than writing all of this boilerplate code. And, since we're creating this driver in the testbench itself, there is no package to reference. So you also need to get rid of all `<pkg_name>` things.

You'll want to extend this not from `uvm_component` but from the generic clock driver instead, which being in a separate package will need the scope operator (`cn_pkg::`).

Problem 2-1

Do your best to modify the `clk_duty_cycle_c` to drive a clock with a 75/25 duty cycle. Consider the variables `period`, `init_delay`, `init_value`, and `init_x` as defined in the base class in your algorithm.

You won't be able to simulate it just yet.

(See answer on the following page).

Solution

Here is one possible solution:

```

4.      verif/alutb/clk_duty_cycle.sv
// class: clk_duty_cycle_c
// A clock with a duty cycle of 75/25.
class clk_duty_cycle_c extends cn_pkg::clk_drv_c;
    `uvm_component_utils(clk_duty_cycle_c)

    //-----
    // Group: Methods
    function new(string name="clk_duty_cycle",
                  uvm_component parent=null);
        super.new(name, parent);
    endfunction : new

    //////////////////////////////////////
    // func: run_phase
    virtual task run_phase(uvm_phase phase);
        int uptime = 3*period_ps / 4;
        int downtime = period_ps - uptime;

        // set to initial value
        clk_vi.clk = (init_x)? 'bx : init_value;
        #(init_delay_ps * 1ps);

        forever begin
            clk_vi.clk = 1;
            #(uptime * 1ps);
            clk_vi.clk = 0;
            #(downtime * 1ps);
        end
    endtask : run_phase

endclass : clk_duty_cycle_c

```

Note that all of the empty methods have been deleted. You could keep them there in case you need them later, with no effect on performance. However, one method that you can never remove is the constructor function, `new`. This function must always be present in derived UVM classes.

Using The Factory

The base test creates the testbench clock driver (`tb_clk_drv`) using the factory overridable `create` method:

```

// Create the clock driver
set_config_string("tb_clk_drv", "intf_name", "tb_clk_vi");
set_config_int   ("tb_clk_drv", "period",    2000);
tb_clk_drv = cn_pkg::clk_drv_c::type_id::create("tb_clk_drv", this);

```

This is essentially the same as calling `new()`, but if there were earlier factory overrides, then they would be used. It is usually best to never call `new()` on a UVM object, but to call `create()` instead, as it provides this flexibility for free.

Problem 2-2

Create a test called `duty_cycle_test_c` that overrides the generic clock driver with your new clock driver. Run this test and see your new duty cycle.

Solution

If you were successful, the signal `alutb_tb_top.tb_clk` should now have the new duty cycle. This is one possible solution:

```
5.      verif/alutb/tests/duty_cycle.sv
`ifndef __DUTY_CYCLE_SV__
`define __DUTY_CYCLE_SV__

    `include "base_test.sv"
    `include "clk_duty_cycle.sv"

    // class: duty_cycle_test_c
    // Run using the clk_duty_cycle_c instead of cn_pkg::clk_drv_c
    class duty_cycle_test_c extends base_test_c;
        `uvm_component_utils(duty_cycle_test_c)

        //-----
        // Group: Methods
        function new(string name="duty_cycle_test",
                     uvm_component parent=null);
            super.new(name, parent);
        endfunction : new

        //////////////////////////////////////
        // func: build_phase
        oovirtual function void build_phase(uvm_phase phase);
            set_type_override_by_type(cn_pkg::clk_drv_c::get_type(),
                                   clk_duty_cycle_c::get_type(), .replace(1));
            super.build_phase(phase);
        endfunction : build_phase

    endclass : duty_cycle_test_c

`endif // __DUTY_CYCLE_SV__
```

There are two keys here. The first is that the `clk_duty_cycle.sv` file must be included by this file, because it is a dependency. The second is that the `set_type_override_by_type` function call must precede the call to `super.build()`. Do you know why?

You could also use `set_type_override_by_name` instead when you want to override a specifically named clock.

Problem 2-3

How could you run the earlier `fast_clk` test that you wrote with your new clock driver but *without* needing to modify any code at all?

Solution

UVM offers a set of options that allow you to perform a `set_type_override` from the command-line. Running your `fast_clk` test with this command-line will do it:

```
6.  
verif/alutb> cnmake sim SIMU=VCS TEST=fast_clk  
SIMOPTS+=+uvm_set_type_override=clk_drv_c,clk_duty_cycle_c
```


Conclusion

This lesson introduced what class factories are capable of and had you write a smattering of SystemVerilog. Up next, you will start doing some real work on the ALU testbench.

The ALU Protocol

The ALU module operates as a simple arithmetic logic unit. It has a 1-bit control signal, an 8-bit data input interface, and a 32-bit data output interface with a 1-bit output ready signal:

Input data and control are sampled on the rising edge of `sclk`. The first cycle in an input transaction contains the 8-bit operation on `alu_dat[7:0]` and the `alu_ctl` signal is high. `alu_ctl` is held low for the remainder of the transaction and may not go high again to signal a new transaction until after the `alu_ready` output signal has been



sampled high. The `alu_dat[7:0]` lines contain the remaining operands based on the ten operation types as per this chart:

Operation	Cycle 0	Cycle 1	Cycle 2	Cycle 3	Cycle 4	Formula	Notes
ADD_A_B	8'h0	a[15:8]	a[7:0]	b[15:8]	b[7:0]	$k * (a+b) + c$	
SUB_A_B	8'h1	a[15:8]	a[7:0]	b[15:8]	b[7:0]	$k * (a-b) + c$	$a > b$
SUB_B_A	8'h2	a[15:8]	a[7:0]	b[15:8]	b[7:0]	$k * (b-a) + c$	$b > a$
MUL_A_B	8'h3	a[15:8]	a[7:0]	b[15:8]	b[7:0]	$k * (a*b) + c$	
DIV_A_B	8'h4	a[15:8]	a[7:0]	b[15:8]	b[7:0]	$k * (a/b) + c$	$b \neq 0$
DIV_B_A	8'h5	a[15:8]	a[7:0]	b[15:8]	b[7:0]	$k * (b/a) + c$	$a \neq 0$
INC_A	8'h6	a[15:8]	a[7:0]	N/A	N/A	$k * (a+1) + c$	
INC_B	8'h7	b[15:8]	b[7:0]	N/A	N/A	$k * (b+1) + c$	
CLR_RES	8'h8	N/A	N/A	N/A	N/A	0	Sets previous result to zero
ACCUM	8'h9	a[15:8]	a[7:0]	N/A	N/A	$k*(a+result) + c$	Uses previous result value

As you can see, some transactions take fewer cycles than others. The result is valid on the assertion of the `ready` signal, and the number of cycles that takes also depends on the operation chosen.

The `k` and `c` values are the 8-bit values programmed into the CSR named `ALU_CONST`. The sum of all results since it was last read should be in the CSR `ALU_RESULT[SOR]`, which is a read-clear CSR.

Your job in future lessons will be to verify this simple sub-module using an agent that you create, collect functional coverage, and maybe even find a bug or two.

ALU_CONST = 0x7E8				
Bit	Field Name	Type	Reset	Description
<63:16>	NS	RAZ	NS	Reserved
<15:8>	K_VAL	R/W	1	The constant value that is first multiplied to the results of all ALU computations before the C_VAL is added in.
<7:0>	C_VAL	R/W	0	The constant value that is added to the results of all ALU computations.

ALU_RESULT = 0x7F0				
Bit	Field Name	Type	Reset	Description
<63:32>	NS	RAZ	NS	Reserved
<31:0>	SOR	RC	0	The sum of all results so far.

Lesson 3: Creating a Transaction

It is often considered good OOP practice to begin with figuring out what the data should look like. We'll create a transaction that can be generated randomly by a sequence, sent to the driver via a sequencer, and monitored by a monitor. The predictor will also be able to receive these transaction classes as will a functional coverage monitor. It would also be nice to be able to pack and unpack these transactions off the wire.

We will put all of our ALU-related code in a separate vcomponent. Normally--but not always--a vcomponent is composed of *all* reusable code related to a single testbench, or used in multiple testbenches (i.e. the NCB and RSL agents). **The `alutb` vcomponent is where this should live. But for the purposes of this tutorial, we will create all ALU related code in a separate vcomponent.**

With these goals in mind, let's get started.

Creating a vcomponent

A vcomponent is meant to be a single reusable package that can be individually compiled with the help of other vcomponents. For a detailed discussion on what a vcomponent is and what it isn't, see the [SV Coding Guidelines](#) on the wiki.

To create a vcomponent, you can again use `utg` in the `verif/vcomponents` directory:

```
verif/vcomponents> utg vcomponent -n alu
++ Creating alu_pkg.sv
++ Created vcomponent alu
++ Exiting.
```

This creates the `alu` directory, the `alu_pkg.sv` package file, and an `alu.flist` file for you. Other vcomponents that this package relies upon will be placed ahead of this one in the build order. Then, the files needed to create the ALU package will be ``included` inside the package file. You will not need to touch the `alu.flist` file at this time.

To incorporate this new vcomponent into the ALUTB testbench, you need to modify the `alutb/Makefile` and add the path to this vcomponent's `flist` file to the `FLISTS` variable. The order in which it is added is important based on this vcomponent's dependencies on other vcomponents. This vcomponent will depend upon `uvm`, `cn`, and `global` `flists`, so add it after those. The `alutb` testbench depends upon this vcomponent, so it needs to be added before that `flist` as well.

The `FLIST` variable in the file `alutb/Makefile` should now be set as follows:

```
7.      verif/alutb/Makefile
FLISTS= verif/vcomponents/cn/cn.flist \
        verif/vcomponents/global/global.flist \
```

```

verif/vcomponents/reg/reg.flist \
verif/vcomponents/ncb/ncb.flist \
verif/vcomponents/alutb/alutb.flist \
verif/vcomponents/alu/alu.flist \
verif/alutb/alutb.flist \
verif/alutb/rtl.flist

```

A Sequence Item is a Transaction

The class `uvm_sequence_item` derives directly from `uvm_transaction`, but has the additional property of being able to be driven by a sequence through a sequencer and on to a driver, which are components that we'll discuss in the near future. It is a rule that we must derive our new class from `uvm_sequence_item`.

Create the sequence item file by specifying the `item` template. Call it 'item', and have it saved to the file named `alu_item.sv`:

```

8.
verif/vcomponents/alu> utg item -n item --filename alu_item.sv -q
++ Creating alu_item.sv

```

The first field in our transaction is the operation, which is a good place for an enumerated type. Enumerated types need to be specified before the call to ``uvm_field_enum`, so put the following at the very top of the class:

```

9.      verif/vcomponents/alu/alu_item.sv
//-----
// Group: Types
typedef enum bit [7:0] {ADD_A_B = 0,
                        SUB_A_B = 1,
                        SUB_B_A = 2,
                        MUL_A_B = 3,
                        DIV_A_B = 4,
                        DIV_B_A = 5,
                        INC_A   = 6,
                        INC_B   = 7,
                        CLR_RES = 8,
                        ACCUM   = 9,
                        } operation_e;

```

Notice that we declared it with the `bit [7:0]` syntax to tell SystemVerilog how wide it should be. This will be useful when we pack and unpack this transaction. Also notice that we provided the suffix of `_e` to help identify it as an enumerated type, as required in the [SV Coding Guidelines](#).

Now, we add it to the list of fields with the ``uvm_field_enum` macro:

```

10.      verif/vcomponents/alu/alu_item.sv
`uvm_object_utils_begin(alu_pkg::item_c)
`uvm_field_enum(operation_e, operation, UVM_ALL_ON)
`uvm_object_utils_end

```

We chose to give it the flag `UVM_ALL_ON` because we want UVM to pack it, print it, and all the other fancy stuff that UVM provides for free.

Now, we need to declare the operation as a member field of the class. Put this in the Fields group:

```
11.      verif/vcomponents/alu/alu_item.sv
      //-----
      // Group: Fields
      // var: operation
      rand operation_e operation;
```

We declared it as a random field because we're going to want to produce lots of random transactions later on.

Now, let's do what we need to do for variables `a` and `b`:

```
12.      verif/vcomponents/alu/alu_item.sv
      `uvm_object_utils_begin(alu_pkg::item_c)
      `uvm_field_enum(operation_e, operation, UVM_ALL_ON)
      `uvm_field_int (alpha,                UVM_ALL_ON | UVM_NOPACK)
      `uvm_field_int (beta,                 UVM_ALL_ON | UVM_NOPACK)
      `uvm_object_utils_end
```

We chose to call them `alpha` and `beta` because `a` and `b` are not very good variable names. We also specified that they are not to be packed. This, as you'll see later, is important when we go to pack and unpack these transactions, because `A` and `B` do not appear in all operation types.

Declare the `alpha` and `beta` fields as `rand bit [15:0]`.

Typedefs Rule

At some point in the life of this ALU, somebody's going to want to make a bus wider. Or narrower. Or they're going to add a square root operation. Or they're going to turn it into a router that sorts packets by size. Whatever happens, it will almost certainly change somehow at some point. And when it does, you're going to have to change it in a *lot* of places. Using constants would require you to only make the change in one place. And, defining constants is much safer because it prevents typos from creating a nightmare debugging session.

However, defining constants for bus widths using ``defines` are so passé. Instead, we're going to create a typedef for the result type. All typedefs end with `_t`:

```
13.      verif/vcomponents/alu/alu_pkg.sv
      typedef logic [31:0] result_t;
```

That will also be much easier to type all over the place than this would:

```
logic [`RESULT_WIDTH-1:0] result;
```

You can place your typedefs in the vcomponent's package (above all of the included files that depend on it), or you can create an `alu_types.sv` file and ``include` it in any file


that needs it. The choice is yours. Either way, the type will be local to the vcomponent's package, so there's no concern for namespace pollution.

Create typedefs for any type that will be commonly used within a package, including enumerated types. The `operation_e` enumerated type was placed in the `item_c` class, but it could just as easily have been placed in the package. You may also, for example, create types for the operands `alpha` and `beta`.

Sequence Item U-Turns

We're also going to put the `result` field in the transaction, like so:

```
14.    verif/vcomponent/alu/alu_item.sv
    // var: result
    // This is the result of the operation, filled in by the driver and sent back with the response
    result_t result;
```



When the transaction is sent to the driver by the sequencer, the driver will (eventually) send a response back. Sometimes, a block gives no response (for one-way communications, for example). In cases such as those, it is common for the driver to send the same request item back to the sequencer after the transaction has been completely driven.

In our case, there is a response. But it's just a 32-bit integer. Both request and response types have to be derived from `uvm_sequence_item`, to ensure that they have all the functionality that the sequencer needs to track them. So, our design will simply have the driver put the result back into the request transaction and send it back once it's been received.

Adding Constraints

You can't divide by zero. And the transaction rules don't allow for negative numbers during subtractions, either. Constraints are the way to do this. Fortunately, SystemVerilog really shines in this department.

Problem 3-1

Exercise your SystemVerilog muscles to add constraints such that all randomized transactions follow the protocol.

(Answer next page).

Solution

There are several ways you can approach this, but this is probably the clearest and most concise way:

```
15.      verif/vcomponents/alu/alu_item.sv
      // ensure that all operands have legal values
      constraint protocol_cnstr {
        (operation == DIV_A_B) -> beta != 0;
        (operation == DIV_B_A) -> alpha != 0;
        (operation == SUB_A_B) -> alpha > beta;
        (operation == SUB_B_A) -> beta > alpha;
      }
```

Note that all constraint blocks need a name, and they should end in `_cnstr`. Also, you should put your constraints near the field declarations themselves, in the Fields group.

Printing Your Classes

Printing your UVM classes to a logfile is something you'll want to do all over the place. All UVM classes have a method that returns the class as a string--`sprint`. This function takes an optional argument of type `uvm_printer` which you can use to define what you want your string to look like. This class is called a policy class, and they are sprinkled all over UVM. You then pass this string to one of the ``cn` messaging macros, ``cn_info`, ``cn_dbg`, ``cn_warn`, ``cn_err`, or ``cn_fatal`.

Table Format

Table format is the default, so you don't need to provide a printer class to the `sprint` function. The following code snippet will produce the table format of a class:

```
`cn_info($psprintf("Monitored:\n%s", item.sprint()))
```

And the results are shown below:

%I-(alu_mon.sv: 140) [alutb_env.alu_agent.mon] { 38597ns} Monitored:			
Name	Type	Size	Value
mon_item	alu_pkg::item_c	-	@6203
operation	operation_e	8	DIV_A_B
alpha	integral	16	'hc6eb
beta	integral	16	'h8a74

By default, your printed class uses the table printer, which prints your class in a nice, easy-to-read format. However, it is expensive in simulation time. For the thrifty, consider using the tree format instead.

Tree Format

To print in the tree format, use the following instead.

```
`cn_info($psprintf("Monitored:\n%s", item.sprint(uvm_default_tree_printer)))
```

And it appears like this in the logfile:

```
%I-( alu_mon.sv: 142) [alutb_env.alu_agent.mon ] { 38573ns} Monitored:
mon_item: (alu_pkg::item_c@6195) {
  operation: DIV_A_B
  alpha: 'hcc3d
  beta: 'h38de
}
```

The `uvm_default_*_printer` objects are globally scoped instances of the `uvm_printer` policy classes that can be used when you don't want to type too much. Alternatively, you can create your own tree printer instance and use it over and over again:

```
uvm_tree_printer tp = new();
`cn_info($psprintf("Monitored:\n%s", item.sprint(tp)))
```

The benefit of this approach is that you can tweak a variety of printer knobs to get your class just as pretty as you like. The variety of knobs available is impressive. See the UVM reference manual's description of the `uvm_printer_knobs` class for details.

The same approach can be used for the table and line printer classes, too.

Line Format

The cheapest formatting is the line printer:

```
`cn_info($psprintf("Monitored:\n%s", item.sprint(uvm_default_line_printer)))
```

And its appearance is as shown here:

```
%I-( alu_mon.sv: 142) [alutb_env.alu_agent.mon ] { 38597ns} Monitored:
mon_item: (alu_pkg::item_c@6203) { operation: DIV_A_B alpha: 'hc6eb beta: 'h8a74 }
```

Convert2String

It's very handy to single-line print things--especially small transactions like these. But UVM's standard `uvm_line_printer` format is not terribly readable. The `convert2string()` function is defined by UVM as an empty function and is Cavium's place to do this, and `$psprintf` is the right tool for the job:

```
16.      verif/vcomponents/alu/alu_item.sv
      //////////////////////////////////////
      // func: convert2string
      // Single-line printing
      virtual function string convert2string();
      convert2string = $psprintf("%s A:%04X B:%04X", operation, alpha, beta);
```

```
endfunction : convert2string
```

You can optionally make your function fancier by not printing out the A and B operands for operations that don't use them, but this will suffice for now. Now, your transaction will print out nice and neat and just the way you like it:

```
%I-( alu_drv.sv: 107) [alutb_env.alu_agent.drv ] { 890ns} Monitored: DIV_A_B A:C6EB B:8A74
```

Packing and Unpacking

There are many routes that will get us to our end goal, but this is a great opportunity to learn about how to pack and unpack in UVM. Later, the driver will be able to pack this transaction down to the bytes that it will put on the wire, and the monitor will be able to collect the bytes and unpack them.

When you call the pack function on a transaction, it will automatically pack all of the fields that are marked by the flags in the macros as packable. Because the alpha and beta fields will not always be used, though, we'll need to *conditionally* pack them. To add this functionality to a UVM object, you override the function `do_pack`:

```
17.      verif/vcomponents/alu/alu_item.sv
      //////////////////////////////////////
      // func: do_pack
      virtual function void do_pack(uvm_packer packer);
      super.do_pack(packer);

      if(operation inside {[ADD_A_B : INC_A], ACCUM})
        packer.pack_field_int(alpha, 16);
      if(operation inside {[ADD_A_B : DIV_B_A], INC_B})
        packer.pack_field_int(beta, 16);
      endfunction : do_pack
```

We've chosen to use SystemVerilog's handy `inside` operator to conditionally pack these ints as 16-bit values.

Like the `sprint` function, the `pack` and `unpack` functions take an optional policy class called `uvm_packer`. Here, we let the policy class do the work for us, in case the caller of `pack` or `unpack` chose to modify the default policy to handle things different.

The following example shows how the user can pack into a list of bytes in little-endian format instead of the default big-endian format:

```
uvm_packer little_endian = new();
byte unsigned stream[];

little_endian.big_endian = 0;
item.pack_bytes(stream, little_endian);
```

Problem 3-2

With the `do_pack` method as a boilerplate, write the corresponding `do_unpack` method.

Solution

The solution is fairly straightforward:

```
18.      verif/vcomponents/alu/alu_item.sv
      //////////////////////////////////////
      // func: do_unpack
      virtual function void do_unpack(uvm_packer packer);
          super.do_unpack(packer);

          if(operation inside {[ADD_A_B : INC_A], ACCUM})
              alpha = packer.unpack_field_int(16);
          if(operation inside {[ADD_A_B : DIV_B_A], INC_B})
              beta = packer.unpack_field_int(16);
      endfunction : do_unpack
```

Problem 3-3

Let's now fill in the test `basic_test_c` such that during the `main_phase` it creates a stream of 50 random ALU transactions, prints them out, and packs them into a list of unsigned bytes. Then, unpack this list of bytes into another transaction, and print it out. Finally, compare the two transactions and print an error if they do not compare.

Use ``cn_info` and ``cn_err` to print informational messages and error messages, respectively. You can also use the function `cn_pkg::print_ubyte_array()` to print out your byte array (or `print_byte_array` if you used signed bytes), to aid in debugging your work.

What happens?

Solution

Here is one possible solution:

```
19.    verif/alutb/tests/basic.sv
class basic_test_c extends base_test_c;
    `uvm_component_utils(basic_test_c)

    //-----
    // Group: Methods
    function new(string name="test",
                  uvm_component parent=null);
        super.new(name, parent);
    endfunction : new

    //////////////////////////////////////
    // func: run_phase
    virtual task main_phase(uvm_phase phase);
        byte unsigned stream[];
        alu_pkg::item_c item;
        alu_pkg::item_c unp_item = alu_pkg::item_c::type_id::create("unp_item");

        phase.raise_objection(this);

        repeat(50) begin
            item = alu_pkg::item_c::type_id::create("item");
            item.randomize();
            `cn_info($psprintf("Created ALU transaction: %s", item.convert2string()))
            item.pack_bytes(stream);
            `cn_info($psprintf("Bytes: %s", cn_pkg::print_ubyte_array(stream)))
            unp_item.unpack_bytes(stream);
            `cn_info($psprintf("Unpacked: %s", unp_item.convert2string()))
            if(item.compare(unp_item) == 0)
                `cn_err($psprintf("Miscompare!"))
        end

        phase.drop_objection(this);
    endtask : main_phase

endclass : basic_test_c
```

Your test will almost certainly fail for the increment operations. But why?

Well, we're both randomizing and comparing the `alpha` and the `beta` fields unconditionally, yet these fields will be neither packed nor unpacked, so will be different for the increment operations which do not use them all.

Problem 3-4

This problem can be solved by overriding the `do_compare()` method, in the same manner as the `do_pack()` and `do_unpack()` methods. What else do you need to do?

(See *answer next page*).

Solution

You'll need to add the flag `UVM_NOCOMPARE` to the `alpha` and `beta` field macro calls. If everything worked correctly, your test should PASS. Notice that the `do_compare` function must first `$cast` the object passed to it to the `item_c` type before you can look at its `alpha` or `beta` member fields.

The complete class looks like this:

```
20.      verif/vcomponents/alu/alu_item.sv
// class: item_c
// An ALU Transaction as a sequence item
class item_c extends uvm_sequence_item;
//-----
// Group: Types
typedef enum bit [7:0] {
    ADD_A_B = 0,
    SUB_A_B = 1,
    SUB_B_A = 2,
    MUL_A_B = 3,
    DIV_A_B = 4,
    DIV_B_A = 5,
    INC_A   = 6,
    INC_B   = 7,
    CLR_RES = 8,
    ACCUM   = 9,
} operation_e;

`uvm_object_utils_begin(alu_pkg::item_c)
    `uvm_field_enum(operation_e, operation, UVM_ALL_ON)
    `uvm_field_int (alpha,                  UVM_ALL_ON | UVM_NOPACK | UVM_NOCOMPARE)
    `uvm_field_int (beta,                   UVM_ALL_ON | UVM_NOPACK | UVM_NOCOMPARE)
`uvm_object_utils_end

//-----
// Group: Fields

// var: operation
rand operation_e operation;

// var: A variable
rand bit [15:0] alpha;

// var: B variable
rand bit [15:0] beta;

// ensure that all operands have legal values
constraint protocol_cnstr {
    (operation == DIV_A_B) -> beta != 0;
    (operation == DIV_B_A) -> alpha != 0;
    (operation == SUB_A_B) -> alpha > beta;
    (operation == SUB_B_A) -> beta > alpha;
}

// var: result
// This is the result of the operation, filled in by the driver and sent back with the response
result_t result;

//-----
// Group: Methods
function new(string name="item");
    super.new(name);
endfunction : new

//-----
```

```

// func: convert2string
// Single-line printing
virtual function string convert2string();
    convert2string = $psprintf("%s A:%04X B:%04X", operation, alpha, beta);
endfunction : convert2string

////////////////////////////////////
// func: do_pack
virtual function void do_pack(uvm_packer packer);
    super.do_pack(packer);

    if(operation inside {[ADD_A_B : INC_A], ACCUM})
        packer.pack_field_int(alpha, 16);
    if(operation inside {[ADD_A_B : DIV_B_A], INC_B})
        packer.pack_field_int(beta, 16);
endfunction : do_pack

////////////////////////////////////
// func: do_unpack
virtual function void do_unpack(uvm_packer packer);
    super.do_unpack(packer);

    if(operation inside {[ADD_A_B : INC_A], ACCUM})
        alpha = packer.unpack_field_int(16);
    if(operation inside {[ADD_A_B : DIV_B_A], INC_B})
        beta = packer.unpack_field_int(16);
endfunction : do_unpack

////////////////////////////////////
// func: do_compare
virtual function bit do_compare(uvm_object rhs,
                               uvm_comparer comparer);

    item_c_rhs;
    $cast(_rhs, rhs);
    do_compare = super.do_compare(rhs, comparer);
    if(operation inside {[ADD_A_B : INC_A], ACCUM})
        do_compare &= comparer.compare_field_int("alpha", alpha, _rhs.alpha, 16);
    if(operation inside {[ADD_A_B : DIV_B_A], INC_B})
        do_compare &= comparer.compare_field_int("beta", beta, _rhs.beta, 16);
endfunction : do_compare

endclass : item_c

```

Printing, Comparing, Packing...Recording?

You've already figured out how to customize a UVM object's comparison mechanism by overriding the `do_compare` function. You'll note that it, too, takes an optional policy class. As you've seen, UVM uses these policy classes together with the `do_*` functions and the UVM field flags such as `UVM_NOCOMPARE` to customize your classes to operate exactly the way you want them to.

UVM also includes a recording mechanism, but what does it do? The UVM reference says that it's "vendor-specific." In fact, it can be used by the simulation tool, the waveform dumper, or any other third-party tool that might come along. You could override it to write itself to a database for further analysis, for example.

Adding Coverpoints

In a coverage driven verification environment, *now*--when you've just created the transaction--is the time to create coverpoints. Start collecting functional coverage early in the project, and collect it often. Creating metrics to chart the progress of the verification environment as early as possible is very beneficial to long-term planning. Plus, you may later be pleasantly surprised at what you're already covering now that you won't have to write a test for someday.

Because this is a UVM tutorial, though, we will not be going over how to create covergroups and coverpoints.

Conclusion

In this lesson, we created a transaction object that represents an ALU request, and we derived it from `uvm_sequence_item`, for reasons we'll explore later. We also briefly touched on how to customize the packing, printing, comparing, and recording of any UVM object. The level of customization that UVM provides for this is deep, and you should refer to the UVM Reference Manual when you next want to pursue these topics.

Lesson 4: Creating an Interface

A few fun facts about interfaces that you may have missed during your SystemVerilog 101 training:

- Interfaces are *not* just a bundle of wires. They are really more like classes or modules. This is because they can contain their own tasks and functions, their own member fields, assertions, procedural blocks (initial and always statements), assignments, checkers, enumerated types...the list of things you can stuff into an interface is very long. *And* they can be passed between RTL modules as a synthesizable construct (that bundle of wires analogy). Unlike classes, though, they are not dynamically created at runtime, but statically created as part of the RTL hierarchy just as modules, regs, and wires are.
- Interfaces cannot be defined in a package. As much as you'd like to define your ALU interface in your ALU package, the language won't allow it. Instead, the interface must be *included* and *known* to the compiler before it compiles any packages that depend on it because...
- Testbench code (i.e. that which lives in the `verif/alutb` directory) can access hard-coded paths such as `alutb_tb_top.ncb_i.inb.src[7:0]`. Packages--such as those which reside in `vcomponent` directories--cannot. Doing so would break the reusability aspect of a `vcomponent`, because in the full-chip testbench, the path `alutb_tb_top` does not exist. Instead, *all* signals accessed by testbench code within a package must come via a *virtual interface*.
- The UVM method of getting a class in a package to see a testbench interface is to put a reference to the interface in the UVM resource database. This reference is placed in the RTL-side of the testbench and is then pulled out of the database by the UVM class's `build` phase.
- An interface may have input (or output) signals, typically a clock and a reset signal. These signals are not specific to the interface itself, because they may be shared with other devices (i.e. `sclk`, `srst_n`). If a signal is specific to the interface (i.e. `pcie_clk`, `srio_clk`, etc.) then it should be a member of that interface.
- Interfaces may contain other interfaces. If a single interface has 4 separate 'channels', then it is often convenient to create a single channel interface, and repeat it 4 times in a container interface. The `rsl_intf` defined in `verif/vcomponents/rsl/rsl_intf.sv` is an example of such a hierarchy.

- Clocking blocks offer you the ability in one place to specify the clock edge on which signals will be sampled or driven. By using different clocking blocks for each component that will access a given signal, you can easily change these edges in one place.
- Clocking blocks also solve race conditions between the RTL and the verification environment that were previously solved in SystemVerilog by the `program/endprogram` constructs.
- Modports offer access protection against modules, drivers, and monitors from touching signals that they shouldn't. In a way, they act like protected access in C++. Given that an interface can have a task or function that accesses signals, it could be easy for a monitor that shouldn't touch any signals to accidentally start affecting things.
- You don't *need* to use clocking blocks and modports, as they add additional typing when accessing signals. You'll have to judge whether their benefits outweigh this additional overhead. We'll use both of them in our example just so you can see how they're used.

Problem 4-1

Take a stab at creating an interface for the ALU signals. Use the `utg` template `intf`. The filename should be `alu_intf.sv`, and it should reside in the `verif/vcomponents/alu` directory. The interface itself should be named `alu_intf`. Write a `reset()` function that sets the control and input data signals to zero. A clock and an active-low reset signal should be inputs to the interface.

You can use the files `verif/vcomponents/rsl/rsl_intf.sv` or `verif/vcomponents/ncb/ncb_intf.sv` to help you.

Extra credit: add assertions that ensure that there are no X's when there shouldn't be. The ``cn_err` macro cannot be used in an interface (it can only be used in verification environment code that derives from `uvm_object` because these macros call the `get_full_name` function). You must use the `$cnErr` function instead.

Solution

Admittedly this one was probably a lot to ask you for, and you don't yet have a way to simulate it to see if you're right. But hopefully the effort was worthwhile:

```
21.      verif/vcomponents/alu/alu_intf.sv
// class: alu_intf
interface alu_intf(input logic clk,
                  input logic rst_n);

    //-----
    // Group: Signals

    // var: ctl
    // Asserted only on the first cycle of a new transaction, while dat contains the operation
    logic      ctl;

    // var: dat
    // The input data nibble that contains the operation and operands
    logic [7:0] dat;

    // var: ready
    // The output of the ALU that indicates when the result data is valid
    logic      ready;

    // var: result
    // The 32-bit result data
    logic [31:0] result;

    //-----
    // Group: Clocking blocks

    // var: drv_cb
    // A clocking block that represents how the environment's driver sees the interface
    clocking drv_cb @(posedge clk);
        output    ctl;
        output    dat;
        input     ready;
        input     result;
        input     rst_n;
    endclocking : drv_cb

    // var: mon_cb
    // A clocking block that represents how the environment's monitor sees the interface
    clocking mon_cb @(posedge clk);
        input     ctl;
        input     dat;
        input     ready;
        input     result;
        input     rst_n;
    endclocking : mon_cb

    //-----
    // Group: Modports

    modport drv_mp(clocking drv_cb,
                  import reset);
    modport mon_mp(clocking mon_cb);

    //-----
    // Group: Methods

    // func: reset
    // Convenience function for the driver to reset its outputs
    function void reset();
        ctl = 0;
        dat = 8'b0;
    endfunction
endinterface
```

```

endfunction : reset

//-----
// Group: Assertions

// Ensure that no signals are ever an X when they shouldn't be
always @(posedge clk) begin
    if(rst_n) begin
        assert(!$isunknown(ctl)) else $cnErr("Signal ctl is an X.");
        assert(!$isunknown(dat)) else $cnErr("Signal dat is an X.");
        assert(!$isunknown(ready)) else $cnErr("Signal ready is an X.");
        if(ready)
            assert(!$isunknown(result)) else $cnErr("Signal result is an X.");
    end
end

endinterface : alu_intf

```

The ALU Interface Explained

- Interface signals should be declared as `logic`, not `wire` or `reg`. The reason SystemVerilog needed to create the `logic` type was for this very reason: these signals could be driven either by a continuous assignment (like a `wire`) or with a procedural assignment (like a `reg`), depending upon the implementation to which it is connected. Thus the need for this new data type.
- We created 2 different types of clocking blocks and 2 modports: one for a monitor that drives no signals, and one for an environment driver that drives the `ctl` and `dat` signals. Only the driver will be allowed to call the `reset` function to clear the `ctl` or `dat` signals. If the RTL uses interfaces to wire itself up, more clocking blocks and modports might be added, too.

Connecting the Interface to the DUT

Now that you've defined the interface, you need to let the simulator see it. This is accomplished by adding the file to the `vcomponents/alu/alu.flist` file. It must be added before the ALU environment package, because the package depends on the interface, not the other way around:

```

22.    verif/vcomponents/alu/alu.flist
+incdir+../../verif/vcomponents/alu
+incdir+../../verif/vcomponents/alu/alu_intf.sv
+incdir+../../verif/vcomponents/alu/alu_pkg.sv

```

Now, we will instantiate a single ALU interface in the `alutb_tb_top` module after the NCB interface, and reset it at time zero so that we do not trip our assertions:

```

23.    verif/alutb/alutb_tb_top.sv
// obj: alu_i
// The <alu_intf> instance
alu_intf alu_i(.clk(tb_clk), .rst_n(tb_rst_n));
initial
    alu_i.reset();

```


You probably look at all the extra documentation comments as being unnecessary, and for such a simple testbench they probably are. But when things get more complicated or are written by others, you will appreciate having pretty documentation to look at. These comments help create that documentation.

In this testbench, the `alu` DUT is instantiated in the `alu_wrapper` module. Add this interface as an argument to that module as the other interfaces are:

```
24.      verif/alutb/alu_wrapper.sv
module alu_wrapper(input logic tb_clk,
                  tb_rst_n,
                  ncb_intf ncb_i,
                  alu_intf alu_i);
```

Strictly speaking, the wrapper module is not *necessary*. The purpose of this wrapper module is to consolidate the instantiation of the DUT and all of the wiring that is necessary such that the top-level testbench can be more easily read and understood.

In the future, perhaps our RTL modules will just accept the SystemVerilog interfaces. But this one doesn't. Instead, you need to wire up each signal to the interface directly. Fortunately, emacs' verilog-mode makes this easy with the `AUTO_TEMPLATE`. You can wire up each signal one-by-one, or you can use the regular expression tricks that the NCB interface makes use of, like this:

```
25.      verif/alutb/alu_wrapper.sv
.alu \(.*\)      (alu_i.\1[]),
```

Now, hitting `C-c C-a` wires everything up for you. You will need to do the same in the `alutb_tb_top`. You are using emacs, aren't you?

Assuming you've done everything correctly, you may now run one of your tests as a sanity check. Everything should still PASS because hooking the interface up to the DUT is (theoretically) innocuous.

Storing the Interface in the Database

The `pre_run_test()` function must appear in all testbenches, because it will be called by the code that sits in the `verif/hdl/tbv_common.v` file to ensure that all interfaces are set in the resource database before any agents try to get them. As you can see in the `alutb_tb_top` testbench, several interfaces have already been set into the database.

You'll want to put an instance of the driver (`drv_mp`) and the monitor (`mon_mp`) modports into the resource database, so that the driver and monitor that will be written later can get them. You will also want to virtualize these interfaces with the `virtual` keyword, because that is what the classes will use.

The name for the interfaces here must match the name of the interfaces that the driver and monitor use to get them. As you'll see later, making that name a configuration field of those components will allow you to attach different instances of components to different interfaces, all on a configurable basis.

Problem 4-2

Put virtual driver and monitor modports into the resource database in the `alutb_tb_top.pre_run_test()` function, using the other examples as models. Name the resource for both `alu_pkg::alu_intf`, and name the interfaces `drv_vi` and `mon_vi`, respectively.

Solution

This should do the trick:

```
uvm_resource_db#(virtual alu_intf.drv_mp)::set("alu_pkg::alu_intf", "drv_vi", alu_i.drv_mp);  
uvm_resource_db#(virtual alu_intf.mon_mp)::set("alu_pkg::alu_intf", "mon_vi", alu_i.mon_mp);
```

However, these handy macros simplify the process somewhat:

```
26.      verif/alutb/alutb_tb_top.sv  
`cn_set_intf(virtual alu_intf.drv_mp, "alu_pkg::alu_intf", "drv_vi", alu_i.drv_mp)  
`cn_set_intf(virtual alu_intf.mon_mp, "alu_pkg::alu_intf", "mon_vi", alu_i.mon_mp)
```

Conclusion

In SystemVerilog, testbench-related code is capable of reaching down into the hierarchy and peeking and poking at signals. All reusable components, though, are kept in packages that have no notion of the hierarchy. The bundle of signals that they can see and affect must be kept in an interface. As we'll see later, components in the environment can easily attach to a named interface via the resource database settings we just completed.

Lesson 5: Creating an Agent

In this lesson, we're going to create the ALU agent, driver, monitor, and sequencer. We will learn about the purpose of each of these types of components, and we will discuss in detail how all that configuration *magic* actually works.

Creating Multiple Files

Again, `utg` makes it a snap to create all of these files at one time:

```
verif/vcomponents/alu> utg agent drv mon sqr -n alu -f -q
++ Creating alu_agent.sv
++ Creating alu_drv.sv
++ Creating alu_mon.sv
++ Creating alu_sqr.sv
```

You'll want to add these files as ``includes` to the package. The order in which these files are included is inconsequential. In fact, you can do it in alphabetical order or in chronological order. What matters most is that each file's dependencies are ``included`, that you use ``include` guards, and that you manage any cyclical dependencies (which we'll see later). If you follow these rules of order, you should never have any build ordering issues.

Objects vs. Components

Let's review a few terms and what the purpose of each of these things is before going forward:

Objects

Everything is a `uvm_object`, because it is the root class of all other UVM-based classes. But when we specifically refer to something as an object, we usually mean it is something that is not also a component. It is dynamic in nature, being created and destroyed on-the-fly (like a transaction or a packet).

Component

A component derives from `uvm_object`, but is quasi-static in nature. While the simulator considers them to be dynamically created (as opposed to, say, modules, wires, and registers), they are created by UVM at time zero and live throughout the lifetime of the simulation.

Phases

Components have phases. First, they run their `new()` function as soon as they are created, as all classes do. Later, they run the `build_phase`, the `connect_phase`, the `end_of_elaboration_phase`, and the `start_of_simulation_phase`. All of these phases happen in zero time at time zero. Then, the `run_phase` happens, which has numerous task-based sub-phases, which consume time. See more about phases in Appendix A.

Objects that are *not* components do *not* have phases.

Driver

The driver is a component that receives transactions from the sequencer and wiggles the wires of an interface to stimulate the DUT, following the protocol of the interface it attaches to. It also re-routes responses to transactions back to the sequencer. Ideally, its purpose should be kept to that, and to that only.

Monitor

The monitor is a component that watches the interface it is attached to and reports what happens, checking the interface for protocol violations. It collects activity on the interface in the form of transactions and pushes them out an analysis port. It is agnostic in nature, in that it monitors both the driver and the DUT looking for errors and reporting activity.

Sequencer

A sequencer is a component that arbitrates among multiple streams of stimulus (sequences), sends sequence items to a driver one at a time, and re-routes responses back to the original sequence.

Agent

An agent is merely a container class that holds the sequencer, driver, and monitor. Its primary purpose is to provide a single configurable, re-usable component that you can plop down into your testbench, without having to deal with multiple components. It may hold other components as well, but serves no other specific purpose. It usually features TLM ports and imps (see Lesson 6) that send and receive transactions to other agents or components.

With those definitions in mind, let's take a look at what `utg` created for us. The file `verif/vcomponents/alu/alu_agent.sv` is already chock full of most of the stuff we need. It has one configuration field, `is_active`, which instantiates the driver and sequencer when set to `UVM_ACTIVE`. Otherwise, the agent is `UVM_PASSIVE`, in which case only the monitor is operating. This is typical for block-level agents that monitor interfaces in the full-chip testbench, but never actually drive stimulus. The build phase is where this all happens:

```
virtual function void build_phase(uvm_phase phase);
    super.build_phase(phase);

    mon = mon_c::type_id::create("mon", this);
    if(is_active) begin
        drv = drv_c::type_id::create("drv", this);
        sqr = sqr_c::type_id::create("sqr", this);
    end
endfunction : build_phase
```

A monitor might perform all of the checking for the agent on-the-fly. But agents *sometimes* do not contain scoreboards or more complex predictors, because they might inhibit their reusability. A block-level agent might be reusable from one chip design to the next, but the prediction algorithm might change. Instead, agents might push monitored traffic out of TLM ports and on to other listeners, be they scoreboards, functional coverage collectors, or other agents. We'll discuss prediction more in Lesson 9.

Sequencers communicate stimulus sequence items and responses to the driver via the `seq_item_port` of the driver and the `seq_item_export` of the sequencer. `utg` has already made this connection for you:

```
virtual function void connect_phase(uvm_phase phase);
    super.connect_phase(phase);
    if(is_active)
        drv.seq_item_port.connect(sqr.seq_item_export);
endfunction : connect_phase
```

We'll talk about these ports and exports in the next lesson, but you should be aware that the ports, the driver, and the sequencer, are all parameterized classes. They need to know what kinds of stimulus item (called the request) that the sequencer will be sending to the driver, and what kind of response item the driver will send back to the sequencer. Sometimes, a driver won't have a data item to send back, if for example the interface has one-way communication. In such a case, the driver will often just send back the request itself once it has completed sending it. Other times, the response will be embedded into the transaction. You'll recall that we added the result field to our transaction for just such a purpose.

So, the driver and sequencer need to be parameterized with our transaction item (the `item_c` you created in the last lesson), and a `result_t` type that represents the result. Since we ran `utg` above with `-q`, `utg` did not prompt us for the `reqType` and `rspType` interactively, so we'll have to fill those in ourselves.

Problem 5-1

Modify the driver and sequencer declarations to have the correct request and response pair. After doing so, you should be able to compile successfully.

Solution

The request and response types are the same thing for both the driver and sequencer: `item_c`. The base class `uvm_driver` sets the default response type to be the same as the request type (like a default argument):

```
class uvm_driver #(
    type REQ = uvm_sequence_item,
    type RSP = REQ
) extends uvm_component;
```

So, you only need to provide the request type and the response type is automatically filled in:

```
27.    verif/vcomponents/alu/alu_drv.sv
class drv_c extends uvm_driver#(item_c);
```

And:

```
28.    verif/vcomponents/alu/alu_sqr.sv
class sqr_c extends uvm_sequencer#(item_c);
```

You also may have run into the fact that the `alu_item.sv` needs to be ``included` in both the sequencer and driver files. This is because they now reference `item_c`, which is now a dependency.

Problem 5-2

Instantiate the ALU agent in the `alutb_pkg::env_c`. Compile and simulate to ensure you did it correctly.

Solution

Instantiating the ALU agent in the ALUTB environment first requires that you added the `alu.flist` file before the `alutb.flist` file in your Makefile, because of the ALUTB's dependency on the ALU package.

Then, instantiate an ALU agent:

```
29.      verif/vcomponents/alutb/alutb_env.sv
      // var: alu_agent
      // The ALU agent
      alu_pkg::agent_c alu_agent;
```

Then, add the following to its `build_phase`:

```
30.      verif/vcomponents/alutb/alutb_env.sv
      // create the ALU agent
      set_config_int("alu_agent", "is_active", is_active);
      alu_agent = alu_pkg::agent_c::type id::create("alu_agent", this);
```

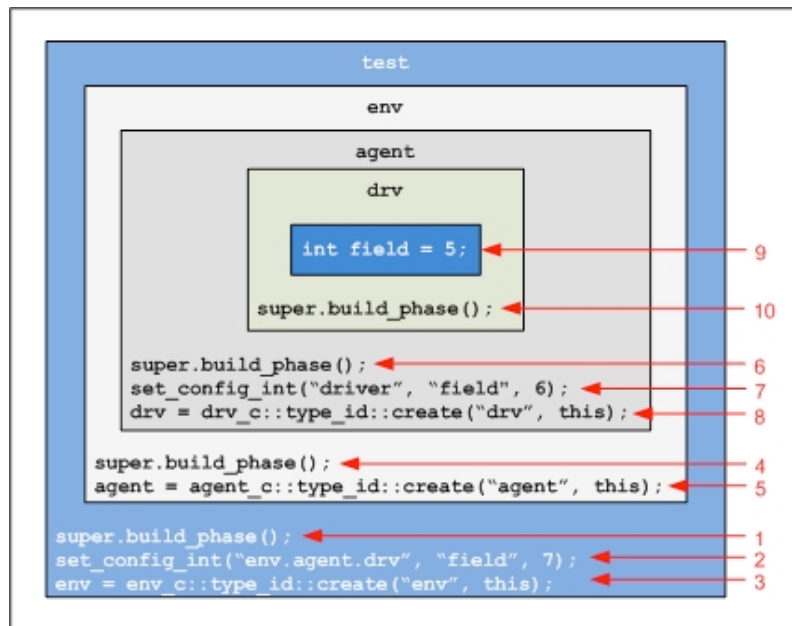
We have configured the ALU's `is_active` to match the one in the ALUTB environment. It's now time to learn about how configurations from the resource database trickle down to children components.

All About Configurations

In the following example, it's important to distinguish between the parent-child relationship of objects within the component hierarchy (i.e., an agent is the parent of a driver) versus the parent-child relationship of class inheritance (i.e., `basic_test_c` derives from `base_test_c`). To help make this clear, we'll refer to *descendants* and *ancestors* to discuss class inheritance, versus *parents* and *children* when discussing component hierarchy.

When a UVM component calls `super.build_phase()`, it calls its ancestor's `build_phase()` function. And that ancestor calls its ancestor's `build_phase()` function, and so on, until the ancestor is the base class `uvm_component`. There, the `build_phase` function looks in the resource database for all of the fields defined by the ``uvm_field` macros. It uses the component's full-name in the hierarchy and the field's name and type to match against items in the database. The field will be eventually set to the matching entry that was placed by the *highest* in the component hierarchy.

The following example hierarchy shows the order in which the functions are executed. At Step 1, only the test instance has been created and newed, and all other components do not exist yet.



1. When UVM enters the build phase, the test's `build_phase` function runs and it calls `super.build_phase()`.
2. The test then puts a value in a resource table that lives in the test that will later match against the `field` in the driver, and gives it a value of 7.
3. The test then creates the `env` instance, which gets newed. All of its fields are set to default values and its `new()` function runs. Its parent is a reference to the test. By creating the `env` component, UVM adds this instance to a list of components that still need to run their `build_phase`.
4. Later, UVM runs the `env`'s `build_phase` function and it also calls `super.build_phase()`.
5. The `env` creates the `agent` instance.
6. Later still, the `agent`'s `build_phase()` function runs and it, too, calls `super.build_phase()`.
7. The `agent` puts a value of 6 in its resource table that will match the driver's field.
8. The driver is created.
9. Upon creation, the field is given the default value of 5, and the driver's `new()` function is then run.
10. Finally, the driver's build phase runs and it, too, calls `super.build_phase()`. Again, this is the `uvm_component::build_phase()` function, and this is where all the work

happens. It calls the `apply_config_settings()` function, which is instructed by the ``uvm_field_int(field, UVM_ALL_ON)` macro to traverse upwards to find matching settings. The first one it finds will be the value of 6 supplied by its own parent, the agent. It then checks the agent's parent, `env`, to see if it has a matching entry. It doesn't, so it keeps going up to the test component, where it finds yet another matching entry which changes the value of the field to 7. Because the test has no parents, that will be the field's final value.

In addition to what was shown in this example is the power of using glob-style **or regular expressions (sure about this?)** for either the hierarchy or the field name. Instead of specifying that the driver's field gets a value of 7, the test could have specified that *any* components within the agent (including the monitor and sequencer) which also have a field named 'field' will also get the value of 7.

This is accomplished instead with this call:

```
set_config_int("env.agent.*", "field", 7);
```

Or, it could say that any integers in the env with a field name starting with "fi" will get a value of 7:

```
set_config_int("env.*", "fi*", 7);
```

As you'll see in Lesson 10, this is how the base test distributes the `cfg` and `reg_block` instances to *all* components within the hierarchy that asks for them:

```
// push the register block and the configurations to all blocks that ask for it
set_config_object("", "reg_block", reg_block, .clone());
set_config_object("", "cfg", cfg, .clone());
```

Care must be exercised, though, if you start instantiating other people's components and they also happen to use field names that have names like `cfg` or `reg_block`, but do not correspond to your testbench's notion of what these objects represent. UVM offers some protection against this by also matching on a field's type, but accidents are still possible.

All About The Class Factory

In Lesson 2, you used the class factory to replace the standard clock driver instance with one of your own. The process is similar to the configuration example above. To replace components, you call one of the four available functions (shown below) to submit a class override to the factory, and later when any sub-component tries to `create()` a class of the base type, the factory searches its database and replaces it with the descendant class you specified.

Component Factory Functions	Purpose
<code>set_type_override_by_type</code>	Replace all classes of the base type with the descendant type.
<code>set_inst_override_by_type</code>	Replace only classes whose name matches a wildcard pattern.
<code>set_type_override</code>	The same function as <code>set_type_override_by_type</code> , but use strings for the class names.
<code>set_inst_override</code>	The same function as <code>set_inst_override_by_type</code> , but use strings for the class names

To replace components, you would call these functions during the build phase of a parent component requesting the override. One important difference between the configuration example, however, is that you may need to call the override function *before* calling `super.build_phase()`.

For example, if you wanted to override the default clock driver from your base test, it stands to reason that you must request the override before creating the clock driver:

```
set_type_override_by_type(cn_pkg::clk_drv_c::get_type(), clk_duty_cycle_c::get_type(), .replace(1));
tb_clk_drv = cn_pkg::clk_drv_c::type_id::create("tb_clk_drv", this);
```

However, if you want to cause the override in a sub-class of `base-test`, such as `duty_cycle_test_c`, you must perform the override before `super.build_phase()`, as we did in Lesson 2. This is because calling `super.build_phase()` calls `base_test_c::build_phase()`, which will create the clock driver.

```
set_type_override_by_type(cn_pkg::clk_drv_c::get_type(), clk_duty_cycle_c::get_type(), .replace(1));
super.build_phase(phase); // <-- calls base_test_c::build_phase and creates the clock driver.
```

If you called it after `super.build_phase()`, the clock driver would have already been created and your factory override request would be too late.

Factory Overrides on Non-Components

Class factory overrides are not limited to just components, though. They do not rely on phases to work properly. You can create your class overrides on transactions, sequences, or any other UVM object. You can even do them after time zero. For example, you could enable sending corrupted CSR transactions after the configuration phase.

So long as the factory override request is made before the create call, the override will hold.

How To See Everything

At some point in your UVM career, you might want to do something with every single component. Or, you might want to print out all of the drivers named “Fred”. Whatever your strange inclinations, UVM offers a way to see everything in the component hierarchy.

`uvm_top` is the globally visible instance of `uvm_root`. No, UVM does not adhere to our clever `_c` naming conventions (let them wallow in their ignorance!). But, because `uvm_top` is globally visible, you can refer to it from anywhere in the codebase. This is awfully nice, because it offers some nifty functions for perusing the hierarchy:

The `find` method will return a handle to the `uvm_component` with the matching hierarchical name:

```
function uvm_component find(string comp_match);
```

The `find_all` method populates a list of components that match your query, and allows for `.` and `?` wildcards.

```
function void find_all(string comp_match,  
    ref uvm_component comps[$],  
    input uvm_component comp = null);
```

Finally, the function `print_topology` can be used to print out the entire hierarchy to the logfile, starting from `uvm_top`. This is a handy thing to do at the start of simulation, after everything has been created. Because it's such a good idea, the `global_env` will do it for you, so long as your debug level is non-zero. You can use the `cnmake` command-line option `TOPO` to specify the depth of the tree to print, with the default being 4.

```
% cnmake sim SIMU=VCS DBG=10 TOPO=8
```

This is essentially the same as printing out all of the configuration classes in `TestBuilder`, except that you don't have to write any `cnPrint` functions.

Conclusion

In this lesson, we learned all about components, their quasi-static nature, the component hierarchy, and how configurations trickle down from the top of the hierarchy. We also learned what goes into an agent and the purpose of all the different types of components. Finally, we discussed the class factory in more detail and a means to access or print all components.

Next, it's time to learn how components *talk* to one another, using Transaction Level Modeling.

Lesson 6: TLM Ports, Imps, and Exports

TLM 1.0 is the UVM equivalent of registered interfaces that were used in TestBuilder. Compared to registered interfaces, it's a more formalized method of passing pointers to classes to call their functions. TLM ports simulate having hierarchical connections between components.

Some of TLM's important features are:

- A TLM port is always the *initiator* of the function call.
- A TLM imp (implementation) is always the *target* of the function call.
- A TLM export provides pass-through functionality, to traverse different levels of hierarchy.
- TLM provides one-to-one, many-to-one, or one-to-many calls.
- TLM initiators can perform a get or a put.
- Calls can be blocking (time-consuming tasks) or non-blocking (zero-time functions).
- Analysis ports provide a one-to-many functionality.

TLM Ports

A TLM port is instantiated as a parameterized class. Like all classes, they must be new'ed. It's unlikely that we'll ever override these classes with the factory, but if you prefer the safety of calling `::create`, that's ok, too. Then, to put a transaction into the port, you call the port's `put()` function.

```
class producer_c extends uvm_component;
  `uvm_component_utils(producer_c)

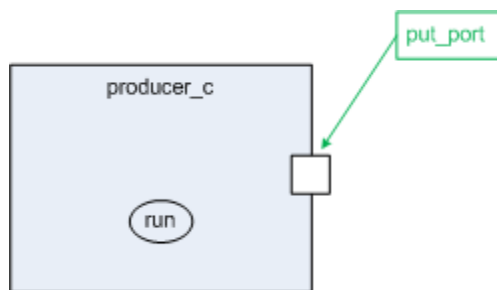
  uvm_blocking_put_port #(item_c) put_port;

  function new(string name, uvm_component parent);
    super.new(name, this);
    put_port = new("put_port", this);
  endfunction : new

  task run();
    item_c my_item;
    forever begin
      (#50ns) my_item = new("my_item");
      put_port.put(my_item);
    end
  endtask : run

endclass : producer_c
```

UVM uses squares in component diagrams to represent ports:



TLM Implementations (Imps)

The `imp` is the target of the `put`. You must instantiate the `imp` just as you did the port, *and* you must write the implementation function:

```
class consumer_c extends ovm_component;
  `uvm_component_utils(consumer_c)


  uvm_blocking_put_imp #(item_c, consumer_c) put_imp;

  function new(string name, uvm_component parent);
    super.new(name, this);
    put_imp = new("put_imp", this);
  endfunction : new

  task put(item_c _item);
    // do something with _item
  endtask : put
endclass : consumer_c
```

The `imp` class is parameterized not only with the item type, but also with the class's type, so that the `imp` knows which class contains the target implementation function.

Imps are represented with circles:



Connecting

Ports are connected to `imps` at a higher level of the hierarchy during the `connect_phase`. The reason this happens during the connect phase instead of during the build phase is because the two components you're trying to connect may not have yet run their build phases. And if that's the case, then they may not have created their ports or `imps` yet.

This is the whole reason for having these separate phases at time zero: all components in the system are *guaranteed* to have run all prior phases.


```
class env_c extends uvm_env;
  `uvm_component_utils(env_c)
  producer_c producer;
  consumer_c consumer;

  virtual function void build();
    super.build();
    producer = producer_c::type_id::create("producer", this);
    consumer = producer_c::type_id::create("consumer", this);
  endfunction : build

  virtual function void connect();
    super.connect();
    producer.put_port.connect(consumer.put_imp);
  endfunction : connect

endclass : env_c
```

Connect calls are always called on the initiator's port, with the receiving imp, port, or export as the argument. This completes the diagram and allows the producer to call another class's put function with a single argument, `item_c`.



TLM Exports

TLM Exports promote an implementation to a higher level in the hierarchy. From another component's point of view, they look exactly like an `imp`, but the real `imp` is buried someplace within the hierarchy. With exports, the external component need not know anything about the lower-level hierarchy.

In the example diagram below, `comp2` and `subcomp2` both have exports that promote the `imp` which is buried within `leaf2`. The `env` which makes the connection between `comp1` and `comp2` need not know the inner workings of either component.

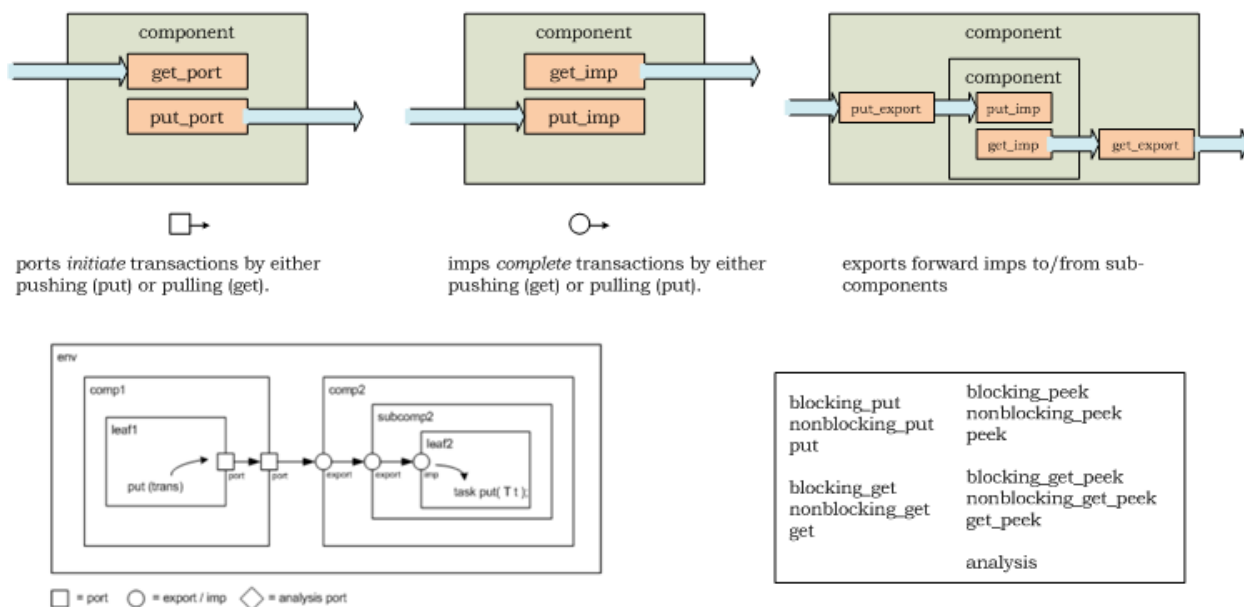
TLM Summary

To summarize:

- TLM has unidirectional interfaces of many forms.

- Blocking ports allow time-consuming tasks to be called.
- Non-blocking ports make zero-time function calls.
- The flow of information may be *from* the port (`put`), or *to* the port (`get`).
- The `port` is *always* the initiator, and the `imp` is *always* the target.

Unidirectional Interfaces



Imp Declarations

You may have noticed something flawed about the consumer class example above. The presence of the `uvm_blocking_put_imp` requires that a task called `put`, accepting an argument of type `item_c`, be present within the class. Astute observers would note that this precludes the class from having any other `put` imps. Surely UVM doesn't think we could live on one `imp` alone!

Fortunately, they don't. And there is a way around this. As usual, UVM chooses macros to solve the problem:

```
`uvm_put_imp_decl(monitored)
`uvm_put_imp_decl(driven)

class consumer_c extends uvm_component;
  uvm_put_imp_monitored#(item_c) monitored_imp;
  uvm_put_imp_driven#(item_c) driven_imp;

  function void put_monitored(item_c t);
    // receives puts coming into monitored_imp
  endfunction

  function void put_driven(item_c t);
```

```
// receives puts coming into driven_imp
endfunction
endclass : consumer_c
```

The ``uvm_put_imp_decl(_something)` macro creates a new class called `uvm_put_imp_something` that allows you to create a function called `put_something()`. This gives you the freedom to create all the imps you need. Preceding the argument with an underscore (`_`) allows you to avoid having to create the poorly-named function `putsomething()`.

The ``uvm_*_imp_decl` macros actually do create a new class for you. You're probably tempted to use the macro just above your class definition, as in this example. However, there may be more than one class in your package that needs to implement a `put`, `get`, or `write` for your given transaction. Using the macro twice, though, just won't do, because the compiler gets understandably cranky when a class is defined more than once. Coming up with a new name each time may not be suitable, either.

Rather than use the macros in only one class file and depend on the build order working out for you, it is advisable to instead place these declarations in the package file.

Compared to TestBuilder


Compared to registered interfaces in C++/TestBuilder, TLM seems to fall short somewhat. Recall that a registered interface permitted a pointer to a pure virtual interface class to be placed in another class. This permitted the other class to call *any* functions that were part of the pure virtual interface class. These functions could have as many parameters as you wished.

On the other hand, multiple inheritance made it fairly straightforward to abuse the registered interface system and put too much functionality into the interface, such that the lines between the two classes could blur and create spaghetti code. TLM, while less convenient, may result in a cleaner separation of architectures.

Many-to-One and One-to-Many

Ports and imps implicitly support many-to-one. That's because multiple ports can be connected to a single imp.

When you hear `analysis_port`, think *broadcast*. Analysis ports permit one-to-many communication and in diagrams are represented with a diamond. These are often the ports and imps of choice for broadcasting information to an indeterminate number of listeners.



We haven't even set up our monitor yet, but we will want for it to broadcast all monitored transactions and result values that it sees to the rest of the world. So, the `mon_c` class will need an analysis port, and the agent should also get one, too. This allows listeners to connect directly to the agent's analysis port without having to dig into the implementation to find out what component broadcasts stuff and what they're called. Providing a clean and consistent API from the agent helps to promote reuse.

Problem 6-1

Create separate `analysis_ports` in the monitor and the agent to broadcast transactions and results to the rest of the world. Connect them together. You don't have any code to test them yet, but make sure you can compile it.

Solution

In the monitor, first declare the ports:

```
31.      verif/vcomponents/alu/alu_mon.sv
      //-----
      // Group: TLM Ports

      // var: monitored_item_port
      // All monitored transactions go out here
      uvm_analysis_port #(item_c) monitored_item_port;

      // var: monitored_result_port
      // All monitored results go out here
      uvm_analysis_port #(result_t) monitored_result_port;
```

Bonus points if you placed these in the TLM Ports group and used a name that ended in `_port`, to comply with the coding standards.

Then, in the `build_phase`:

```
32.      verif/vcomponents/alu/alu_mon.sv
      monitored_item_port = new("monitored_item_port", this);
      monitored_result_port = new("monitored_result_port", this);
```

You can put the exact same lines in the agent, if you desire. You also would need to remember to ``include` the `alu_item.sv` file as it is now a dependency of these two files (just like with the driver and sequencer earlier).

Finally, in the `connect_phase` of the agent, you make the connection unconditionally (not dependent on `is_active`), because the monitor will be present in all types of testbenches:

```
33.      verif/vcomponents/alu/alu_agent.sv
      // connect to the monitor's analysis ports
      mon.monitored_item_port.connect(monitored_item_port);
      mon.monitored_result_port.connect(monitored_result_port);
```

Broadcasting to a Listener

Now let's write something out of the analysis port and create something to listen for it. Take your test code from `verif/alu/tests/basic.sv` and move it to the `main_phase` of the ALU monitor. After creating a transaction, call the port's `write()` function with the transaction as the argument. This will send it out the analysis port.

```
34.      verif/vcomponents/alu/alu_mon.sv
      virtual task main_phase(uvm_phase phase);
      byte unsigned stream[];
      alu_pkg::item_c item;
      alu_pkg::item_c unp_item;

      phase.raise_objection(this);

      repeat(50) begin
          item = alu_pkg::item_c::type_id::create("item");
```

```

        item.randomize();
        monitored_item_port.write(item);
        `cn_info($sprintf("Created ALU transaction: %s", item.convert2string()))
        item.pack_bytes(stream);
        `cn_info($sprintf("Bytes: %s", cn_pkg::print_ubyte_array(stream)))
        unp_item = alu_pkg::item_c::type_id::create("unp_item");
        unp_item.unpack_bytes(stream);
        `cn_info($sprintf("Unpacked: %s", unp_item.convert2string()))
        if(item.compare(unp_item) == 0)
            `cn_err($sprintf("Miscompare!"))
    end

    phase.drop_objection(this);
endtask : main_phase

```

In the ALUTB environment, where the ALU agent is instantiated, create a `uvm_subscriber` using `utg`, naming it `alu_item_subscriber`. A subscriber is a specialized UVM parameterizable component that comes pre-packaged with an `analysis_imp` that can accept whatever object it is parameterized for. All you need to do is fill in the `write()` function.

Here is how you create the subscriber component with `utg`:

35.

```

verif/vcomponents/alutb> utg subscriber -n alu_item_subscriber -f
++ Creating alutb_alu_item_subscriber.sv
alu_item_subscriber subscriber: Enter substitution for <description>: Listens for all monitored ALU
transactions.
alu_item_subscriber subscriber: Enter substitution for <subscription_type>: alu_pkg::item_c

```

Take special note that we used the scope operator to specify that the transaction is coming from the `alu_pkg`, because this component will be a member of the `alutb_pkg`.

Problem 6-2

Modify the `alu_item_subscriber` to print out the transactions that are written to it. Then, instantiate the component in the ALUTB environment and connect it to the ALU agent's transaction analysis port. Simulate and see that your subscriber is seeing all of the transactions. The subscriber uses an export called `analysis_export` that you will connect to.

Solution

Monitors don't often generate transactions, but this is just a temporary self-test. After removing some boilerplate code, the whole `alu_item_subscriber_c` class would look like this:

```
36.      verif/vcomponents/alu/alu_item_subscriber.sv
// class: alu_item_subscriber_c
// Print out all ALU transactions.
class alu_item_subscriber_c extends uvm_subscriber#(alu_pkg::item_c);
    `uvm_component_utils(alutb_pkg::alu_item_subscriber_c)

    //-----
    // Group: Methods
    function new(string name="alu_item_subscriber",
                  uvm_component parent=null);
        super.new(name, parent);
    endfunction : new

    ///////////////////////////////////////////////////////////////////
    // func: write
    // Receives the alu_pkg::item_c
    virtual function void write(alu_pkg::item_c t);
        `cn_info($sprintf("Received this ALU transaction: %s", t.convert2string()))
    endfunction : write
endclass : alu_item_subscriber_c
```

And the ALUTB environment would need to ``include` the `alutb_alu_item_subscriber.sv` file, instantiate and create the class as before, and then connect it during the `connect_phase` like this:

```
37.      verif/vcomponents/alu/alu_env.sv
alu_agent.monitored_item_port.connect(alu_item_subscriber.analysis_export);
```

If everything was done correctly, you should see your subscriber happily printing all the transactions it sees.

Leave it in there for now, but **remove the generated transactions from the monitor**, since up next we're going to write the real stimulus sequences.

Conclusion

In this lesson we learned all about how different UVM components *talk* to one another with ports, imps, and exports. We learned how to instantiate these TLM objects and how to connect them up. We also created a subscriber that can act as a benign listener.

There are two other interesting aspects of TLM that were not covered here:

- TLM can be used to bridge across different languages. Thus, an analysis port can send out all monitored transactions to a UVM subscriber, to a TestBuilder TVM, or to a SystemC implementation...all without having to change the UVM source code.

- TLM2 is new to UVM, but not new to the rest of the world. TLM2 offers ports that are bidirectional--called sockets. It also provides a way to generically model memory transactions using the `uvm_tlm_generic_payload` class. While the rest of UVM does not use TLM2 (yet!), later revisions of UVM will likely be driving towards it.

Lesson 7: Beginning Sequences

Sequences are intended to drive the stimulus items to the driver. They are dynamic objects that are created and destroyed (as opposed to components, which are quasi-static). They can act as generators to create new stimulus items. They can create other sequences and form a protocol stack. They can last only a short amount of time to perform a few operations. Or, they can be procedural blocks like an interrupt handler that last the whole life of the simulation.

The number and variety of use cases for sequences is fairly large, and their complexity may at first seem daunting to new users. This lesson is only the first lesson on sequences, to keep the topic more digestible. In Lessons 12 and 14, you will learn more details about sequences and sequencers.

A Simple Generator

Let's create a simple sequence to generate some random transactions. Be sure to first remove the generator code we placed in the monitor's `run_phase`.

Most sequences are short pieces of code. And as we'll see later you can create a special exerciser sequence that mixes and matches all (or some) of the sequences you write. For this purpose, it is often handy to put all of these sequences in one file called a sequence library file. Other sequences that you write may be larger and more specialized and may live in their own file.

Create a sequence library with `utg`. Sequences expect to eventually propagate sequence items (like our `item_c` class) as requests to a driver, and to optionally receive sequence items as responses. Since we unified our result data into the transaction class, it's ok if they're one and the same:

```
verif/vcomponents/alu> utg seq_lib -n alu -f
++ Creating alu_seq_lib.sv
alu seq_lib: Enter substitution for <description>:
alu seq_lib: Enter substitution for <reqType>: item_c
alu seq_lib: Enter substitution for <rspType>: item_c
```

Your sequence library now contains 2 pre-made classes for you. The first is a generic sequence called `alu_seq_c`, and the second is a sequence "library" sequence called `lib_seq_c`. This special library sequence is an exerciser that randomly picks a sequence that is registered to it and drives it. We won't need that for now, so comment it out.

A sequence is a class with a special time-consuming task called a `body()`. Technically, it has a `pre_body()` task as well as a `post_body()` task, but right now we'll just use the `body()` task to drive in a transaction:


```

38.      verif/vcomponents/alu/alu_seq_lib.sv
      //////////////////////////////////////
      // func: body
      virtual task body();
      item_c item = item_c::type_id::create("item");
      start_item(item);
      item.randomize();
      finish_item(item);
      endtask : body

```

Step-by-step, this is what this task does:

- Create a new transaction.
- Call `start_item`. This does a few things, but primarily waits until the sequencer is ready to send it.
- Now, randomize it. By randomizing after we've waited for the sequencer to accept it, we've achieved *late binding*. This allows us to add additional constraints, for example, based on the current conditions.
- Call `finish_item`. This pushes it to the sequencer that this sequence is operating on.

To see that the sequence is being driven to the driver, change the run phase of your ALU driver to get the next item, print it out, and then tell the sequencer that it's done:

```

39.      verif/vcomponents/alu/alu_drv.sv
      //////////////////////////////////////
      // func: run_phase
      virtual task run_phase(uvm_phase phase);
      // constantly poll for new transactions, printing them out
      forever begin
      seq_item_port.get_next_item(req);
      `cn_info($psprintf("Driving: %s", req.convert2string()))
      seq_item_port.item_done(req);
      end
      endtask : run_phase

```

Now, only one last thing remains. You have to create the sequence and *start* it on the sequencer. There are multiple ways of doing this, but putting it in the basic test's `main_phase` is the most straightforward for now:

```

40.      verif/alutb/tests/basic.sv
      //////////////////////////////////////
      // func: main_phase
      virtual task main_phase(uvm_phase phase);
      alu_pkg::alu_seq_c alu_seq = new("basic_seq");
      alu_seq.randomize();

      phase.raise_objection(this);
      `cn_info("Starting alu_seq.")
      alu_seq.start(alutb_env.alu_agent.sqr);
      phase.drop_objection(this);
      endtask : main_phase

```

Running the basic test should produce the desired results. When a sequence's `start` task is called, it will bind itself to the sequencer specified in its argument, and then call its `pre_start`, `pre_body`, `body`, `post_body`, and finally its `post_start` tasks.

At this point, you're probably a bit underwhelmed. After all, that's an awful lot of code to write just to send one measly transaction. But things will get more interesting soon.

`uvm_do Macros

For one thing, we can clean up our sequence a little bit. UVM provides a set of macros to create, randomize, and perform a sequence or sequence item from another sequence. These are the ``uvm_do` macros, and it's *very important* to know what each of them does before using them. But first, let's shorten our `alu_seq_c` body a bit:

```
41.      verif/vcomponents/alu/alu_seq_lib.sv
      virtual task body();
          item_c item;
          `uvm_do(item)
      endtask : body
```

Much shorter. ``uvm_do` does pretty much everything we had before. Sometimes, you want to randomize your sequence item with constraints. In this case:

```
42.      verif/vcomponents/alu/alu_seq_lib.sv
      `uvm_do_with(item, { operation == MUL_A_B; alpha == 7; beta == 5; })
```

``uvm_do_with` takes a constraint block and applies it during the randomization call. Try running this code to ensure you get the correct results.

Other times, you want to simply create the transaction, set it precisely, and then send it:

```
43.      verif/vcomponents/alu/alu_seq_lib.sv
      `uvm_create(item)
      item.operation = item_c::MUL_A_B;
      item.alpha = 7;
      item.beta = 5;
      `uvm_send(item)
```

This is less concise than the ``uvm_do_with` call, but it's more efficient because the transaction doesn't go through the randomization process.

UVM has other macro facilities to send things with a higher priority, to send them on a different sequencer (for sequences attached to virtual sequencers, to be discussed later), for randomizing and sending, and for combinations of all the above. **These macros can only be called from a sequence, not from within the component hierarchy.**

Learn what macros are available from the table below and use the right one.

Macro	Purpose
<code>`uvm_create</code>	Calls <code>::create</code> on the given sequence, assigning it to this sequence's sequencer.
<code>`uvm_do</code>	Creates, starts, randomizes, and finishes the given sequence, on this sequence's sequencer.
<code>`uvm_do_pri</code>	The same as <code>`uvm_do</code> , but assigns a priority.
<code>`uvm_do_with</code>	The same as <code>`uvm_do</code> , but randomizes with the given constraint block.
<code>`uvm_do_pri_with</code>	The same as <code>`uvm_do</code> , but assigns a priority and randomizes with the given constraint block.
<code>`uvm_create_on</code>	Calls <code>::create</code> on the given sequence, assigning it to this given sequencer.
<code>`uvm_do_on</code>	Creates, starts, randomizes, and finishes the given sequence, on the given sequencer.
<code>`uvm_do_on_pri</code>	The same as <code>`uvm_do_on</code> , but assigns a priority.
<code>`uvm_do_on_with</code>	The same as <code>`uvm_do_on</code> , but randomizes with the given constraint block.
<code>`uvm_do_on_pri_with</code>	The same as <code>`uvm_do_on</code> , but assigns a priority and randomizes with the given constraint block.
<code>`uvm_send</code>	Calls <code>start</code> and <code>finish</code> on a previously created sequence, without randomization.
<code>`uvm_send_pri</code>	Calls <code>start</code> and <code>finish</code> on a previously created sequence, without randomization, but with the specified priority.
<code>`uvm_rand_send</code>	Calls <code>start</code> , <code>randomize</code> , and <code>finish</code> on a previously created sequence.
<code>`uvm_rand_send_pri</code>	The same as <code>`uvm_rand_send</code> , but with the specified priority.
<code>`uvm_rand_send_with</code>	The same as <code>`uvm_rand_send</code> , but randomizes with the given constraint block.
<code>`uvm_rand_send_pri_with</code>	The same as <code>`uvm_rand_send_with</code> , but with the specified priority.

Getting a Response

Our driver doesn't do much with it yet, but it will soon enough. For the sake of our future examples, though, let's talk about how the driver would send back a response.

The response that the driver sends back to the sequencer will end up in the sequence that called it when the sequence calls its `get_response` function:

```
44.      verif/vcomponents/alu/alu_seq_lib.sv
      virtual task body();
      item_c item;
      `uvm_create(item)
      item.operation = item_c::MUL_A_B;
      item.alpha = 7;
      item.beta = 5;
      `uvm_send(item);
      get_response(rsp);
      `cn_info($psprintf("Got response result: %08X", rsp.result))
      endtask : body
```

What is `rsp`? Well, a sequence has two pre-defined variables in its base class--`req` and `rsp`--which are of the request and response types, respectively. In our case, they're both transactions of type `item_c`. Let's modify the ALU driver's `run_phase` to send back a real result whenever it sees a multiply operation:

```
45.      verif/vcomponents/alu/alu_drv.sv
      forever begin
      seq_item_port.get_next_item(req);
      `cn_info($psprintf("Driving: %s", req.convert2string()))
      if(req.operation == item_c::MUL_A_B)
      req.result = result_t'(req.alpha * req.beta);
      seq_item_port.item_done(req);
      end
```

Again, in our ALU example, our response is just the request with the result field filled in. This is a common application, but other cases may be different. A response might be a different packet or transaction. But it must be of type `uvm_sequence_item`, and there's no sense in creating a new one just to hold a 32-bit unsigned integer.

With any luck, you should see your sequence returns the number 0x23:

```
%I-(alu_seq_lib.sv: 51) [alutb_env.alu_agent.sqr.basic_seq] { 78ns} Got response result: 00000023
```

Problem 7-1

Spice things up a bit by adding a random count variable to your sequence. Constrain it to between 1 and 100, and perform that many random transactions, getting each one's response. Also, rename your sequence to `exer_seq_c`, since `alu_seq_c` isn't very descriptive. In addition to renaming it in `basic_test_c`, you'll also need to randomize it before calling its `start` function.

What happens if you don't call `get_response()`?

Solution

```

46.      verif/vcomponents/alu/alu_seq_lib.sv
// class: exer_seq_c
// Runs <count> transactions
class exer_seq_c extends uvm_sequence #(item_c, item_c);
    `uvm_object_utils_begin(alu_pkg::exer_seq_c)
        `uvm_field_int(count, UVM_ALL_ON | UVM_DEC)
    `uvm_object_utils_end

    //-----
    // Group: Fields

    // var: count
    // The number of random transactions to perform
    rand int count;
    constraint reasonable_cnstr { count inside {[1:100]}; }

    //-----
    // Group: Methods

    function new(string name="alu_seq");
        super.new(name);
    endfunction : new

    //////////////////////////////////////
    // func: body
    virtual task body();
        item_c item;

        repeat(count) begin
            `uvm_do(item)
            get_response(rsp);
            `cn_info($psprintf("Got response result: %08X", rsp.result))
        end
    endtask
endclass : exer_seq_c

```

If you do not fetch the responses, the sequence fills up with responses and starts complaining about a response queue overflow. You can manage the depth of your sequence's response queue by calling `set_response_queue_depth` and `get_response_queue_depth`, or you can turn off the error reporting with `set_response_queue_error_report_disabled`. Most often, though, it's prudent to just get the responses.

Response Queue Handler

Sequences can also launch multiple requests at a time:

```

fork
    `uvm_do(item1)
    `uvm_do(item2)
join

```

In this case, the response for `item2` could come before the response for `item1`. Or, depending your architecture, responses can come out-of-order by their very nature. As an alternative to calling `get_response(rsp)` for each request, you can set up a response handler.

To do so, you create the function `response_handler` that will be called whenever a response comes in. Because this is a virtual function from the base sequence class, which doesn't know what type of response to expect, the function will receive a generic `uvm_sequence_item`, which you can then `$cast` to your own response type. To enable this function to be called, you call `use_response_handler` with a value of 1. You can enable or disable this setting on-the-fly, although this may not be a good idea.

For example:

```
class my_seq_c extends uvm_sequence#(item_c);
  `uvm_object_utils(my_seq_c)

  //-----
  // Group: Fields

  // var: outstanding_requests
  // An assoc. array of all pending requests
  item_c outstanding_requests[int];

  //-----
  // Group: Methods
  function new(string name="my_seq");
    super.new(name);
    use_response_handler(1);
  endfunction : new

  //////////////////////////////////////
  // func: response_handler
  // Matches ID of each outstanding request
  virtual function void response_handler(uvm_sequence_item response);
    item_c resp, request;

    // cast the response to our transaction type
    $cast(resp, response);

    // look it up and compare
    request = outstanding_requests[resp.id];
    if(!request || !request.compare(resp))
      `cn_err("Response miscompare!")
    else
      outstanding_requests.delete(resp.id);
  endfunction : response_handler

  //////////////////////////////////////
  // func: body
  // Launch 5 transactions
  task body();
    item_c item[5] = new[5];
    fork
      foreach(item[x])
        fork
          automatic item_c this_item = item[x];
          `uvm_do(this_item)
          outstanding_requests[this_item.id] = this_item;
        join_none
      join
    endtask : body
  endclass : my_seq_c
```

There are some interesting SystemVerilog constructs above that you may not have seen before:

```
item_c outstanding_requests[int];
```

This declares an associative array (or hash, map, or dictionary, if you prefer) of transactions, keyed by an integer.

```
use_response_handler(1);
```

This tells the sequence that its `response_handler` function should be called for each response that is received.

```
// look it up and compare
request = outstanding_requests[resp.id];
if(!request || !request.compare(resp))
    `cn_err("Response miscompare!")
else
    outstanding_requests.delete(resp.id);
```

Sequence items come built-in with a unique `id` field. The first line above gets the request from the associative array, using the `id`. The second line uses the short-circuit condition of the OR operator. It first checks to see if it was found and if not it prints an error. Otherwise, it checks to see if it mis-compares. The last line deletes the outstanding request from the associative array.

```
fork
    foreach(item[x])
        fork
            automatic item_c this_item = item[x];
            `uvm_do(this_item)
            outstanding_requests[this_item.id] = this_item;
        join_none
    join
```

This seemingly complicated bit of code sends 5 transactions in parallel, pushing each one to the `outstanding_requests` hash by its `id`. Presumably, these IDs are unique. The nested fork-join construct is explained in more detail in Appendix B.

Complex Routines

At this point, sequences must seem pretty boring. But sequences get interesting when they do more than just one transaction. Our little ALU isn't capable of too many exotic functions, but there's enough there that we can do something more complicated.

Factorials

Let's create a sequence that performs the operation:

$$f(x) = x! = 1 * 2 * 3 * \dots * x$$

The first thing to do is create a new sequence. With `-c`, `utg` will just give you the sequence without the file header and you can cut-and-paste it into your `alu_seq_lib.sv` file:

```

verif/vcomponents/alu> utg seq -n factorial -c
factorial seq: Enter substitution for <reqType>: item_c
factorial seq: Enter substitution for <rspType>: item_c
class factorial_seq_c extends uvm_sequence #(item_c, item_c);
    `uvm_object_utils_begin(alu_pkg::factorial_seq_c)
    `uvm_object_utils_end

    //-----
    // Group: Fields

    //-----
    // Group: Methods

    function new(string name="factorial_seq");
        super.new(name);
    endfunction : new

    //////////////////////////////////////
    // func: body
    virtual task body();
    endtask : body
endclass : factorial_seq_c

```

Our factorial sequence should have one random field, operand, which is the number that we want to compute the factorial of. Then, we'll loop from 1 to operand, collecting the result response each time, and present the final answer. We should also squirrel away the answer in the class as a member field, so that the basic test can then print the final result.

Here is the final sequence:

```

47.    verif/vcomponents/alu/alu_seq_lib.sv
class factorial_seq_c extends uvm_sequence #(item_c, item_c);
    `uvm_object_utils_begin(alu_pkg::factorial_seq_c)
    `uvm_field_int(operand, UVM_ALL_ON)
    `uvm_field_int(result, UVM_ALL_ON)
    `uvm_object_utils_end

    //-----
    // Group: Fields

    // var: operand
    // The value to perform the factorial on
    rand bit [15:0] operand;
    constraint operand_cnstr { operand <= 9; }

    // var: result
    // The final result
    bit [15:0] result = 1;

    //-----
    // Group: Methods
    function new(string name="factorial_seq");
        super.new(name);
    endfunction : new

    //////////////////////////////////////
    // func: body
    // Loop from 1..operand and multiply all the numbers together
    virtual task body();
        item_c item;
        byte num;

        for(num = 1; num <= operand; num++) begin
            `uvm_do_with(item, { operation == MUL_A_B; alpha == num; beta == local::result; })
            get_response(rsp);
            result = rsp.result;
            `cn_info($sprintf("num=%0d, result=%0d", num, result))
        end
    endtask
endclass

```



```
end
`cn_info($psprintf("%0d! = %0d", operand, result))
endtask : body
endclass : factorial_seq_c
```

If you're a math whiz, you've already figured out that our ALU has a bit of a limitation. Since it can only take 16-bit operands, our largest factorial operand is a meager 9. And the result we keep must be constrained to those 16 bits. Nonetheless, it's a good example of how sequences can be used to convert multiple transactions into one larger transaction.

Problem 7-2

Modify your `basic_test_c` to create and send in a factorial sequence with an operand of 9, and print the result. What would happen if instead of the variable name `answer`, we chose the name `result` instead? How would you solve this problem?

Solution

The first part of this should have been straightforward:

```
48.      verif/alutb/tests/basic.sv
////////////////////////////////////
// func: main_phase
virtual task main_phase(uvm_phase phase);
    alu_pkg::factorial_seq_c factorial_seq = new("factorial_seq");
    factorial_seq.operand = 9;
    phase.raise_objection(this);
    factorial_seq.start(alutb_env.alu_agent.sqr);
    `cn_info($psprintf("The factorial of 9 is %d", factorial_seq.result))
    phase.drop_objection(this);
endtask : main_phase
```

The problem with naming the answer `result` instead is that it will conflict with the local name `result` in the transaction. When the constraint is applied, it is local to the `item_c` class, which already has a `result` field. Thus, the `beta == result;` constraint refers to the transaction's version of `result`. Worse, because we left `result` uninitialized, VCS will complain that there are X's or Z's in a constraint value.

You might be tempted to use `beta == this.result;` to distinguish the two, but since the constraint block is in the `item_c` scope, `this` also refers to the transaction and not the sequence.

Fortunately, the architects of SystemVerilog anticipated this problem and created the `local::` scope operator. This does exactly what you want:

```
49.      verif/vcomponents/alu/alu_seq_lib.sv
`uvm do with(item, { operation == MUL A B; alpha == num; beta == local::result; })
```

Problem 7-3

Write a sequence, `summation_seq_c`, that performs this operation, and prints the results:

$$\sum_{k=x}^y k = x + (x+1) + (x+2) + \dots (y-1) + y$$

So things don't get too crazy, constrain its `x` and `y` variables to be between 1 and 20.

Solution

You could use the ADD_A_B operation, but this solution uses the ACCUM operation.

```
50.      verif/vcomponents/alu/alu_seq_lib.sv
// class: summation_seq_c
// Sends in ALU transactions that add up the values from op_x to op_y, storing the final answer in
// result
class summation_seq_c extends uvm_sequence #(item_c, item_c);
    `uvm_object_utils_begin(alu_pkg::summation_seq_c)
        `uvm_field_int(op_x, UVM_ALL_ON | UVM_DEC)
        `uvm_field_int(op_y, UVM_ALL_ON | UVM_DEC)
        `uvm_field_int(result, UVM_ALL_ON | UVM_DEC)
    `uvm_object_utils_end

    //-----
    // Group: Fields

    // vars: op_x, op_y
    // Operands for this summation function
    rand bit [15:0] op_x;
    rand bit [15:0] op_y;

    constraint operands_cnstr { op_x < op_y; }
    constraint reasonable_cnstr {
        op_x inside {[1:20]};
        op_y inside {[1:20]};
    }

    // var: result
    // The final result
    result_t result = 0;

    //-----
    // Group: Methods
    function new(string name="summation_seq");
        super.new(name);
    endfunction : new

    //////////////////////////////////////
    // func: body
    // Loop from x..y, adding them up
    virtual task body();
        item_c item;
        byte num;
        `uvm_do_with(item, { operation == CLR_RES; })
        get_response(rsp);
        for(num = op_x; num <= op_y; num++) begin
            `uvm_do_with(item, {
                operation == ACCUM;
                alpha == num;
            })
            get_response(rsp);
        end
        result = rsp.result;
    endtask : body
endclass : summation_seq_c
```

You also would need to change your driver to provide a result when the operations are either ACCUM or CLR_RES:

```
51.      verif/vcomponents/alu/alu_drv.sv
virtual task run_phase(uvm_phase phase);
    result_t prev_result;

    // constantly poll for new transactions, printing them out
```

```

    forever begin
        seq_item_port.get_next_item(req);
        `cn_info($psprintf("Driving: %s", req.convert2string()))
        case(req.operation)
            item_c::MUL_A_B: req.result = result_t'(req.alpha * req.beta);
            item_c::ACCUM : req.result = result_t'(prev_result + req.alpha);
            item_c::CLR_RES: req.result = 0;
        endcase
        prev_result = req.result;
        seq_item_port.item_done(req);
    end
endtask : run_phase

```

You may have missed that you'll want to send in a CLR_RES transaction first, to make sure the ALU does not already contain a result. However, as you'll see soon, when other sequences start interfering this will become a small problem.

Problem 7-4

Let's write one more, just to make sure you've got the hang of it. This sequence will create a random array of words and add them up. We'll call it `sum_array_seq_c`. To create a randomly sized array, it's always important to constrain the array's size to a reasonable value, otherwise it will probably randomize itself to an extremely large number and crash the simulator.

```

52.    verif/vcomponents/alu/alu_seq_lib.sv
    // var: data
    // An array of words to be summed
    rand bit [15:0] data[];

    // keep it to a reasonable size
    constraint data_cnstr { data.size() inside {[1:50]}; }

```

Now, when we randomize our sequence, it will create a randomly sized array of random 16-bit words. To declare our data array as a UVM configuration field, we use the following macro:

```

53.    verif/vcomponents/alu/alu_seq_lib.sv
    `uvm_field_array_int(data, UVM_ALL_ON | UVM_DEC)

```

This line declares an array of *numbers*. UVM's field macros just lump all numbers as ints, they do not distinguish between bytes, ints, or 42-bit values. Declaring this field allows you to pack, unpack, and print it correctly (if you wanted to).

Write the `body()` task of the `sum_array_seq_c` sequence to send transactions into the ALU that will sum the data array.

Solution

Just like in our last sequence, we'll need to first clear the result state, and then push in each piece of data in turn. Be sure to modify your `basic_test_c` class to randomize your new sequence and send it.

```
54.    verif/vcomponents/alu/alu_seq_lib.sv
    virtual task body();
        item_c item;
        result_t result;

        `uvm_do_with(item, { operation == CLR_RES; })
        get_response(rsp);
        foreach(data[x]) begin
            `uvm_do_with(item, { operation == ACCUM; alpha == data[x]; })
            get_response(rsp);
        end
        result = rsp.result;
        `cn_info($psprintf("Sum of this array = %0x", result))
    endtask : body
```

We've also used SystemVerilog's handy `foreach` construct. This neatly replaces iterators from C++ and creates two implicit variables: `x` and `data[x]`. These implicit variables are local in scope only to the `foreach` block.

Sequencing other Sequences

So far, our little sequences have just been producing sequence items of our transaction type, `item_c`. But, the `uvm_sequence` class is derived from `uvm_sequence_item`. Therefore, sequences can also *do* other sequences. This allows you to have layers of sequences.

When you call ``uvm_do` (or any of its variants) on a sequence, things are just a little bit different. First of all, you don't need to (in fact, you shouldn't) call `get_response`. That's because the ``uvm_do` macros recognize that you are using a sequence and they call their start tasks, instead of `start_item`. When ``uvm_do` completes, your sequence is already done, and whatever response data that's stored in the sequence is available for your inspection.

Problem 7-5

With everything you've learned so far, this next one should present some fresh challenges. Use our `factorial_seq_c` and our `sum_array_seq_c` to provide a `sum_of_factorials_seq_c` that solves this equation:

$$\sum_{k=x}^y k = x! + (x+1)! + (x+2)! + \dots (y-1)! + y!$$

Solution

Here is an implementation of the sum-of-factorials equation using the `sum_array_seq_c` sequence. Hopefully yours looks similar:

```
55.      verif/vcomponents/alu/alu_seq_lib.sv
class sum_of_factorials_seq_c extends uvm_sequence #(item_c);
    `uvm_object_utils_begin(alu_pkg::sum_of_factorials_seq_c)
        `uvm_field_int(op_x, UVM_ALL_ON | UVM_DEC)
        `uvm_field_int(op_y, UVM_ALL_ON | UVM_DEC)
        `uvm_field_int(result, UVM_ALL_ON | UVM_DEC)
    `uvm_object_utils_end

    //-----
    // Group: Fields

    // vars: op_x, op_y
    // Operands for this summation function
    rand bit [15:0] op_x;
    rand bit [15:0] op_y;

    constraint operands_cnstr {
        op_x < op_y;
        op_x inside {[1:9]};
        op_y inside {[1:9]};
    }

    // var: result
    // The final result
    result_t result;

    //-----
    // Group: Methods
    function new(string name="sum_of_factorials_array_seq");
        super.new(name);
    endfunction : new

    //////////////////////////////////////
    // func: body
    virtual task body();
        int num;
        bit [15:0] data[];
        factorial_seq_c fact_seq;
        sum_array_seq_c sum_seq;
        int idx;

        // fill the data array with all the factorials
        data = new[(op_y - op_x + 1)];
        idx = 0;
        for(num = op_x; num <= op_y; num++) begin
            `uvm_do_with(fact_seq, { operand == num; })
            data[idx] = fact_seq.result;
            idx++;
        end

        // now sum the array
        `uvm_create(sum_seq)
        sum_seq.data = data;
        `uvm_send(sum_seq)
        result = sum_seq.result;
        `cn_info($psprintf("The sum of factorials from %0d to %0d is %0d.",
                           op_x, op_y, result))

    endtask : body
endclass : sum_of_factorials_seq_c
```

A few important notes about this implementation:

```
data = new[(op_y - op_x + 1)];
```

This is how you size a dynamic array--by newing it with the number of entries. Using a queue would be a nice alternative, but the `sum_array_seq_c` expects a dynamic array of unsigned bytes, so that's what we'll give it.


```
`uvm_create(sum_seq)
sum_seq.data = data;
`uvm_send(sum_seq)
```

We do not want to create a random sequence with random data. We want to use a very specific array, therefore we cannot use ``uvm_do`. Instead, we call ``uvm_create` to perform both the `new` operation and to wait for the sequencer to accept it. Then, we assign the data array to our local array. Because it's an array, this is merely copying a reference and so doesn't have the overhead of copying each byte at a time. Finally, we call ``uvm_send` to send it to the sequencer. When this call returns, the sequence has completed and we can fetch the answer.

Sequence Hierarchy

Allowing sequences to 'do' other sequences is the basis of a sequence hierarchy. More complex examples might involve a complete protocol hierarchy. A hefty example is an independent TCP or UDP session, living inside IPv4 packets, driven into a device using SRIO messages, all for the sake of performing a compression operation.

What's the point of all this? The point is that the compression sequence does not need to know anything about the underlying sequences, which interface it went over, or how it was broken up into many packets over a long period of time. It just deals with sending in a compression request, and when it's done, it has the compression result.



The Sequencer

Take a look at your `alu_sqr.sv` sequencer file. Note that there's really nothing in it. It's actually doing quite a bit, though. Where sequencers really rock is when they're handling multiple streams of stimulus at once. So far, we've been sending it one sequence at a time. Pretty boring. Let's instead use `fork..join` to start multiple sequences at the same time. Change your basic test's `main_phase` to something more like this:

```
56.    verif/alutb/tests/basic.sv
virtual task main_phase(uvm_phase phase);
    alu_pkg::sum_of_factorials_seq_c sof_seq = new("sof");
    alu_pkg::sum_array_seq_c sum_array_seq = new("sum_array");
    alu_pkg::exer_seq_c exer_seq = new("exer_seq");

    phase.raise_objection(this);
    fork
        begin
            sof_seq.randomize();
            `cn_info($psprintf("Starting:\n%s", sof_seq.sprint()))
            sof_seq.start(alutb_env.alu_agent.sqr);
            `cn_info($psprintf("The sum-of-factorials from %0d to %0d is %d",
                                sof_seq.op_x, sof_seq.op_y, sof_seq.result))
        end

        begin
            sum_array_seq.randomize() with {data.size() > 8; };
            `cn_info($psprintf("Starting:\n%s", sum_array_seq.sprint()))
            sum_array_seq.start(alutb_env.alu_agent.sqr);
            `cn_info($psprintf("The sum is %0d", sum_array_seq.result))
        end

        begin
            exer_seq.randomize();
            exer_seq.start(alutb_env.alu_agent.sqr);
            `cn_info($psprintf("exer sequence completed after %0d transactions.", exer_seq.count))
        end

        begin
            alu_pkg::factorial_seq_c fact_seq = new("fact_seq");
            fact_seq.randomize();
            fact_seq.start(alutb_env.alu_agent.sqr);
        end
    join

    phase.drop_objection(this);
endtask : main_phase
```

The `fork..join` concept is one you'll soon use a lot of. Here, we've spawned four separate threads that will start four separate randomized sequences on the `alutb_env.alu_agent.sqr` at the same time. The sequencer now has four sequence threads arbitrating for its attention. It has a built-in arbitrator that has several different modes of operation:

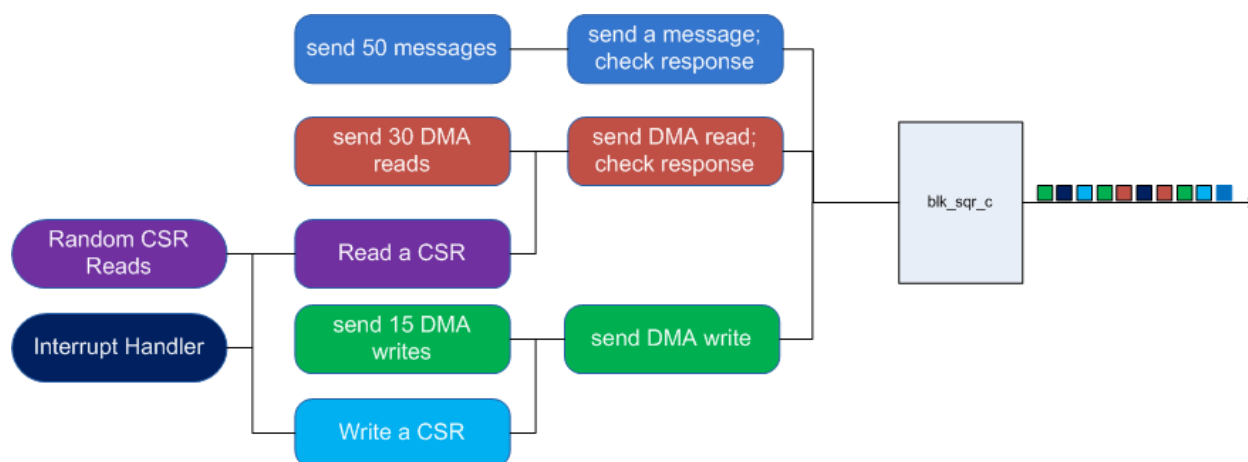
Arbitration Type	Purpose
SEQ_ARB_FIFO	Requests are granted in FIFO order (default)
SEQ_ARB_WEIGHTED	Requests are granted randomly by weight

Arbitration Type	Purpose
SEQ_ARB_RANDOM	Requests are granted randomly
SEQ_ARB_STRICT_FIFO	Requests at highest priority granted in fifo order
SEQ_ARB_STRICT_RANDOM	Requests at highest priority granted in randomly
SEQ_ARB_USER	Arbitration is delegated to the user-defined function, <code>user_priority_arbitration</code> . That function will specify the next sequence to grant.

The sequencer chooses to arbitrate among the sequences when the currently chosen sequence either waits for a delay or sends an item. When the sequence calls `start_item()`, it puts its hat in the ring to try to win the next arbitration, and the item's priority and other factors will determine the winner. The arbitration occurs when the driver calls `get_next_item()`. If the sequence is chosen, then the `start_item()` task finishes and the sequence's thread takes over, performing randomization on the item, or whatever else you wish, and eventually sending it along to the driver.

This interleaving of sequence items is the whole point of sequences and sequencers. While they have other benefits, this is what turns generic operations into a tangled web of stimulus that is better at finding bugs than directed test cases are.

Below is a more complex example of how sequencers and their sequences can be viewed. On the left is a hierarchy of independent sequences, each one minding its own business and doing its own thing. They can each call lower-level sequences. In the middle, the sequencer performs the arbitration and interleaving. On the rightmost edge is the random stream of transactions themselves.



And it does it all for free, which actually has two benefits:

- First, you don't have to write this complex arbitrator yourself, for every testbench, and for every transaction type.
- Second, these sequences essentially operate in a vacuum. They effectively separate the generation of stimulus from the underlying implementation, allowing you to focus on one thing at a time. Creating a sum-of-factorials sequence does not need to take into account how the factorial sequence was implemented. As we'll see later, this can lead down the path of *prediction encapsulation*.

Locking and Grabbing

If you now run your test, you'll find that all your sequences are quite jumbled together. However, if you used your calculator to verify the output in the logfile, you'll see this leads to a problem. In order to sum an array of values using the clear and accumulate operations, these sequences need to be atomic. Otherwise, the stored result value in the ALU will become mixed up with all the other calculations.

However, a sequence can get exclusive access to its sequencer via `lock` or `grab`. When a sequence locks its sequencer, it first waits for it to win arbitration, and then it blocks other sequences until it calls `unlock`. Calling `grab` does the same thing, but it prioritizes the lock so that it will win arbitration next so long as other sequences do not already have exclusivity on the sequencer. You would use `grab` instead of `lock` for higher priority sequences--such as an interrupt handler.

Problem 7-6

Change the two sequence bodies of `summation_seq_c` and `sum_array_seq_c` to lock the sequencer before doing the `CLR_RES` operation and to unlock it after the last `ACCUM` operation. Verify that your results are now correct.

Solution

The body of your `summation_seq_c` should now look something like this:

```
57.      verif/vcomponents/alu/alu_seq_lib.sv
      virtual task body();
      item_c item;
      byte num;
      lock();
      `uvm_do_with(item, { operation == CLR_RES; })
      get_response(rsp);
      for(num = op_x; num <= op_y; num++) begin
          `uvm_do_with(item, {
              operation == ACCUM;
              alpha == num;
          })
          get_response(rsp);
      end
      result = rsp.result;
      unlock();
      endtask : body
```

And your `sum_array_seq_c` will have this in its body task:

```
58.      verif/vcomponents/alu/alu_seq_lib.sv
      virtual task body();
      item_c item;

      lock();
      `uvm_do_with(item, { operation == CLR_RES; })
      get_response(rsp);
      foreach(data[x]) begin
          `uvm_do_with(item, { operation == ACCUM; alpha == data[x]; })
          get_response(rsp);
      end
      unlock();
      result = rsp.result;
      `cn_info($psprintf("Sum of this array = %0x:", result))
      endtask : body
```

Prediction Encapsulation

One of the most complex tasks in creating a testbench is prediction. The RTL does what it does with this jumble of stimulus we've created, and the predictor has the enormous task of determining if it's correct or not.

We will later write a predictor that sees each individual transaction and determines if the ALU generated the correct result. This is a fairly straightforward mechanism that should work, for such a simple design. But how do we know that larger operations are succeeding?

Take, for example, our sequence that calculates the sum of an array. Relying only on our basic predictor, the testbench would presume that since each of the accumulate operations gave the correct result that the entire operation is correct. But before locking the sequencer, both your summation sequences were getting the wrong answer. The testbench would not have found this.

The sequence itself, though, has all the facilities to reach the correct value and check the result, because it has all the data right there and it receives the final answer. By encapsulating a predictor right in the sequence, our sequence becomes self-checking. What's more, our sequence knows nothing of all the other traffic flying around. Imagine writing the scoreboard that has to identify which operations went with which. You might have to embed a tag or some other identifier within each transaction so that the predictor can determine where they all go.

Just as a demonstration of how prediction encapsulation can work, let's use the `sum_array_seq_c` as an example.

```

59.    verif/vcomponents/alu/alu_seq_lib.sv
virtual task body();
    item_c item;
    result_t exp_result = data.sum();

    lock();
    `uvm_do_with(item, { operation == CLR_RES; })
    get_response(rsp);
    foreach(data[x]) begin
        `uvm_do_with(item, { operation == ACCUM; alpha == data[x]; })
        get_response(rsp);
    end
    unlock();
    result = rsp.result;
    `cn_info($psprintf("Sum of this array = %0x:", result))
    if(result != exp_result)
        `cn_err($psprintf("The result %0x did not match the expected value %0x.", result,
                           exp_result))
endtask : body

```

We've embedded a checker to ensure that the entire operation completes successfully, and we've used SystemVerilog's handy `sum()` function on a dynamic array to quickly calculate the answer. This will ensure that the DUT can correctly calculate the sums of an array.

Have we achieved *prediction nirvana*!?!?

Well, maybe yes and maybe no. For one thing, not all sequences will be re-usable at a higher level of abstraction, like in the full-chip testbench. If you are not able to carry a complex prediction algorithm forward to other testbenches, then that's not a good place to put it.

For another, a great many things can happen to your transactions along their journey that this sequence might not be privy to. These circumstances may or may not be manageable. For example, what would happen if a reset occurred in the middle of this sequence? While this situation may be manageable, others may not be so simple.

How would you check this answer *without* embedding the prediction into the sequence? The long and involved answer may lead you to believe that could be the better approach.

Conclusion

Sequences and sequencers offer dramatically more features than presented here. You'll see more of them later. What we know now should suffice to provide enough interesting stimulus for the driver we're about to write.

Lesson 8: Drivers and Monitors

We already created the empty driver class way back in Lesson 5. In this lesson, we're going to hook it up to the interface we created in Lesson 4, receive requests from the sequencer, drive the transactions, collect the results, and send them back to the sequencer. Then, we'll create a monitor in the same way, and use analysis ports to announce what happens on the interface.

Fetching the Interface

You'll recall that during Lesson 4 we instantiated the `alu_intf` in the top-level testbench and pushed virtual references to its modports into the resource configuration database with the names `alu::drv_vi` and `alu::mon_vi`. The suffix `_vi` indicates that it's a virtual interface.

Well, we're trying to make a reusable agent, and the interfaces won't always have those names, so let's create a configuration string in the driver called `intf_name`:

```
60.      verif/vcomponents/alu/alu_drv.sv
class drv_c extends uvm_driver#(item_c);
    `uvm_component_utils_begin(alu_pkg::drv_c)
        `uvm_field_string(intf_name, UVM_ALL_ON)
    `uvm_component_utils_end

    //-----
    // Group: Configuration Fields

    // var: intf_name
    // The name of the virtual interface that we'll hook up to
    string intf_name = "drv_vi";
```

As you can see, we have used the macro ``uvm_field_string` to declare this configuration field, so that it will auto-populate itself when its `super.build()` is called. If no other classes have added this configuration to the database for this driver, then it will take the default name of `"drv_vi"`, which is what we want in the first place. This is how you give a configuration its default value.

Next, we need to declare the virtual interface as a field in the driver, so that it can access the signals:

```
61.      verif/vcomponents/alu/alu_drv.sv
//-----
// Group: Fields

// var: drv_vi
// Virtual interface to drive on
virtual alu_intf.drv mp drv_vi;
```

We've placed this definition in the fields group, to let readers know that this is not something they'll be touching. All they need to do is set the `intf_name` string correctly.

Also, the type of `drv_vi` must match **exactly** what was used to store this interface in the database in `alutb_tb_top`:

```
`cn_set_intf(virtual alu_intf.drv_mp, "alu_pkg::alu_intf", "drv_vi" , alu_i.drv_mp);
```

This will give us access to the driver modport of this interface. Now, it needs to be fetched from the interface.

Problem 8-1

Fetch the interface from the resource database and assign it to `drv_vi`. In the driver's `run_phase`, call the interface's `reset()` function.

You can also **remove** the call to `reset()` that we left in the testbench earlier:

```
initial
    alu_i.reset();
```

Afterwards, you should be able to simulate successfully.

Solution

Like the ``cn_set_intf` macro that was used earlier, there is a corresponding ``cn_get_intf`:

```
62.      verif/vcomponents/alu/alu_drv.sv
      virtual function void build_phase(uvm_phase phase);
      super.build_phase(phase);
      // get the interface
      `cn_get_intf(virtual alu_intf.drv_mp, "alu_pkg::alu_intf", intf_name, drv_vi)
      endfunction : build_phase
```

The macro unrolls to this:

```
if(!uvm_resource_db#(virtual alu_intf.drv_mp)::get("alu_pkg::alu_intf", intf_name, drv_vi))
  `cn_fatal($psprintf("%s virtual interface not present.", intf_name))
```

Note that it's important to flag a fatal error if the interface is not found, because that would indicate a configuration error, and there is no purpose in continuing. The call to get the interface must take place *after* the call to `super.build_phase()` because if it happened before then the `intf_name` variable would not have been updated by the database.

Your `run_phase` should now look like this to ensure that the assertions in the interface do not fire.

```
63.      verif/vcomponents/alu/alu_drv.sv
      virtual task run_phase(uvm_phase phase);
      drv_vi.reset();
      endtask : run_phase
```

Getting the Next Item

As you saw in Lesson 7, drivers have a built-in port called `seq_item_port` through which the sequencer pushes `uvm_sequence_items`. It's a special port of type `uvm_seq_item_pull_port` that is parameterized to your request and response types. Like the sequences, the driver also has built-in fields `req` and `rsp` corresponding to its current request and response.

To get the next request, you call a blocking task :

```
seq_item_port.get_next_item(req);
```

It will return when the next request is available. Or, you can call the function `try_next_item(req)` to see if one is available in zero-time.

Our ALU design operates in a pull-mode. The driver attempts to fetch the next item from the sequencer, and the sequencer's arbitration scheme selects a sequence, and so on. Both driver and sequencer are also offered in push-mode, where the sequencer pushes sequence items to the driver.

Driving a Transaction

There are myriad approaches to doing this, but we went through the exercise of learning how to pack and unpack transactions, so we might as well use that method.

UVM's current packing algorithms are offered in three flavors: give me a list of bits, give me a list of bytes, and give me a list of ints (32-bit numbers). Since this is the fun part, you can do it.

Problem 8-2

Write a task in the driver called to fetch and drive transactions. Name it `driver`, and change the `run_phase` task to simply call this task.

Hint:

You can use the following algorithm:

- Call the interface's reset function, to ensure that X's do not get into the design.
- Run forever.
 - Get the next transaction.
 - Pack the transaction into an array of bytes.
 - For each byte:
 - Wait 1 clock cycle on the interface.
 - Assert the `ctl` line only if this is the very first cycle.
 - Drive the byte of data.
 - Wait 1 clock, then clear the bus.
 - Wait for the `ready` line to go high, then wait 1 clock.
 - Fetch the result from the `result` bus.
 - Assign the result to the original request's result field.
 - Call `item_done` and send back the original request.

Solution

The challenge here was probably dealing with the clocking block. While it may make for extra typing here, clocking blocks are worth the effort if the clock edges may change or if you later want to add sampling delays. Clocking blocks in an interface define what the sampling clock edge is, so the notion of `@(posedge clk)` has been abstracted away. If an interface suddenly became a dual-data rate interface, the only place you would need to change any code would be in the clocking block.

```
64.    verif/vcomponents/alu/alu_drv.sv
// func: driver
// Drive transactions by packing into an array of bytes, then sending two 4-bit cycles
// for each byte. Then wait for the response and send it back.
task driver();
    byte unsigned stream[];

    @(posedge drv_vi.drv_cb.rst_n);

    forever begin
        seq_item_port.get_next_item(req);

        req.pack_bytes(stream);
        foreach(stream[x]) begin
            @(drv_vi.drv_cb);

            drv_vi.drv_cb.ctl <= (x == 0)? 1'b1 : 1'b0;
            drv_vi.drv_cb.dat <= stream[x];
        end

        // wait 1 clock, then clear the bus
        @(drv_vi.drv_cb);
        drv_vi.reset();

        // wait for result
        @(posedge drv_vi.drv_cb.ready);
        @(drv_vi.drv_cb);
        req.result = drv_vi.drv_cb.result;
        seq_item_port.item_done(req);
    end
endtask : driver
```

If you happened to leave the checker in the `sum_array_seq_c`, then hopefully the test still passes.

A few notes on the above design:

```
@(drv_vi.drv_cb);
```

Because of the clocking block, the method to wait 1 clock is to just wait on the clocking block.

```
drv_vi.drv_cb.ctl <= (x == 0)? 1'b1 : 1'b0;
drv_vi.drv_cb.dat <= stream[x];
```

The virtual interface that the driver has is a modport of an existing interface. Recall that this modport, `drv_mp`, has only two features: `drv_cb` and the `reset` task:

```
modport drv_mp(clocking drv_cb,  
              import reset);
```

Therefore, all signal accesses (both read and write) must take place through the clocking block, which also specifies the direction (input or output) that the driver is permitted to access. There is no such signal as `drv_vi.dat`. There is only `drv_vi.drv_cb.dat`. The same can be said for `rst_n`.

Monitoring Activity

As its name implies, the monitor is a component that watches the interface. It checks for any protocol violations and collects all activity it sees, broadcasting these as transactions to all listeners through an analysis port. Unlike the driver and the sequencer, a monitor is always present in an agent whether it is active or not.

The monitor in this agent will look very much like the driver, except the exact opposite. Instead of watching the `drv_cb` clocking block, it will watch the `mon_cb`. Instead of getting transactions from the sequencer, it will broadcast the ones that it sees. And instead of packing and driving transactions, it will collect data and unpack into a transaction.

Problem 8-3

Write the ALU monitor. Use the hints below if you need to.

Hint:

As with the driver, we must start by making sure we get the interface. Then, set up the `run_phase` to launch two tasks: `monitor_item` and `monitor_result`. Use the `fork...join_any` construct so that the tasks are disabled on reset.

The `monitor_result` task should be straightforward: constantly wait for the rising edge of the `ready` signal, and broadcast the result bus out the `monitored_result_port` by calling its `write()` function.

The `monitor_item` task is a little bit trickier. Because the protocol offers no end-of-transaction signal, it must be inferred from the operation. Wait for the rising edge of `ctl`, grab the data, and based on the operation that's being driven, collect the correct number of data cycles in a dynamic array of bytes. Report an error if the `ctl` signal goes high on any cycle other than the first one.

When that's done, you should be able to create a new transaction, unpack those bytes into it, and write it out of the `monitored_item_port`. You might also want to add debug messages to see what it's doing.

Solution

Here is one possible implementation of the ALU monitor, with some explanations following:

```

65.      verif/vcomponents/alu/alu_mon.sv
`include "alu_item.sv"

// class: mon_c
// Monitors an ALU bus and reports activity.
class mon_c extends uvm_monitor;
    `uvm_component_utils_begin(alu_pkg::mon_c)
        `uvm_field_string(intf_name, UVM_ALL_ON)
    `uvm_component_utils_end

    //-----
    // Group: Configuration Fields

    // var: intf_name
    // The name of the virtual interface that we'll hook up to
    string intf_name = "mon_vi";

    //-----
    // Group: TLM Ports

    // var: monitored_item_port
    // All monitored transactions go out here
    uvm_analysis_port #(item_c) monitored_item_port;

    // var: monitored_result_port
    // All monitored results go out here
    uvm_analysis_port #(result_t) monitored_result_port;

    //-----
    // Group: Fields

    // var: mon_vi
    // Virtual interface to monitor
    virtual alu_intf.mon_mp mon_vi;

    //-----
    // Group: Methods
    function new(string name="mon",
                  uvm_component parent=null);
        super.new(name, parent);
    endfunction : new

    //-----
    // func: build_phase
    virtual function void build_phase(uvm_phase phase);
        super.build_phase(phase);

        // get the interface
        `cn_get_intf(virtual alu_intf.mon_mp, "alu_pkg::alu_intf", intf_name, mon_vi)

        monitored_item_port = new("monitored_item_port", this);
        monitored_result_port = new("monitored_result_port", this);
    endfunction : build_phase

    //-----
    // func: run_phase
    virtual task run_phase(uvm_phase phase);
        forever begin
            fork
                monitor_item();
                monitor_result();
                @(negedge mon_vi.mon_cb.rst_n);
            join
        end
    endtask

```

```

        join_any

        `cn_info("Stopping monitor due to reset.")
        disable fork;

        @(posedge mon_vi.mon_cb.rst_n);
    end

    endtask : run_phase

    //////////////////////////////////////
    // func: monitor_item
    // Watch and broadcast the transactions on the bus
    virtual task monitor_item();
        int num_clocks;
        byte unsigned data[];
        item_c item;

        forever begin
            // wait for the rising edge of the control
            @(posedge mon_vi.mon_cb.ctl);

            // determine how many clocks are in this itemaction
            case(mon_vi.mon_cb.dat)
                // 5-cycle transactions
                item_c::ADD_A_B, item_c::SUB_A_B, item_c::SUB_B_A,
                item_c::MUL_A_B, item_c::DIV_A_B, item_c::DIV_B_A :
                    num_clocks = 5;

                // 3-cycle transactions
                item_c::INC_A, item_c::INC_B, item_c::ACCUM :
                    num_clocks = 3;

                // 1-cycle transactions
                item_c::CLR_RES :
                    num_clocks = 1;

                default:
                    `cn_err($sprintf("Unknown operation type: %02X", mon_vi.mon_cb.dat))
            endcase

            // collect the data for each cycle
            data = new[num_clocks];
            for(int clk=0; clk < num_clocks; clk++) begin
                data[clk] = mon_vi.mon_cb.dat;
                @(mon_vi.mon_cb);
                if(mon_vi.mon_cb.ctl == 1)
                    `cn_err($sprintf("The CTL signal is high during a transaction that should have been
%0d clocks.", num_clocks))
            end

            // create the transaction, unpack into it, and write it out the monitored_item_port
            item = item_c::type_id::create("mon_item");
            item.unpack_bytes(data);
            `cn_info($sprintf("Monitored: %s", item.convert2string()))
            monitored_item_port.write(item);
        end
    endtask : monitor_item

    //////////////////////////////////////
    // func: monitor_result
    // Monitor the ready and result signal and broadcast it out the monitored_result_port
    virtual task monitor_result();
        forever begin
            @(posedge mon_vi.mon_cb.ready);
            `cn_info($sprintf("Monitored Result: %08X", mon_vi.mon_cb.result))
            monitored_result_port.write(mon_vi.mon_cb.result);
        end
    endtask : monitor_result

endclass : mon_c

```

A few explanations may be in order:

```
@(posedge mon_vi.mon_cb.ct1);
```

Here, we are waiting for the positive edge of the `ct1` signal. Because of the clocking block, there is no concern that this event will be triggered on the rising edge of `ct1` but before the DUT drives the `dat` lines. Why is that? Because the event will sample on the clock specified by the clocking block--in this case, the positive edge of `clk`. If the clocking block had specified the negative edge of `clk`, you would find that the `posedge` event triggers on the negative edge of the `clk`, even though the `ct1` signal is driven on the rising edge.

```
`cn_err($sprintf("Unknown operation type: %02X", mon_vi.mon_cb.dat))
```

A key element to any monitor is protocol checking. The interface itself already checks for X's. Here, we ensure that only legal values are specified on the first cycle. Later, we check that the `ct1` signal does not go high any other time during the transaction. And, at the end of the transaction, we guarantee that the data collected has a legal format because it can successfully unpack. If the unpack were to fail, it would report an error for us.

```
data = new[num_clocks];
```

If you're coming from a C/C++ background, you might be wary by this line, because there are no free or delete operators. And this sits in a forever loop. This looks like a memory leak, but it's not. SytemVerilog has automatic garbage-collection, so when nobody else is referencing this data, it gets freed automatically. As it happens, this data is local in scope to this task, and it's copied byte-by-byte into the transaction during the unpack operation. Therefore, the bytes pointed to by the `data` variable will have zero references the next time the `data` variable is newed.

```
item = item_c::type_id::create("mon_item");
```

Like the `data` variable, the `item` variable is local in scope to the task, and we keep creating new ones. Are these also garbage-collected when the `item` variable creates a new one? Well, yes and no. If nobody were listening on the analysis port, then they would be garbage collected. In this testbench, though, we have the monitor broadcasting to the agent, which broadcasts to the `alu_item_subscriber_c` class we implemented in Lesson 6. This subscriber doesn't hold onto the transaction, though. It merely prints it out, so these transactions will be garbage collected. Other subscribers may store them in a scoreboard, or some other temporary location, and their reference count will not go to zero, so they will not be freed.

```
item.unpack_bytes(data);
```

The `unpack_bytes` function is used in the same manner as the `pack_bytes` function we used in the driver. The `unpack_bytes` function will report an error on its own if it is

unsuccessful, and it returns an `int` which is the number of bits that were unpacked. We don't need this information, though, so we don't look at it.

```
monitored_item_port.write(item);
```

Here we are writing that transaction out of the analysis port. Elsewhere, we write the result value out of the other analysis port. Since these are connected to the agent's analysis ports, they will likewise be broadcast out of the agent.

Conclusion

Drivers and monitors are a snap in SystemVerilog because of all the built-in constructs like events, queues, dynamic and associative arrays, and easy signal access. Hooking them up the UVM way wasn't too difficult, either.

It's probably time we got around to checking that the ALU is actually doing some math, though.

Lesson 9: Writing a Predictor

We've been so focused on creating, driving, and monitoring the stimulus that we haven't bothered to check if it's correct or not. We'll be getting to reading and writing CSRs in Lesson 10, so for now consider the K and C values to be 1 and 0 respectively.

UVM offers a variety of options when constructing your predictor. Some are more appropriate than others, depending on the nature of the prediction being made.

But first, some definitions:

Intra-Agent Prediction

Refers to the prediction within a single agent, requiring only one monitor to perceive all necessary information.

Inter-Agent Prediction

Predictions of DUT activity between different agents. Multiple agents are required to stimulate and monitor the traffic conditions.

A block-level testbench may have only one or both of these. The ALU agent is an example of a testbench that requires only one intra-agent predictor, because one monitor is sufficient to monitor both the transactions and the results.

A PCIE agent might also have an intra-agent predictor, tracking requests and responses within the PCIE framework. But that agent will likely communicate to an outside predictor that coordinates with another agent to check traffic at a higher level of abstraction.

Monitor Prediction

For such a small testbench, you're probably tempted to put all of the prediction in the monitor. And for something this simple, that decision is probably justified. Typically, though, this is not such a great idea. Separating the low-level pin details and interface protocol checking from the higher-level prediction algorithms is often considered the best practice.

Scoreboards

UVM provides a `uvm_scoreboard` class, but it might interest you to know that in UVM 1.1 that class is just an empty component and serves no useful purpose. (As a matter of fact, so are the agent, environment, test, and a variety of other UVM classes). The reasons these "empty" classes exist are that they help identify the purpose of the derived class

you're creating and they serve as placeholders for future functionality that may come along someday.

Meantime, if the nature of your prediction algorithm requires a scoreboard, this is the component you should create. Pipe monitored requests and responses into the scoreboard component via TLM interfaces and use some of SystemVerilog's handy data types such as queues or associative arrays and away you go.

PW Scoreboard


The folks at Paradigm Works, Inc., created and released a more fully-functional open source scoreboard that appears in the vcomponents directory as `pw_scoreboard`. They are lobbying for UVM to adopt it as the standard `uvm_scoreboard` class, and it does have some compelling features.

It's a parameterized component that you can drop in with the transaction types you hope to check, and it automatically instantiates the analysis exports for you. The API is described in more detail in its [User's Guide](#).

Comparators

UVM offers several built-in comparator components that operate as simplistic versions of the `PW scoreboard`. The `uvm_in_order_comparator`, `uvm_in_order_built_in_comparator`, and `uvm_in_order_class_comparator` all compare two streams of data objects, both of the same type. This may be useful in situations where what goes in must come out, but for most applications probably will not be sufficient.

One place where these comparators can be useful is as an agent self-check. Placing a comparator in the agent between the driver and the monitor ensures that both components are in sync with one another.



The `uvm_algorithmic_comparator` offers similar functionality, but is parameterized to work with different transaction classes as the two streams and also takes a class that is used to transform the first class into the predicted second class.

This sounds perfect for our ALU, which would have transaction sequences going in one port, and results going in the other, with a simple math function in between to predict expected results. What could be easier?

Unfortunately, the class that UVM 1.1 offers has one small bug--it cannot take an integer type as one of the streams. We would have to enclose the result value into a separate results class containing just a 32-bit value. Not an overwhelming limitation, though.

Prediction Encapsulation

In Lesson 7, we saw how the sequences themselves could perform their own prediction, under the right circumstances. This was great for checking the result of a longer computation because it isolated the prediction from the interleaved traffic that is common on interfaces.

Roll-Your-Own

Creating your own predictor component will often be the avenue of choice, and it's the one we'll choose here, though not because it's the most prudent. Rather, we'll roll our own because it's the situation you'll most often find yourself in. Also, it will later provide us with the ability to demonstrate using your predictor as a reference model in place of the real RTL.

This one should be a piece of cake for you.

Problem 9-1

Instantiate an ALU predictor (`pred_c`) that derives from `uvm_component` in the ALU agent and provide it with two analysis imps, one taking an `item_c` and the other taking a `result_t` value. Use the ``uvm_analysis_imp_decl` macros to permit you to create both `write()` functions.

Hook up the predictor's imps to the monitors ports and fire away.

Solution

Hopefully this was not too challenging for you. As discussed in Lesson 6, these can be placed in the `alu_pkg.sv` file:

```
66.      verif/vcomponents/alu/alu_pkg.sv
      `uvm_analysis_imp_decl(_item)
      `uvm_analysis_imp_decl(_result)
```

It's not too distressing if you placed them in the same file as the predictor, but it's a good habit to always put them in the package so they are declared for all classes.

```
67.      verif/vcomponents/alu/alu_pred.sv
class pred_c extends uvm_component;
  `uvm_component_utils(pred_c)
  //-----
  // Group: TLM Ports

  // var: monitored_item_imp
  uvm_analysis_imp_item #(item_c, pred_c) monitored_item_imp;

  // var: monitored_result_imp
  uvm_analysis_imp_result #(result_t, pred_c) monitored_result_imp;

  //-----
  // Group: Fields

  // var: result
  // The result of the monitored transaction is stored here and checked with the received result
  result_t result = 0;

  //-----
  // Group: Methods
  function new(string name="pred",
               uvm_component parent=null);
    super.new(name, parent);
  endfunction : new

  //-----
  // func: build_phase
  function void build_phase(uvm_phase phase);
    super.build_phase(phase);

    monitored_item_imp = new("monitored_item_imp", this);
    monitored_result_imp = new("monitored_result_imp", this);
  endfunction : build_phase

  //-----
  // func: write_item
  // Accepts ALU transactions and sets the next expected result
  virtual function void write_item(item_c _item);
    case(_item.operation)
      item_c::ADD_A_B : result = _item.alpha + _item.beta;
      item_c::SUB_A_B : result = _item.alpha - _item.beta;
      item_c::SUB_B_A : result = _item.beta - _item.alpha;
      item_c::MUL_A_B : result = _item.alpha * _item.beta;
      item_c::DIV_A_B : result = _item.alpha / _item.beta;
      item_c::DIV_B_A : result = _item.beta / _item.alpha;
      item_c::INC_A   : result = _item.alpha + 1;
      item_c::INC_B   : result = _item.beta + 1;
      item_c::CLR_RES : result = 0;
      item_c::ACCUM   : result += _item.alpha;
    endcase

    `cn_dbg(30, $sprintf("Calculated result %08X on item: %s", result, _item.convert2string()))
```

```

endfunction : write_item

////////////////////////////////////
// func: write_result
// Called when a result is monitored
virtual function void write_result(result_t _result);
    if(_result != result)
        `cn_err($sprintf("Actual result: %08X != Expected result: %08x",
                           _result, result))
    endfunction : write_result
endclass : pred_c

```


Hopefully by now, your only real challenge was having multiple impls appear in a component. This required the declaration macros, the new `uvm_analysis_imp` types that these macros created, and the naming of your implementation functions `write_item` and `write_result`.

Reference Modeling


It won't always be the case, but sometimes your predictor behaves very much like the DUT itself. Sometimes it's a C or SystemC reference model; or the RTL hasn't actually been written yet and you'd like something to test against; or you want to try some early performance modeling to prove the architecture; or you want to run very fast sims without the RTL to test that your functional coverage is adequate.

Whatever your reasoning, sometimes you want a *reference model*. The predictor designed above is an excellent candidate for that.

Let's review what the architecture looks like right now. Our ALU agent--when it is active--looks like this:



In order to create a reference model, all we would really need to do is this:



We could modify our predictor to contain a `seq_item_port` just like the driver, so that it can accept transactions. Then, in the main or run phase, we could have it get the next item, calculate the result just as it was doing before, and send it back to the sequencer as a response. We would also want to add analysis ports that broadcast the transactions and results as if a monitor were present. And we could have it wait for a computational delay that models the real hardware.

The rest of the verification environment is completely isolated from this change, and for all intents and purposes the RTL might as well be present.

How much practical use this method would be is a matter of debate. But in this testbench, it is so easy that it begs to be written.

Problem 9-2

Modify the ALU predictor to behave like a reference model when a configuration field `ref_model` is set. Modify the ALU agent to connect itself like a reference model when its `ref_model` field is set. You may need to reset the ALU interface at time zero to avoid the x-checkers.

Create a new test, `ref_model_test_c`, that turns these bits on.

Solution

The `alu_agent_c` modifications are fairly straightforward:

```
68.      verif/vcomponents/alu/alu_agent.sv
      `uvm_field_int(ref_model, UVM_ALL_ON)
      ...
      // var: ref_model
      // When set to 1, the predictor operates in reference mode and the monitor/driver are not enabled
      bit ref_model = 0;
```

In the build phase, ensure that the predictor's value of `ref_model` matches the agent's by default. This way, the test(s) only have to modify the agent.

```
69.      verif/vcomponents/alu/alu_agent.sv
      //////////////////////////////////////
      // func: build_phase
      virtual function void build_phase(uvm_phase phase);
      super.build_phase(phase);
      set_config_int("pred", "ref_model", ref_model);
```

Also, don't create either the driver or the monitor when in reference model mode:

```
70.      verif/vcomponents/alu/alu_agent.sv
      if(!ref_model)
        mon = mon_c::type_id::create("mon", this);
      if(is_active) begin
        if(!ref_model)
          drv = drv_c::type_id::create("drv", this);
        sqr = sqr_c::type_id::create("sqr", this);
      end
```

The connect phase should match the TLM connections of the diagram.

```
71.      verif/vcomponents/alu/alu_agent.sv
      virtual function void connect_phase(uvm_phase phase);
      super.connect_phase(phase);

      if(!ref_model) begin
        ...
      end else begin
        // as a reference model
        pred.seq_item_port.connect(sqr.seq_item_export);
        pred.monitored_item_port.connect(monitored_item_port);
        pred.monitored_result_port.connect(monitored_result_port);
      end
      endfunction : connect_phase
```

The changes needed to the predictor are shown below in boldface:

```
72.      verif/vcomponents/alu/alu_pred.sv
      class pred_c extends uvm_component;
      `uvm_component_utils_begin(pred_c)
      `uvm_field_int(ref_model, UVM_ALL_ON)
      `uvm_component_utils_end

      //-----
      // Group: Configuration Fields

      // var: ref_model
      // When set, operates in reference model mode
      bit ref_model = 0;
```

```

//-----
// Group: TLM Ports

// var: monitored_item_imp
uvm_analysis_imp_item #(item_c, pred_c) monitored_item_imp;

// var: monitored_result_imp
uvm_analysis_imp_result #(result_t, pred_c) monitored_result_imp;

// var: seq_item_port
// As a reference model, pulls transactions from the sequencer
uvm_seq_item_pull_port #(item_c) seq_item_port;

// var: monitored_item_port
// As a reference model, drives out the transactions that were "driven"
uvm_analysis_port #(item_c) monitored_item_port;

// var: monitored_result_port
// As a reference model, drives out the results that were "seen"
uvm_analysis_port #(result_t) monitored_result_port;

//-----
// Group: Fields

// var: result
// The result of the monitored transaction is stored here and checked with the received result
result_t result = 0;

// var:
//-----
// Group: Methods
function new(string name="pred",
             uvm_component parent=null);
    super.new(name, parent);
endfunction : new

////////////////////////////////////
// func: build_phase
function void build_phase(uvm_phase phase);
    super.build_phase(phase);

    if(ref_model) begin
        seq_item_port = new("seq_item_port", this);
        monitored_item_port = new("monitored_item_port", this);
        monitored_result_port = new("monitored_result_port", this);
    end else begin
        monitored_item_imp = new("monitored_item_imp", this);
        monitored_result_imp = new("monitored_result_imp", this);
    end

end

endfunction : build_phase

////////////////////////////////////
// func: main_phase
task main_phase(uvm_phase phase);
    item_c item;

    if(ref_model) begin
        forever begin
            seq_item_port.get_next_item(item);

            `cn_dbg(30, $sprintf("REF_MODEL: %s", item.convert2string()))

            // create a delay that models the transmission of the transaction
            #5ns;
            monitored_item_port.write(item);

            // calculate result
            write_item(item);
        end
    end
endtask

```



```

        // create a delay that models the ALU calculation speed
        #15ns;

        // send back the result
        write_result(result);
        item.result = result;
        seq_item_port.item_done(item);
        monitored_result_port.write(result);
    end
end
endtask : main_phase

////////////////////////////////////
// func: write_item
// Accepts ALU transactions and sets the next expected result
virtual function void write_item(item_c _item);
    case(_item.operation)
        item_c::ADD_A_B : result = _item.alpha + _item.beta;
        item_c::SUB_A_B : result = _item.alpha - _item.beta;
        item_c::SUB_B_A : result = _item.beta - _item.alpha;
        item_c::MUL_A_B : result = _item.alpha * _item.beta;
        item_c::DIV_A_B : result = _item.alpha / _item.beta;
        item_c::DIV_B_A : result = _item.beta / _item.alpha;
        item_c::INC_A   : result = _item.alpha + 1;
        item_c::INC_B   : result = _item.beta + 1;
        item_c::CLR_RES : result = 0;
        item_c::ACCUM   : result += _item.alpha;
    endcase

    `cn_dbg(30, $psprintf("Calculated result %08X on item: %s", result, _item.convert2string()))
endfunction : write_item

////////////////////////////////////
// func: write_result
// Called when a result is monitored
virtual function void write_result(result_t _result);
    if(_result != result)
        `cn_err($psprintf("Actual result: %08X != Expected result: %08x",
            _result, result))
    endfunction : write_result

endclass : pred_c

```

The big change is the `main_phase`, which now looks similar to the one in the driver. We could modify it to choose different delays based on the transaction type. Or, we could pull in the virtual interface that the driver uses and wait on clock edges instead.

You already know how to create the test, but just in case you've forgotten, it's here:

```

73.    verif/alutb/tests/ref_model.sv
// class: ref_model_test_c
// Test the ALU using the predictor as a reference model
class ref_model_test_c extends basic_test_c;
    `uvm_component_utils(ref_model_test_c)

    //-----
    // Group: Methods
    function new(string name="test",
        uvm_component parent=null);
        super.new(name, parent);
    endfunction : new

    //////////////////////////////////////
    // func: build_phase
    virtual function void build_phase(uvm_phase phase);
        super.build_phase(phase);
        set_config_int("alutb_env.alu_agent", "ref_model", 1);
    endfunction : build_phase

```

```
endclass : ref_model_test_c
```

An alternative to this approach would be to leave the ALU agent largely the same, but disconnect the RTL from the interface. Allow the driver to drive the transaction onto the interface and the monitor to see it. Then, when the monitor tells the predictor that a transaction occurred, have the predictor drive the result directly onto the interface some time later. This approach leads to fewer changes to the code, but incurs the overhead of having a driver and monitor and dissociates itself from operating only at the transaction level.

Conclusion

UVM offers a variety of choices for simple prediction schemes. TLM interfaces provide a lot of flexibility and can allow you to run without RTL at all, provided you plan ahead. Ultimately, though, prediction algorithms will be simple because of the rich set of aggregate data types and other functions that are available in the SystemVerilog language.


Lesson 10: Configuration Registers

The CSRs that are specific to this testbench are described in the file `rtl/alu/alu.csr`. In there, you'll find a CSR titled `ALU_CONST`, and it contains a read/writable field called `K_VAL`. Let's write a test that specifically configures that CSR field to a value between 9 and 20.

Register Organization

Let's first understand where this register is in the hierarchy. Here is a picture of the testbench hierarchy that gets you to that field:

This picture shows that the test component instantiates two objects, `cfg` and `reg_block`. The test then has the `cfg`'s instance of `reg_block` reference the one in the test. Because it's a reference, there is only one instance of the `reg_block` in the simulation, and no extra memory is consumed. This is accomplished with the following code:



```
// create reg_block
if(reg_block == null) begin
    reg_block = alu_csr_pkg::alu_ncb_reg_block_c::type_id::create("reg_block", this);
    reg_block.configure(null, "aluth_tb_top.dut_wrapper.dut");
    reg_block.build();
    reg_block.lock_model();
end

// set configuration's reference to reg_block
cfg.reg_block = reg_block;
```

The `reg_block` is derived from `uvm_reg_block` and inside of this lives the `reg_file`, which is derived from `uvm_reg_file`. To understand this better, let's observe the following definitions from the UVM 1.1 Reference Manual:

Register Block

A register block represents a design hierarchy. It can contain registers, register files, memories and sub-blocks. A block has one or more address maps, each corresponding to a physical interface on the block.

Register Address Map

An address map is a collection of registers and memories accessible via a specific physical interface. Address maps can be composed into higher-level address maps.

Register File

A register file is a collection of register files and/or registers used to create regular repeated structures. Register files are usually instantiated as arrays.

Register

A register represents a set of fields that are accessible as a single entity. A register may be mapped to one or more address maps, each with different access rights and policy.

Register Adapter

This class defines an interface for converting between the generic register operation (`uvm_reg_bus_op`) and a specific bus transaction.

In short, when you hear “Register Block” think design hierarchy (such as MIO, IOB, or GMX). When you hear “Address Map” think interface (such as NCB, RSL, or PCIE). When you hear “Register File” think of repeatable structures such as per-port statistics registers, or physical and virtual functions.

In our simple testbench, the register block created by the CSR generation scripts contains the registers in the `alu` block. The register block is located in `vcomponents/reg/obj/unit/regs__alu.sv`. It contains the register classes and the register block associated with the NCB address map. In other blocks, it may be the case that multiple different interfaces can access registers, in which case they would have more than one map.

References for Everyone!

As is usually the case, many different components and objects will want to be able to access the CSRs and the testbench’s configuration class. As we saw in Lesson 5, UVM’s configuration and resource database gives us a way to push references with minimal code. We accomplish this in the following manner:

1. Push the register block into the resource database with the field name “`reg_block`.” Push it to all components by setting the `inst_name` argument to “*”. Since we want each component to have a reference to the original register block, set the `clone` argument to zero.

2. For any component that needs access to the register block, instantiate a field of that type with the name “`reg_block`.” Use the ``uvm_field` macros to declare this field as a member of this component.
3. Each component will get a reference to this one instance when it calls `super.build_phase()`.

The base test uses this method to push both its `cfg` field and the `reg_block` field to any component that wishes to see them with this bit of code:

```
// push the register block and the configurations to all blocks that ask for it
set_config_object("", "reg_block", reg_block, .clone());
set_config_object("", "cfg", cfg, .clone());
```

Configuring Your Block

You’d be excused if you missed it while flipping through the reference manual, but UVM’s register block class contains a nice little function called `update()`. What this function does is perform the *minimum* number of CSR writes to get the CSRs in the DUT to match the testbench’s register block. So, the process to randomize and configure a block’s CSRs is straightforward:

1. Randomize the register block. The CSR script has declared all registers as `rand` member fields, and all register fields as `rand` member fields, so randomizing the entire block will randomize all of the register fields.
2. Call `reg_block.update()`, preferably during the configure phase.

That was easy. Maybe too easy. It turns out that UVM will write the CSRs in address order. Maybe that’s what works for your block (it works great for the `alu`). But it might require a more complicated configuration routine than that.

In that case, you can add modify the `configure_phase` of a component, and you can organize the register writes in whatever manner you choose. The registers themselves also have a task called `update()`, and by calling this it will perform the CSR write (if and only if one is needed).

Register blocks and registers both also have a corresponding task called `mirror()`, which will perform CSR reads to get the environment’s version of the registers to match the one in the DUT.

The kval Test

Let’s put all of this together and try to accomplish the goals of this lesson. Then we’ll take a look at how all of it worked.

Create a test class called `kval_test_c` and extend it from the `basic_test_c` class. In it, add this constraint:

```
74.      verif/alutb/tests/kval.sv
constraint kval_cnstr {
    reg_block.ALU_CONST.K_VAL.value inside {[9:20]};
}
```

Up until now the CSRs have contained their innocuous values, due to this constraint in the `alutb_pkg::cfg_c` class:

```
// Constrain K_VAL and C_VAL to both be innocuous
constraint innocuous_cnstr {
    reg_block.ALU_CONST.K_VAL.value == 1;
    reg_block.ALU_CONST.C_VAL.value == 0;
}
```

If we were to run this test now, there would be a conflict in the constraints that the simulator would be unable to solve. Go ahead and try it.

Problem 10-1

You could eliminate this problem by removing this constraint altogether. How could you fix it *without* removing the constraint or altering the configuration class?

Solution

Your first instinct might have been to use the factory. In that case, you could derive a class from `alutb_pkg::cfg_c`, and in its `new` function set the `constraint_mode` of the `innocuous_cnstr` to zero.

However, the `cfg` instance is accessible from the test. So long as you turn off the constraint mode before the configuration class is randomized, then the constraint will not be applied. How you do that, though, is the hard part.

```
virtual function void build_phase(uvm_phase);
    cfg.innocuous_cnstr.constraint_mode(0);
    super.build();
endfunction : build_phase
```

If you do this, you'll find that the `cfg` instance hasn't been created yet, so you're trying to access a NULL object.

```
virtual function void build_phase(uvm_phase);
    super.build();
    cfg.innocuous_cnstr.constraint_mode(0);
endfunction : build_phase
```

If you do this, you'll find that the base test's build phase has already randomized the knob. This cannot be moved to the connect or end of elaboration phases, because it's necessary to randomize it in the build phase so that the rest of the components will build based on any configurations that might appear in `cfg`.

Fortunately, the base test was designed to solve this problem. The `cfg` randomization takes place in a virtual function called `randomize_cfg`, which you can override to do something before or after the randomization. So, in the `kval_test_c` class:

```
75.      veril/alutb/tests/kval.sv
      //////////////////////////////////////
      // func: randomize_cfg
      // Turn off the innocuous constraint
      virtual function void randomize_cfg();
          cfg.innocuous_cnstr.constraint_mode(0);
          super.randomize_cfg();
      endfunction : randomize_cfg
```

Naturally, your test will fail because your predictor doesn't know anything about the K or C configurations, yet. We'll deal with that in a little bit.

How It All Works

How did the register block know how to create NCB stores and loads for us? How does it know that the register isn't on the RSL interface instead?

The answer to these questions lie in the register block, the register maps, and the register adapters.

Investigate the `alu_csr_pkg::alu_ncb_reg_block_c::build()` function (in `verif/vcomponents/reg/obj/unit/reg__alu.v`), which configures and builds each of the CSRs and builds the NCB address map. The address map will be assigned a reference to the NCB sequencer and the NCB register adapter by the testbench.

This happens in the `base_test_c::connect_phase` function. It will be these classes that turn generic read and write transactions that the UVM register model knows about into the NCB stores and loads that we need.

The UVM 1.1 Class Reference and the User's Guide (section 5.5) have lots more detail on all of these.

Question

Why couldn't the register adapter have been specified in the register file or register block instead of in the base test?

Answer

Associating a CSR with an interface is a testbench function. In the full-chip testbench, these CSRs would not use the NCB to be written because that is an internal interface. Thus, this code does not belong in the vcomponent, which is meant for code that can be re-used at higher levels.

Fixing Our Predictor

Fixing the predictor to be CSR aware should be straightforward, as the predictor just needs to be aware of the CSR values at the time it makes the prediction. So, we can instantiate the register block in `alu_pred_c` and make it a configurable reference:

```
76.      verif/vcomponents/alu/alu_pred.sv
      `uvm_component_utils_begin(pred_c)
      `uvm_field_object(alutb_reg_block, UVM_REFERENCE)
      `uvm_field_int(ref_model,          UVM_ALL_ON)
      `uvm_component_utils_end

      //-----
      // Group: Configuration Fields

      // var: alutb_reg_block
      // Auto-generated Register Block
      alu_csr_pkg::alu_ncb_reg_block_c alutb_reg_block;
```

We chose to provide it with the macro flag `UVM_REFERENCE`, which will ensure that the object is populated by reference. It also means that the whole register file will not be printed out if we ever `sprint()` this predictor.

Problem 10-2

Now, modify the `write_item` function to consider both the K and C values that are present in the register block. Also, you'll need to populate the register block in the predictor.

Solution

Take special note of the fact that we get the current values from `uvm_reg_field` objects by looking at their `value` field.

```
77.      verif/vcomponents/alu/alu_pred.sv
virtual function void write_item(item_c _item);
    bit [7:0] k_val = alutb_reg_block.ALU_CONST.K_VAL.value;
    bit [7:0] c_val = alutb_reg_block.ALU_CONST.C_VAL.value;

    case(_item.operation)
        item_c::ADD_A_B : result = k_val * (_item.alpha + _item.beta) + c_val;
        item_c::SUB_A_B : result = k_val * (_item.alpha - _item.beta) + c_val;
        item_c::SUB_B_A : result = k_val * (_item.beta - _item.alpha) + c_val;
        item_c::MUL_A_B : result = k_val * (_item.alpha * _item.beta) + c_val;
        item_c::DIV_A_B : result = k_val * (_item.alpha / _item.beta) + c_val;
        item_c::DIV_B_A : result = k_val * (_item.beta / _item.alpha) + c_val;
        item_c::INC_A   : result = k_val * (_item.alpha + 1) + c_val;
        item_c::INC_B   : result = k_val * (_item.beta + 1) + c_val;
        item_c::CLR_RES : result = 0;
        item_c::ACCUM   : result += _item.alpha;
    endcase

    `cn_dbg(30, $psprintf("Calculated result %08X on item: %s", result, _item.convert2string()))
endfunction : write_item
```

Also note that when this function runs, it will fetch the predicted values based on the most recently completed CSR write. If a field is marked as being volatile, it means that it may change values between accesses (like a status register or a statistics counter).

One place to push the register block down to the ALU predictor would be in the build phase of the ALUTB environment class, which instantiates the ALU agent:

```
78.      verif/vcomponents/alutb/alutb_env.sv
      set_config_object("alu_agent.pred", "alutb_reg_block", reg_block, .clone(0));
```

Register Callbacks

There are many scenarios where you would want the environment to be alerted whenever a read or write occurs on a CSR. Here are just a few:

- Aliased registers. When one register is written, another register takes a new value.
- Expecting interrupts. When a CSR is written to, the write may conditionally cause an interrupt that needs to be expected.
- Soft resets. Writing a CSR causes a software reset, and some agents and predictors must be reset, too.
- Writes to the registers on-the-fly cause predictions to shift accordingly.

For cases such as these, you want to use a register callback. Register callbacks follow the same pattern that other callbacks in UVM follow:

1. Derive a class from `uvm_reg_cbs`.
2. Fill in the tasks and/or functions that supply the code you want to run.
3. Instantiate your callback class somewhere in the environment.
4. Register your callback class on the CSR(s) that it should apply to.

You may register multiple callback classes or multiple instances of a callback class on a given CSR.

The following is an example of a callback that prints the read value anytime the `ALU_RESULT[SOR]` field is read from.

```
79.
// class: alu_result_reg_cb_c
class alu_result_reg_cb_c extends uvm_reg_cbs;
  `uvm_object_utils(alu_result_reg_cb_c)

  //-----
  // Group: Methods
  function new(string name="alu_result_reg_cb");
    super.new(name);
  endfunction : new

  //////////////////////////////////////
  // func: post_read
  // Print to the logfile anytime the SOR CSR is read from.
  virtual task post_read(uvm_reg_item rw);
    `cn_info($psprintf("Read from SOR: %08X", rw.value[0]));
  endtask : post_read
endclass : alu_result_reg_cb_c
```

Then, the callback class is created and registered with the `ALU_RESULT` CSR anyplace and at any time in the environment.

```
80.
alu_result_reg_cb_c alu_result_reg_cb = new("alu_result_reg_cb");
uvm_reg_cb::add(reg_block.ALU_RESULT, alu_result_reg_cb);
```

Built-In Register Sequences

For free, UVM's register package includes a variety of register test sequences to completely exercise the CSR and memory space. These sequences can be the basis of your very first tests of a block to ensure that all CSR bugs are discovered and fixed. What's more, the sequences will automatically adapt to changes to the CSRs over time without any need to be modified.

The sequences available test the DUT's ability to withstand resets, bit-bashing, shared accesses (registers that can be accessed by more than one interface), memory walking patterns, and back-door HDL accesses.

Each sequence can be configured via the UVM resource database to turn on or turn off testing of specific CSRs or blocks.

Conclusion

UVM's register package offers significantly more functionality than what is presented here. UVM also has backdoor accesses, FIFO-like CSRs, memories, indirect registers, and more. But the information presented in this Lesson should provide a baseline for future lessons where these CSRs will be used.

Lesson 11: Writing Tests

When an environment is written properly, writing tests to hit specific regions of the design or areas of concern is merely a matter of adjusting a few knobs, tweaking a few configurations, or creating a few factory overrides. The base test instantiates everything, but all of the other tests should be easy to write.

CFG Classes

UVM uses the word configuration when referring to the configuration and resource database that is the backbone for storing values that trickle down to sub-components. However, we're going to talk about the random configurations that make each simulation job a little different.

There are myriad classes in a vcomponent where you can put random variables. Transactions, drivers, agents, sequences...all of the classes we've discussed so far can have random variables that change how they behave from one run to the next. And that's part of the problem: if you're writing a predictor that needs to be aware of all of these different values, then the predictor must have references to each of these components to have visibility into these values.

The K and C CSR values from Lesson 10 are a good example. We had to share the register file with the predictor so it could get access to these values. Far more complex scenarios can be imagined, though.

In our environment, all of these random variables--and the CSR register block--are bundled into a `cfg` class that is instantiated in the base test. A vcomponent's `cfg` class is similar to one of UVM's policy classes. Every class that instances this `cfg` class with the correct name will automatically get the reference to it during the build phase:

```
set_config_object("", "cfg", cfg, .clone(0));
```

What's more, all tests that you write can just add constraints directly into the `cfg` member field. The ALUTB package contains a `cfg_c` class that, admittedly, doesn't do a whole lot just yet. We're going to add some constraints to it to modify the values of K and C during the configuration phase and attempt to hit a wide swath of values.

Problem 11-1

Remove the `innocuous_cnstr` constraint from `alutb_pkg::cfg_c`. Create a configuration *knob* in the ALUTB configuration class that permits values of K and C to be

within the following values. By default, constrain the knob to always choose `INNOCUOUS` mode:

Setting	K	C	Frequency
INNOCUOUS	1	0	20%
SMALL	2..5	2..10	50%
LARGE	6..50	11..128	15%
XLARGE	51..255	129..255	5%
UNLIMITED	1..255	0..255	10%

Then, write a test called `exer_test_c` that varies the knob settings according to the frequencies shown above. You should find that it passes consistently.

Solution

After removing the ALU base constraints, you start by creating an enumerated type that represents the knob. By placing this type in the configuration class, you do not pollute the namespace of the environment. This should go at the top of the `cfg_c` class:

```
81.    verif/vcomponents/alu/alu_cfg.sv
class cfg_c extends uvm_component;
    //-----
    // Group: Types

    // enum: alu_const_knob_e
    // Used to constrain the K and C values for ALU
    typedef enum { INNOCUOUS, SMALL, LARGE, XLARGE, UNLIMITED } alu_const_knob_e;
```

Then, you create a random instance of your new enumerated type, and constrain it to the value `INNOCUOUS`.

```
82.    verif/vcomponents/alu/alu_cfg.sv
`uvm_component_utils_begin(alutb_pkg::cfg_c)
`uvm_field_object(reg_block, UVM_REFERENCE)
`uvm_field_enum(alu_const_knob_e, alu_const_knob, UVM_ALL_ON)
`uvm_component_utils_end

//-----
// Group: Configuration Fields

// var: reg_block
// Register block for this environment
rand alu_csr_pkg::alu_ncb_reg_block_c reg_block;

// var: alu_const_knob
// Constrains the K_VAL and C_VAL
rand alu_const_knob_e alu_const_knob;
constraint const knob_cnstr { alu_const_knob == INNOCUOUS; }
```

Note the syntax of the ``uvm_field_enum` macro requires that you provide the type as the first argument. Now, add implication constraints to K and C that match the table:

```
83.    verif/vcomponents/alu/alu_cfg.sv
// constrain K_VAL based on alu_const_knob
constraint kval_cnstr {
    alu_const_knob == INNOCUOUS -> (reg_block.REG_FILE.ALU_CONST.K_VAL.value == 1);
    alu_const_knob == SMALL      -> (reg_block.REG_FILE.ALU_CONST.K_VAL.value inside {[2:5]});
    alu_const_knob == LARGE      -> (reg_block.REG_FILE.ALU_CONST.K_VAL.value inside {[6:50]});
    alu_const_knob == XLARGE     -> (reg_block.REG_FILE.ALU_CONST.K_VAL.value inside {[51:255]});
    alu_const_knob == UNLIMITED -> (reg_block.REG_FILE.ALU_CONST.K_VAL.value inside {[0:255]});
}

// constrain C_VAL based on alu_const_knob
constraint cal_cnstr {
    alu_const_knob == INNOCUOUS -> (reg_block.REG_FILE.ALU_CONST.C_VAL.value == 0);
    alu_const_knob == SMALL      -> (reg_block.REG_FILE.ALU_CONST.C_VAL.value inside {[2:10]});
    alu_const_knob == LARGE      -> (reg_block.REG_FILE.ALU_CONST.C_VAL.value inside {[11:128]});
    alu_const_knob == XLARGE     -> (reg_block.REG_FILE.ALU_CONST.C_VAL.value inside {[129:255]});
    alu_const_knob == UNLIMITED -> (reg_block.REG_FILE.ALU_CONST.C_VAL.value inside {[0:255]});
}
```

The `exer_test_c` class uses a distribution constraint to randomize the knob. Notice that you can add a constraint *through* the class instance hierarchy. This is essential. If you

couldn't do this in SystemVerilog, the only other way to do it would be by extending the class with another one that contained the constraint, and then overriding that type using the factory. You'll find this method to be a lot less typing.

Here is the complete `exer_test_c` class:

```
84.      verif/alutb/tests/exer.sv
// class: exer_test_c
// Turns on all possible values of K_VAL and C_VAL
class exer_test_c extends basic_test_c;
    `uvm_component_utils(exer_test_c)

    //-----
    // Group: Configuration Fields
    constraint const_knob_cnstr {
        cfg.alu_const_knob dist { alutb_pkg::cfg_c::INNOCUOUS :/ 20,
                                alutb_pkg::cfg_c::SMALL      :/ 50,
                                alutb_pkg::cfg_c::LARGE       :/ 15,
                                alutb_pkg::cfg_c::XLARGE      :/ 5,
                                alutb_pkg::cfg_c::UNLIMITED   :/ 10
        };
    }

    //-----
    // Group: Methods
    function new(string name="alu_exer",
                uvm_component parent=null);
        super.new(name, parent);
    endfunction : new
endclass : exer_test_c
```

The `exer` test now overwrites the configuration class's `const_knob_cnstr` constraint with its new one. Because the enumerated type is local to the scope of the configuration class, (which is local to the scope of the ALUTB package), you need to use the scoping operator twice.

Test Overrides

Sometimes you want to create a sequence for a semi-directed test case. Perhaps you want to create an error case that you don't want intermingling with your other sequences. Or you want your driver to behave differently for a single test.

Once again, this is where the factory comes in. It is common practice to create a derived class of something else in the environment and place it in the test file, then use a factory override for that test only.

Problem 11-2

Create a new test, `div0_test_c`, that sets the denominator of a divide operation to zero 50% of the time. Do this by overriding the `alu_pkg::item_c` with a new transaction type, `div0_item_c`. See what happens.

Solution

Here is one solution to the problem:

```

85.      verif/alutb/tests/div0.sv
`include "basic.sv"

//*****
// class: div0_item_c
// Causes a divide-by-zero on 50% of all divide operations
class div0_item_c extends alu_pkg::item_c;
    `uvm_object_utils(div0_item_c)

    constraint protocol_cnstr {
        (operation == DIV_A_B) -> beta dist { 0 :/ 50,
                                                [1:'hffff] :/ 50};
        (operation == DIV_B_A) -> alpha dist { 0 :/ 50,
                                                [1:'hffff] :/ 50};
        (operation == SUB_A_B) -> alpha > beta;
        (operation == SUB_B_A) -> beta > alpha;
    }

    //-----
    // Group: Methods
    function new(string name="div0_item");
        super.new(name);
    endfunction : new
endclass : div0_item_c

//*****
// class: div0_test_c
// Uses the div0_item_c class instead
class div0_test_c extends basic_test_c;
    `uvm_component_utils(div0_test_c)

    //-----
    // Group: Methods
    function new(string name="div0_test",
                  uvm_component parent=null);
        super.new(name, parent);
    endfunction : new

    //////////////////////////////////////
    // func: build_phase
    virtual function void build_phase(uvm_phase phase);
        super.build_phase(phase);
        set_type_override_by_type(alu_pkg::item_c::get_type(), div0_item_c::get_type());
    endfunction : build_phase

endclass : div0_test_c

```

Some notes:

```

constraint protocol_cnstr {
    (operation == DIV_A_B) -> beta dist { 0 :/ 50,
                                          [1:'hffff] :/ 50};
    (operation == DIV_B_A) -> alpha dist { 0 :/ 50,
                                          [1:'hffff] :/ 50};
    (operation == SUB_A_B) -> alpha > beta;
    (operation == SUB_B_A) -> beta > alpha;
}

```

This changes the constraints to have a zero in the denominator 50% of the time for division operations. Note how we've used the distribution constraint on top of an

implication constraint. It was also necessary to put the subtraction constraints in because we cannot partially override a constraint block.

And finally, the factory override method should now be familiar to you:

```
set_type_override_by_type(alu_pkg::item_c::get_type(), div0_item_c::get_type());
```

If your first inclination was to only perform the factory override 50% of the time, there's really no effective way to do that. Using constraints and overriding all transactions is the way to go.

If you did everything correctly, you should find that the RTL puts X's out as the result and the test fails. Congratulations! You've found your first bug. The designer agrees and decides to use a denominator of 1 whenever a zero is seen instead.

You can patch the RTL with the `utut` script, like this:

```
verif/alutb> utut fix 1
```

Question

Now if you run your test again, you'll find that the test passes. But if you wrote your predictor the same way as the solution from Lesson 9, there's a problem in there someplace, can you find it and fix it?

Hint:

Look in the logfile. It's a problem that also plagues designers.

Answer

Only if you looked at your logfile would you likely find the problem. It's that the predictor we wrote also doesn't have the fix and is calculating the result to be X. But the result is no longer X and yet the test still passes. What's going on here?

There are actually two problems. The first is that the predictor is predicting X, and the second is that it's not flagging an error.

If you recall, our comparison took place in the TLM imp function, `write_result()`:

```
virtual function void write_result(result_t _result);
    if(_result != result)
        `cn_err($psprintf("Actual result: %08X != Expected result: %08x",
                           _result, result))
endfunction : write_result
```

Here, we've used an `if` statement, and one of the operands is an X. Just as in RTL, verification people have to beware that the `if` statement does not propagate an X. The above test will not succeed, and no error gets printed.

Of course, we need to modify the calculations to now match the RTL. But we want to make sure that we don't miss this issue again next time.

To fix this you could use the `else` clause:

```
if(_result == result)
;
else
    `cn_err($psprintf("Actual result: %08X != Expected result: %08x",
                      _result, result))
```

But this is a pretty terrible solution. Fortunately, though, the `assert` statement does not have the same affliction as the `if` statement. This is the better solution:

```
86.      verif/vcomponents/alu/alu_pred.sv
      assert(_result == result) else
          `cn_err($psprintf("Actual result: %08X != Expected result: %08x",
                              _result, result))
```

And to fix the ALU predictor, you would change the two cases as follows:

```
87.      verif/vcomponents/alu/alu_pred.sv
      item_c::DIV_A_B : result = k_val * (_item.alpha / (_item.beta? _item.beta : 1)) + c_val;
      item_c::DIV_B_A : result = k_val * (_item.beta / (_item.alpha? _item.alpha : 1)) + c_val;
```

Conclusion

While the rest of our jobs may be overly complicated, writing tests should be kept simple. Vcomponents should each have a policy class called `cfg`, instantiated in the base test and distributed to its components through the configuration database. Tests should be created

with the template generator. Tweak it with some configurations and knobs, and maybe a factory override or two, and you're done.

Lesson 12: Advanced Sequences I

Sequences and sequencers are so involved, they deserve three separate lessons! This lesson will discuss how to better launch sequences to get a more random mix of traffic, how we can use the resource database to configure sequences, and what a persistent sequence is.

Default Sequences

So far we've seen two methods that we can use to launch sequences onto a sequencer. From a component, we can create a sequence and call its start task, with a reference to the sequencer as an argument. From another sequence, we can have it launch another sequence by calling one of the ``uvm_do` macros.

A second way that you can launch sequences from a component is by calling the sequencer's task `execute_item()` with an instance of the sequence. This isn't a whole lot different from calling the sequence's start task, so there's not much gain here.

But there's a third way as well. When any phase of a sequencer is started, its function `start_phase_sequence` is run. If there is a configuration item in the database called "default_sequence" for this sequencer's phase, then that sequence is run automatically.

For example, to automatically launch a `sum_array_seq_c` during the main phase, you would put this in your test's build phase:

```
uvm_config_db#(uvm_object_wrapper)::set(this, "alutb_env.alu_agent.sqr.main_phase",  
                                         "default_sequence", alu_pkg::sum_array_seq_c::type_id::get());
```

This will create, randomize, and launch a sequence of this type for you when the main phase starts. If you want a very specific sequence to run, then the call is slightly different:

```
alu_pkg::sum_of_factorials_seq_c sof_seq = new("sof");  
sof_seq.randomize() with {op_x == 1; op_y == 5;};  
uvm_config_db #(uvm_sequence_base)::set(this, "alutb_env.alu_agent.sqr.main_phase",  
                                         "default_sequence", sof_seq);
```

Here, notice that we did not just 'set' `op_x` and `op_y`. We needed to randomize it for this to work, because UVM will randomize it for us if it hasn't already been randomized, which would then squash our settings.

Regardless of the method you choose, you're probably wondering where all the benefit is. Sure, it allows you to use the configuration database to start these sequences, but it only allows us to launch one sequence per phase. What good is that?

Well, if that sequence is one that happens to launch many other sequences, then it can be used to concisely do a whole lot more than what our basic test's `main_phase` currently does. That's where library sequences come into play.

Library Sequences

Before we get too far into library sequences, it should be mentioned that these are new to UVM 1.1 and are not considered "production-level". As such, they are also not particularly well documented yet, either.

A library sequence is a sequence that picks a random sequence, launches it, then picks another, and another, and another, until a specified count is reached. This is an ideal method for creating highly random stimulus as it allows you to create a new sequence to test some feature, write a test that focuses on that sequence, and then add it to one or more library sequences to see how well it works with others.

You might recall that when you created `alu_seq_lib.sv` with `utg` that it created a class called `lib_seq_c`, and we had you comment it out. Those few lines of code are a library sequence. You can create as many library sequences as you wish. And each library can have as many sequences as you wish. To place a sequence into a library, you use the macro ``uvm_add_to_seq_lib`. For example, to add our factorial sequence to the `lib_seq_c` library, you would add the following to the `factorial_seq_c` class:

```
88.      `verif/vcomponents/alu/alu_seq_lib.sv
class factorial_seq_c extends uvm_sequence #(item_c, item_c);
    `uvm_object_utils_begin(alu_pkg::factorial_seq_c)
        `uvm_field_int(operand, UVM_ALL_ON)
        `uvm_field_int(result, UVM_ALL_ON)
    `uvm_object_utils_end
    `uvm_add_to_seq_lib(factorial_seq_c, lib_seq_c);
```

This macro must specify the type of the sequence and which library you want it to go into. A sequence can call this macro more than once to be placed into different libraries. You can create different library sequences to mix different types of sequences. Your test can then set a library sequence as the default sequence for a particular phase, such as the main phase.

Problem 12-1

Add all of the sequences in your library sequence to the `lib_seq_c`.

Clear out the run phase of your basic test, and launch your new library sequence instead.

Solution

To launch the sequences from your basic test, you should have altered its build phase as follows.

```
89.      verif/alutb/tests/basic.sv
      virtual function void build_phase(uvm_phase phase);
      super.build_phase(phase);

      uvm_config_db#(uvm_object_wrapper)::set(this, "alutb_env.alu_agent.sqr.main_phase",
      "default_sequence", alu_pkg::lib_seq_c::type_id::get());
      endfunction : build_phase
```

When you run this test, you should see 10 different random sequences chosen. Of course, most of those sequences spawn other sequences and sequence items, so a lot more traffic will be seen.

Now, all of your tests that derive from `basic_test_seq_c` will run your library sequence by default.

Library Sequence Configuration

Our library sequence only runs 10 random sequences. And it's always random. Maybe you want a more interesting pattern to choose from? Once again, UVM uses the policy class pattern to help out.

The `uvm_sequence_library_cfg` class is the policy class that you can use to customize your library sequence. It offers three different options: `selection_mode`, `min_random_count`, and `max_random_count`.

The min and max random count options do just what you think. When the sequence starts, it chooses a random number between these two numbers, and then sends that many sequences.

The selection mode is an enumerated type that offers you four choices as to how sequences will be chosen:

uvm_sequence_lib_mode	Method
UVM_SEQ_LIB_RAND	Select randomly (default)
UVM_SEQ_LIB_RANDC	Random cyclic selection
UVM_SEQ_LIB_ITEM	Send only items. Does not send sequences.
UVM_SEQ_LIB_USER	Applies a user-defined selection algorithm.

Random cyclic is a term we haven't used so far in this tutorial. It tells SystemVerilog to choose randomly from a list of items, but don't choose an item again until the list has been exhausted.

Sending only sequence items is a very useful scenario. Essentially, the library sequence will just send random transactions, ignoring whatever sequences you put into it. Sequence items, though, get responses. And unfortunately, if you choose this selection mode, you will find that the library sequence's response queue fills up. To manage this, you would need to implement a response handler function that was described in Lesson 7.

The user-defined algorithm has you create a function (`select_sequence`) that returns an unsigned integer between zero and its argument, `max`. At your disposal is the member field `sequences`, which is a queue of all the different sequences in the library, stored as `uvm_object_wrapper` classes. Calling `get_type_name()` on each member gives the sequence's name as a string. The following example shows how you might distribute sequences in a library unevenly:

```
90.      verif/vcomponents/alu/alu_seq_lib.sv
class lib_seq_c extends uvm_sequence_library #(item_c);
    `uvm_object_utils(alu_pkg::lib_seq_c)
    `uvm_sequence_library_utils(lib_seq_c)

    //-----
    // Group: Fields

    // var: selector
    // Distributes sequence selection unevenly when UVM_SEQ_LIB_USER is used
    rand int unsigned selector;
    constraint selector_cnstr { selector dist { 0 :/ 40, 1 :/ 20, 2 :/ 10, 3 :/ 30 }; }

    //-----
    // Group: Methods
    function new(string name="seq_lib");
        super.new(name);
        init_sequence_library();
    endfunction : new

    //////////////////////////////////////
    // func: select_sequence
    virtual function int unsigned select_sequence(int unsigned max);
        randomize() with { selector < max; };
        return selector;
    endfunction : select_sequence

endclass : lib_seq_c
```

Library sequences offer several more features that allow you to modify the list of sequences dynamically, such as `add_sequence`, `remove_sequence`, and `add_typewide_sequence`, which allows you to add a sequence type to all your library sequences at one time.

Persistent Sequences

A persistent sequence is one that starts during some phase (at time zero, perhaps) and remains active for the remainder of the simulation. Why would you want such a thing?

Well, we know that a sequencer can have multiple active sequences, so it's not as though your persistent sequence is the only thing happening. Typically, your sequence will wait some period of time, or wait for some event to happen, and then will spring into action. Here are a few possible examples of persistent sequences:

- **Random Read Sequence.** Picks a random register to safely read, reads it, then waits for a random period of time.
- **Interrupt Handler.** Waits for some interrupt event to occur, locks down the sequencer, and begins the process of discovery and clearing of interrupts.
- **Sequence Adapter.** Let's say that you had one sequencer that was spitting out messages and was feeding requests into your ALU agent for mathematical operations. The writer of that agent doesn't know about our `item_c` class, so is sending their own sequence items that need to be converted into `item_c`.

The ALU model has a second CSR called `ALU_RESULT`. This CSR has a 32-bit read-clear field called `SOR` which contains the sum of all results since it was last read. We're going to add a sequence that reads and clears the `SOR` register every time we believe that it should exceed `32'h100_000`, and we will check to ensure that it does.

Here's how we'll do it:

1. Pipe all monitored results into the ALU sequencer and put each one in a mailbox.
2. Run a sequence starting at time zero that pulls from the mailbox and accumulates the expected `SOR` value.
3. If the expected `SOR` value exceeds `32'h100_0000`, then read the CSR to clear it.
4. Ensure that the `SOR` field read a value greater than `32'h100_0000`.
5. Clear the expected `SOR` value that's been accumulated so far.

There are a whole slew of new concepts above.

The first one is the mailbox, which is a handy SystemVerilog construct. It is a parameterized object that has `put` and `get` tasks to let you stuff objects in one place and pull them out of another. It also has zero-time functions `try_put` and `try_get`. UVM

offers you the choice of using a TLM FIFO here as well, but the mailbox will work fine for our purposes.

The next concept is having the sequence pull from its sequencer's mailbox. How can a sequence get a handle to a mailbox in the sequencer? Better yet: how will it read from a CSR?

The answer is sequencer references.

Sequencer References

The `uvm_sequence` base class comes built-in with a reference back to the sequencer it's running on. This reference is called `m_sequencer`. Unfortunately, because it is declared in the base class as being of the type `uvm_sequence`, it's not very useful by itself. After all, the base class doesn't have a mailbox of results.

But, we could create our own reference and cast the `m_sequencer` back to the `alu_pkg::sqr_c` that we know it to be, like this:

```
sqr_c my_sqr;  
$cast(my_sqr, m_sequencer);
```

It turns out that this is a pretty common thing to do, so UVM automated it. We have written plenty of sequences so far that haven't needed to reference the sequencer, so it's an optional macro that you add to the class:

```
`uvm_declare_p_sequencer(sqr_c)
```

This macro call creates a reference, `p_sequencer`, that points to our sequencer and knows what kind it is.

Before, all of our sequences were detached from the component hierarchy. If we wanted to interact with any other components in the environment, we were out of luck. But the `p_sequencer` handle *anchors* our sequence into the hierarchy. This allows us to do *lots* of nifty new things with sequences.

In addition to the tiny bit of overhead, the only drawback to declaring a `p_sequencer` handle with the macro is that now our sequence can only be run on this specific sequencer type.

Sequences can be run on different sequencers?!?! Yes, they can. And they often do. That's a topic we'll explore some more in Lesson 14. Meantime, let's see how this would all fit together. It's somewhat complicated at first, so we'll do it in stages.

Problem 12-2

Add an analysis imp to your sequencer that accepts items of type `result_t`. Also, create a mailbox that holds the same type. Connect it to the analysis port from the monitor. Print out the result each time one is received so you know that it's working.

Solution

To add the analysis imp and mailbox to the sequencer, you would put these in all the right places. In `sqr_c`:

```
91.      verif/vcomponents/alu/alu_sqr.sv
      uvm_analysis_imp #(result_t, sqr_c) monitored_result_imp;
      mailbox #(result_t) monitored_result_mbox;
```

Just as with TLM ports, you also need to `new()` your mailboxes. They accept an argument that would limit the number of entries, but the default of zero permits an infinite number, and that should work fine for us:

```
92.      verif/vcomponents/alu/alu_sqr.sv
      virtual function void build_phase(uvm_phase phase);
      super.build_phase(phase);
      monitored_result_imp = new("monitored_result_imp", this);
      monitored_result_mbox = new();
      endfunction : build_phase
```

Now, the `write` function for the analysis imp:

```
93.      verif/vcomponents/alu/alu_sqr.sv
      virtual function void write(result_t _result);
      monitored_result_mbox.try_put(_result);
      `cn_info($sprintf("Received result %08X", _result))
      endfunction : write
```

The one and only challenge here was the use of `try_put`. Because `write` is always a function, and because a `put` to a mailbox could consume time (if it were full, for example), then we must use `try_put`. And because it's an infinitely large mailbox, `try_put` will always succeed.

Connecting to the monitor's port in the agent should also be straightforward:

```
94.      verif/vcomponents/alu/alu_agent.sv
      if(!ref_model) begin
        if(is_active) begin
          drv.seq_item_port.connect(sqr.seq_item_export);
          mon.monitored_result_port.connect(sqr.monitored_result_imp);
        end
      end
```

Be sure to only call this when both the monitor and the sequencer exist, which would be when the agent is active and not acting as a reference model.

Handling Cyclical Dependencies

In the next problem, you're going to write a sequence that declares its sequencer using ``uvm_declare_p_sequencer`. But we're also going to start it from that very sequencer. This will create a cyclical dependency, because they both must know about each other.

Fortunately, SystemVerilog supports forward references and late binding. Simply adding a forward reference above your sequences in `alu_seq_lib.sv` tells the compiler: “It’s ok, I know what I’m doing here.”

```
95.    verif/vcomponents/alu/alu_seq_lib.sv  
    // Forward declaration of sequencer  
    typedef class sqr_c;
```

Now, the compiler sees the ``uvm_declare_p_sequencer(sqr_c)` and knows that `sqr_c` refers to some class, which will later be declared.

Problem 12-3

So, let’s write our persistent sequence, `sor_clr_seq_c`. It should constantly try to pull results from its sequencer’s mailbox and accumulate them. Hold off on actually performing the CSR read for now. When the sum reaches or exceeds the value `'h100_0000`, announce with a print that it’s time to read the CSR. Then clear your accumulated value.

Have the sequencer itself start this sequence. Run the basic test to ensure that it’s doing its job.

Solution

Here is the sequence that does the job:

```

96.      verif/vcomponents/alu/alu_seq_lib.sv
class sor_clr_seq_c extends uvm_sequence#(item_c);
  `uvm_object_utils(alu_pkg::sor_clr_seq_c)
  `uvm_declare_p_sequencer(sqr_c)

  //-----
  // Group: Methods
  function new(string name="sor_clr_seq");
    super.new(name);
  endfunction : new

  //////////////////////////////////////
  // func: body
  virtual task body();
    result_t accum_result = 0;
    result_t new_result;

    forever begin
      // get the next monitored transaction and add it to the accumulated results
      p_sequencer.monitored_result_mbox.get(new_result);
      accum_result += new_result;
      `cn_dbg(30, $psprintf("accum_result = %8X", accum_result))

      if(accum_result >= 32'h100_0000) begin
        `cn_info("Time to read the RESULT_CSR!")
        // clear out accum_result
        accum_result = 0;
      end
    end
  endtask : body
endclass : sor_clr_seq_c

```

The solution declares the `p_sequencer` handle using the macro. It runs in a `forever` loop and calls `get` to fetch the latest result from the sequencer, using our `p_sequencer` handle.

And here is how the sequence is launched from the sequencer's `run_phase`:

```

97.      verif/vcomponents/alu/alu_sqr.sv
virtual task run_phase(uvm_phase phase);
  sor_clr_seq_c sor_clr_seq = sor_clr_seq_c::type_id::create("sor_clr_seq");
  sor_clr_seq.start(this);
endtask : run_phase

```

When you start a sequence, you have to give it a handle to the sequencer that it's supposed to run on, which just like in C++ is represented by the `this` field.

Problem 12-4

Now, replace your print statement with the actual read of the CSR. Follow that up with a check that ensures that the sum actually does exceed ``h100_0000`. How do you go about accessing the CSR from a sequence?

Solution

If you did all of that correctly, it's entirely possible that your test will fail. We'll discuss why it might fail and what can be done about it soon. But first a solution.

The CSR register blocks and files are all within the component hierarchy. So to use these, the sequence will have to reach down into its own sequencer via the `p_sequencer` handle. We were able to get the register block in the predictor, so we should be able to do that here, too. Start by declaring the register block in the sequencer:

```
98.    verif/vcomponents/alu/alu_sqr.sv
    // var: reg_block
    // Auto-generated Register block
    alu_csr_pkg::alu_ncb_reg_block_c reg_block;
```

Now that the register block is in the sequencer, the sequence can easily access it. Here is what was done to its body task. All of the changes are in bold:

```
99.    verif/vcomponents/alu/alu_seq_lib.sv
    virtual task body();
    result_t accum_result = 0;
    result_t new_result;
    uvm_status_e status;
    uvm_reg result_reg = p_sequencer.reg_block.ALU_RESULT;
    uvm_reg_data_t sor_value;

    forever begin
        // get the next monitored transaction and add it to the accumulated results
        p_sequencer.monitored_item_mbox.get(new_result);
        accum_result += new_result;
        `cn_dbg(30, $psprintf("accum_result = %8X", accum_result))

        if(accum_result >= 32'h100_0000) begin
            result_reg.read(status, sor_value);

            // ensure status was ok
            if(status == UVM_NOT_OK) begin
                `cn_err("Unable to read from ALU_RESULT register.")
            end else begin
                // ensure that current value exceeds 'h100_0000
                `cn_dbg(30, $psprintf("SOR read as %08X", sor_value))
                if(sor_value < 'h100_0000)
                    `cn_err($psprintf("Read from ALU_RESULT[SOR] but its value was %08X", sor_value))

                // clear out accum_result
                accum_result = 0;
            end
        end
    end
    endtask : body
```

We've been using them all along, but it's worth repeating: SystemVerilog uses implicit handles all over the place. This bit of code doesn't create anything new, it just assigns a handle to the register:

```
uvm_reg result_reg = p_sequencer.aluth_reg_block.ALU_RESULT;
```

If instead of calling `read` you called `mirror` and then look at the CSR's value, you would have seen a value of zero, which is not the value that was read. This is because the `mirror` task performs the read, then updates its expected value of the CSR. Because it's declared as read-clear ("RC"), it will now be zero.

Calling `read`, though, actually gives you the value that was read. And, it's always safe to check that the status of the read was OK.

Now, if you did all of this correctly, your test may still fail--under the right circumstances--with an error like this one:

```
%E-( _alu_seq_lib.sv: 312) [alutb_env.alu_agent.sqr.sor_clr_seq]({__1008ns) Read from
ALU_RESULT[SOR] but its value was 000215b3
```

How is it that we did not read out something greater than 'h100_0000'?

The answer lies in the ordering of sequences and the way the RTL was written. We submitted a read to the sequencer that handles NCB transactions, but these take a certain amount of time to occur. In the meantime, other ALU traffic is going on. If a read happens in the exact same cycle that a new result occurs, then the SOR field will not go to zero, but will go to the value of the new result, which is what you would want in an accumulating CSR. Because the reads take an indeterminate amount of time, multiple results may have been put into the mailbox before the first read completes, and has cleared out the CSR value. Then, your sequence reads all of these mailbox entries in zero time and quickly accumulates beyond the threshold, but is now out of sync with the real value.

The solution to this problem is to stop the sequencer from issuing new ALU traffic while an NCB read is happening. Previously, we had called `lock()` and `unlock()` to give our sequence exclusive access. But, locking a sequencer actually isn't even fast enough, because `lock` can be preempted by other traffic--potentially *lots* of other traffic. When our sequence needs to read from the CSR, it needs to do so before any more ALU requests go through. The `grab()` and `ungrab()` functions will do the trick, because these are prioritized over all other traffic.

Also, because the read consumes time, it is possible that any sequences which had launched prior to the `grab()` would now have their results in the results mailbox, these should be cleared out, too. Otherwise, our accumulated result will not be correct.

If you add these to your sequence just before and just after the read, you will find that your test now passes. Below is the complete solution:

```
100.    verify/vcomponents/alu/alu_seq_lib.sv
    virtual task body();
        result_t accum_result = 0;
        result_t new_result;
```



```

uvm_status_e status;
uvm_reg result_reg = p_sequencer.reg_block.ALU_RESULT;
uvm_reg_data_t sor_value;

forever begin
    // get the next monitored transaction and add it to the accumulated results
    p_sequencer.monitored_result_mbox.get(new_result);
    accum_result += new_result;
    `cn_dbg(30, $psprintf("accum_result = %8X", accum_result))

    if(accum_result >= 32'h100_0000) begin
        grab();
        result_reg.read(status, sor_value);

        // ensure status was ok
        if(status == UVM_NOT_OK) begin
            `cn_err("Unable to read from ALU_RESULT register.")
        end else begin
            // ensure that current value exceeds 'h100_0000
            `cn_dbg(30, $psprintf("SOR read as %8X", sor_value))
            if(sor_value < 'h100_0000)
                `cn_err($psprintf("Read from ALU_RESULT[SOR] but its value was %8X", sor_value))

            // clear out accum_result
            accum_result = 0;

            // empty out mailbox in case any have been added since the read started
            while(p_sequencer.monitored_result_mbox.try_get(new_result))
                ;
        end
        ungrab();
    end
end
endtask : body

```

Sequence Configurations

Let's say that we didn't always want to perform the read on such a specific value as 'h100_0000. Maybe we want to make our sequence *configurable*.

We can configure components through the resource database, so why not sequences as well? Well, sequences don't have a build phase like components do, so they don't benefit from the magic call to `super.build_phase()` to automatically populate their configurable fields.

You could put explicit set and get calls in your code, like this set in the `build_phase` of the basic test:

```

101.   verif/alutb/tests/basic.sv
      uvm_resource_db#(alu_pkg::result_t)::set("alu::sor_clr_seq", "trigger_value", 'h300_0000);

```

And then by placing the get call in the `sor_clr_seq_c`:

```

102.   verif/vcomponents/alu/alu_seq_lib.sv
      trigger_value = uvm_resource_db#(result_t)::get_by_name("alu::sor_clr_seq",
                                                             "trigger_value").read();

```

While this works, these calls are not very attractive. Alternatively, you can have this trigger value be a standard configuration field in your *sequencer*, and then access it from the sequence using the `p_sequencer` handle.

This gives you a simple and consistent method for setting knobs and configurations for your sequences. Just be sure not to over-pollute the sequencer with fields that sequences may conflict upon.

Conclusion


Would it surprise you to realize that we've only really just scratched the surface of what sequences and sequencers are capable of? Wait until they're *virtualized*!

Lesson 13: A New Testbench

It shouldn't surprise you that our ALU module doesn't live all by itself. All block-level testbenches are to be reused in a higher level environment, and the ALU is no different. Fortunately, the ALU subsystem isn't a whole lot more complicated, and most of the work has already been done for you. Integrating your ALU vcomponent with another vcomponent in a new testbench is the objective of this lesson.

The ALU Framer

This diagram shows the complete ALU block, which consists of the unit you have been verifying, together with its cousin the framer block. ALU results are pushed into the framer's FIFO one at a time. A frame length (frame_len, valid on the clock cycle in which frame_len_val is high) is specified to the framer. When the depth of the



FIFO reaches the frame length, all of the results are pushed out as a complete frame, one 32-bit result word at a time. The framer can only accept one frame length request at a time. Once the frame has been completed, a new frame length can be specified.

The FIFO is 32 quadwords deep, the maximum length of a frame. If it begins to fill up, a frame_bp signal will be asserted high, indicating that no more ALU transactions should be sent in.

Creating the FRM Testbench

The objective of this lesson is merely to create the new testbench. Fortunately, a script will get you most of the way there. And, the framer vcomponent is already available to you in `verif/vcomponents/frm`.

But first, let's see what the environment will look like by the time we're done. We're going to look at it using the topology report that you may have seen at the beginning of your logfiles. Some lines have been deleted for the sake of brevity:

Name	Type	Size	Value
<unnamed>	uvm_root	-	@162
uvm_test_top	basic_test_c	-	@596
env	frm_pkg::env_c	-	@669
alu_agent	alu_pkg::agent_c	-	@738
drv	alu_pkg::drv_c	-	@761
mon	alu_pkg::mon_c	-	@753
monitored_result_port	uvm_analysis_port	-	@919
monitored_item_port	uvm_analysis_port	-	@910
pred	pred_c	-	@948
sqr	alu_pkg::sqr_c	-	@787
item logger	vm_component	-	@931
frm_agent	frm_pkg::agent_c	-	@712
drv	frm_pkg::drv_c	-	@1134
mon	frm_pkg::mon_c	-	@1126
sqr	frm_pkg::sqr_c	-	@1160
ncb_env	ncb_pkg::env_c	-	@724
cfg	frm_pkg::cfg_c	-	@622
reg_block	alutb_pkg::reg_block_c	-	@623
global env	global_pkg::env_c	-	@614
watchdog	global_pkg::watchdog_c	-	@2790
tb_clk_drv	clk_drv_c	-	@683
cfg	frm_pkg::cfg_c	-	@622
reg_block	alutb_pkg::reg_block_c	-	@623
reg_block	alutb_pkg::reg_block_c	-	@623
ALU_CONST	alu_const_reg_c	-	@629
C_VAL	uvm_reg_field	...	RW ALU_CONST[7:0]=8'h2a (Mirror: 8'h00)
K_VAL	uvm_reg_field	...	RW ALU_CONST[15:8]=8'h67 (Mirror: 8'h00)
RSVD0	uvm_reg_field	...	RO ALU_CONST[63:16]=48'h000000000000
ALU_RESULT	alu_result_reg_c	-	@634
SOR	uvm_reg_field	...	RC ALU_RESULT[31:0]=32'h00000000
RSVD0	uvm_reg_field	...	RO ALU_RESULT[63:32]=32'h00000000
NCB	uvm_reg_map	-	@625

UVM's handy topology report is an easy way to see the structure of the testbench, how random configurations were chosen, and what the values of all the CSRs will be after configuration.

As you can see, it looks pretty similar to the ALUTB environment, with the addition of the framer agent, also written for you.

UVM New Testbench (untb)

Like utg, the untb script will generate a lot of the boilerplate code for you. It's quite simple to use:

```
verif> untb frm
++ Creating Testbench verific/frm
++ Creating frm_tb_top.sv
++ Creating base_test.sv
++ Creating basic.sv
```

That was simple. Unfortunately, there's still quite a bit of work to do before you'll be able to run it.

The script creates the standard Makefile, flists, the top-level testbench, and the base test and basic test. The first thing to do is add the vcomponents this testbench will need. You'll want to change the Makefile so that the FLISTS variable is set like this:

```
103.    verif/fm/Makefile
FLISTS= verif/vcomponents/cn/cn.flist \
        verif/vcomponents/global/global.flist \
        verif/vcomponents/ncb/ncb.flist \
        verif/vcomponents/alu/alu.flist \
        verif/vcomponents/alutb/alutb.flist \
        verif/vcomponents/fm/fm.flist \
        verif/fm/fm.flist \
        verif/fm/rtl.flist
```

You'll also want to change the CSR packages list to point to the one in the ALUTB vcomponent, since we're just reusing that one:

```
104.    verif/fm/Makefile
CSR_PKGS = $(ROOT_DIR)/verif/vcomponents/reg/obj/unit/regs_alu.sv
```

The `rtl.flist` file can just be copied from the one in `verif/alutb/rtl.flist` since both testbenches will be using the same RTL.

Problem 13-1

Modify the top-level testbench, `fm_tb_top.sv`, to create all interfaces and instantiate the DUT as needed. You may use the `alu_wrapper.sv` file from the ALUTB testbench, or you can instantiate the `alu` directly in the top-level. The testbench will be largely the same as the one used by ALUTB with the addition of the `fm_intf` interface.

You will also need to make some changes to the base test. Namely, the `reg_block` will be in the `alutb_pkg`, the NCB interface names must be assigned, and the register block must be tied to the NCB register adapter (see the ALUTB base test's `connect_phase`).

Solution

In the top-level testbench, you need to instantiate the NCB interface, the ALU interface, and the ALU wrapper. You also need to push the interfaces into the configuration database.

```

105.    verif/fm/fm tb top.sv
// obj: ncb_i
// The <ncb_intf> instance.
ncb_intf ncb_i(.clk(tb_clk), .reset_n(tb_rst_n));

// obj: alu_i
// The <alu_intf> instance
alu_intf alu_i(.clk(tb_clk), .rst_n(tb_rst_n));

//-----
// Group: DUT
// (Instantiate the DUT and other modules here)
alu_wrapper alu_wrapper(*AUTOINST*/
                        // Interfaces
                        .ncb_i          (ncb_i),
                        .alu_i          (alu_i),
                        .frm_i          (frm_i),
                        // Inputs
                        .tb_clk         (tb_clk),
                        .tb_rst_n       (tb_rst_n));

//-----
// Group: Procedural Blocks
function void pre_run_test();
    uvm_config_db#(virtual cn_clk_intf) ::set(uvm_root::get(), "", "tb_clk_vi", tb_clk_i);

    // NCB Interfaces
    uvm_config_db#(virtual ncb_inb_intf.drv_mp)::set(uvm_root::get(), "", "ncb::inb_drv_vi",
        ncb_i.inb.drv_mp);
    uvm_config_db#(virtual ncb_inb_intf.mon_mp)::set(uvm_root::get(), "", "ncb::inb_mon_vi",
        ncb_i.inb.mon_mp);
    uvm_config_db#(virtual ncb_outb_intf.drv_mp)::set(uvm_root::get(), "", "ncb::outb_drv_vi",
        ncb_i.outb.drv_mp);
    uvm_config_db#(virtual ncb_outb_intf.mon_mp)::set(uvm_root::get(), "", "ncb::outb_mon_vi",
        ncb_i.outb.mon_mp);
    uvm_config_db#(virtual ncb_outb_intf.outb_bp_mp)::set(uvm_root::get(), "",
        "ncb_outb_bp_mon_vi",
        ncb_i.outb.outb_bp_mp);
    uvm_config_db#(virtual ncb_inb_intf.arb_drv_mp)::set(uvm_root::get(), "", "ncb::arb_drv_vi",
        ncb_i.inb.arb_drv_mp);
    uvm_config_db#(virtual ncb_inb_intf.arb_mon_mp)::set(uvm_root::get(), "", "ncb::arb_mon_vi",
        ncb_i.inb.arb_mon_mp);

    // ALU Interfaces
    uvm_config_db#(virtual alu_intf.drv_mp) ::set(uvm_root::get(), "", "alu::drv_vi",
        alu_i.drv_mp);
    uvm_config_db#(virtual alu_intf.mon_mp) ::set(uvm_root::get(), "", "alu::mon_vi",
        alu_i.mon_mp);

    // FRM Interfaces
    uvm_config_db#(virtual frm_intf.drv_mp) ::set(uvm_root::get(), "", "frm_drv_vi", frm_i.drv_mp);
    uvm_config_db#(virtual frm_intf.mon_mp) ::set(uvm_root::get(), "", "frm_mon_vi", frm_i.mon_mp);
endfunction : pre_run_test

// proc: Reset Busses
// Clear the ncb/rsl blocks
initial begin
    // clear out the ncb
    ncb_i.reset();
    alu_i.reset();
    frm_i.reset();

```

```
end
```

In the `alu_wrapper.sv` file, you need to have the `frm_intf` interface as an input, and wire it up to the `alu` block. In the base test, the register block comes not from the `frm_pkg`, but the `alutb_pkg`. In its build phase, you need to set the interface name for the NCB interface:

```
106.  verif/frm/alu_wrapper.sv  
      // set up the frm env  
      set_config_string("env.ncb_env",    "intf_name",    "ncb_vi");
```

Lesson 14: Advanced Sequences II

Virtual sequencers solve the important problem of how to create a sequence that controls more than one agent at a time. We will discuss these as well as more exotic uses of sequences such as adapter sequences.

Virtual Sequences and Sequencers

We know that a sequencer arbitrates among sequences and forwards sequence items--usually to a driver. Virtual sequencers are just a little bit different.

Virtual Sequencer

A virtual sequencer is one that does not send sequence items. Instead, it holds references to other sequencers and forwards sequences onto those sequencers. Virtual sequencers are not parameterized with a request or response type.

Likewise:


Virtual Sequence

A virtual sequence is any sequence that is running on a virtual sequencer. It is not parameterized to send requests and responses through a `seq_item_port`. It may push sequence items or other sequences, but it always pushes them to a sequencer that is referenced by the virtual sequencer on which it is running.

The current `frm_pkg::env_c` environment contains the `frm_agent` and the `alu_agent`, but does not contain the virtual sequencer (`vsqr`). Our goal is to create this sequencer, and then run virtual sequences on it that are capable of touching both agents.

As can be seen from the diagram, the `vsqr` component references the individual sequencers in `frm_agent` and `alu_agent`.

These are referred to as its subsequencers. The references should be set during the environment's connect phase. You can also add the register block to the virtual sequencer, in case any of the sequences running on it need to access CSRs.



Problem 14-1

Create the `frm_pkg::vsqr_c` component using `utg` (with the template named `vsqr`), populate the references to the `frm_agent` and `alu_agent` sequencers, and the `reg_block` instance, and instantiate it within `frm_pkg::env_c`.

Solution

utg provides most of what you need. All that needs to be done is to add the references.

```
107.    verif/vcomponents/frm/frm_vsqr.sv
`include "frm_sqr.sv"

// class: vsqr_c
// Virtual sequencer holding references to frm_sqr and alu_sqr
class vsqr_c extends uvm_sequencer;
    `uvm_component_utils_begin(frm_pkg::vsqr_c)
        `uvm_field_object(frm_sqr,          UVM_REFERENCE)
        `uvm_field_object(alu_sqr,          UVM_REFERENCE)
        `uvm_field_object(alutb_reg_block, UVM_REFERENCE)
    `uvm_component_utils_end

    //-----
    // Group: Sequencer references

    // var: frm_sqr
    sqr_c frm_sqr;

    // var: alu_sqr
    alu_pkg::sqr_c alu_sqr;

    //-----
    // Group: Fields

    // var: reg_block
    // alutb register block (reference to the one in cfg)
    alu_csr_pkg::alu_ncb_reg_block_c alutb_reg_block;

    //-----
    // Group: Methods
    function new(string name="vsqr",
                  uvm_component parent=null);
        super.new(name, parent);
    endfunction : new
endclass : vsqr_c
```

Note that the build phase is not used to create the subsequencers.

In the `frm_pkg::env_c`, the virtual sequencer is instantiated in the usual manner and the connect phase is used to perform the assignments. If either of the sequencers are not present (which would happen if either agent were accidentally set to `UVM_PASSIVE`), then a fatal error is reported:

```
108.    verif/vcomponents/frm/frm_env.sv
virtual function void connect_phase(uvm_phase phase);
    super.connect_phase(phase);

    if(!frm_agent.sqr || !alu_agent.sqr)
        `cn_fatal("frm_agent.sqr or alu_agent.sqr are not present!")

    vsqr.frm_sqr = frm_agent.sqr;
    vsqr.alu_sqr = alu_agent.sqr;
endfunction : connect_phase
```

Virtual Sequences

Virtual sequences are generally only different from other sequences in three ways.

1. First, they launch their sequences and sequence items using the ``uvm_do_on` family of macros, specifying the reference to the sequencer on which to launch the item as their second argument.
2. Second, in order to refer to their sequencer's subsequencers they must call the macro ``uvm_declare_p_sequencer`.
3. Third, library sequences do not support virtual sequences (as of UVM 1.1).

Uncomment the example virtual sequence in `frm_vseq_lib.sv`. It launches a random frame sequence item and an `alu_pkg::exer_seq_c` in parallel, with the number of ALU transactions matching the expected number in the frame.

```

109.    verif/vcomponents/frm/frm_vseq_lib.sv
class basic_vseq_c extends uvm_sequence;
    `uvm_object_utils(frm_pkg::basic_vseq_c)
    `uvm_declare_p_sequencer(vsqr_c)

    //-----
    // Group: Methods
    function new(string name="basic_vseq");
        super.new(name);
    endfunction : new

    //////////////////////////////////////
    // func: body
    virtual task body();
        frame_c frame;
        alu_pkg::exer_seq_c alu_exer_seq;

        // create and randomize to see how many ALU transactions to send
        `uvm_create_on(frame, p_sequencer.frm_sqr)
        frame.randomize();

        `cn_info($sprintf("Sending this frame: %s", frame.convert2string()))

        fork
            begin : send_frame
                `uvm_send(frame);
                get_response(rsp);
                `cn_info($sprintf("Frame completed: %s", rsp.convert2string()))
            end

            `uvm_do_on_with(alu_exer_seq, p_sequencer.alu_sqr, { count == frame.frame_len; })
        join

    endtask : body
endclass : basic_vseq_c

```

Some explanations are below:

```
alu_pkg::exer_seq_c alu_exer_seq;
```

Instead of creating our own sequence which runs a number of random ALU transactions, we can just re-use sequences from other vcomponents that this vcomponent is aware of.

```
`uvm_create_on(frame, p_sequencer.frm_sqr)
frame.randomize();
```

We want to randomize the frame ourselves so that we know how many ALU transactions to send in. Since a sequence's sequencer must be determined at the time of its creation, UVM provides the ``uvm_create_on` macro. Because we declared the `p_sequencer` variable with ``uvm_declare_p_sequencer`, that is how we reference the `frm_sqr` instance.

```
`uvm_send(frame);
```

Since the frame has already been created and randomized, we do not use ``uvm_do_on`, but instead we just want to launch it.

```
get_response(rsp);
`cn_info($sprintf("Frame completed: %s", rsp.convert2string()))
```

We sent in an object derived from sequence item (`frame_c`), so we need to get the response. In this agent, the response frame will not return until the frame with all its data has been pushed out of the framer block.

```
`uvm_do_on_with(alu_exer_seq, p_sequencer.alu_sqr, { count == frame.frame_len; })
```

Here, we are doing a sequence, specifying the non-virtual sequencer to do it on, and supplying the correct count of ALU transactions to run. UVM offers the ``uvm_do_on_with` macro for just such an occasion.

Problem 14-2

Nothing about the framer block says that the frame length has to be specified before or after the ALU transactions associated with it. In the file `verif/vcomponents/frm_vseq_lib.sv`, write a virtual sequence called `basic_delay_vseq_c` that varies which one comes first. Write another sequence, `exer_vseq_c`, that executes a random number of these. Set this as the default sequence of the test in `verif/frm/tests/basic.sv`.

See if you can find anything wrong with the simulation.

Solution

This implementation of `basic_delay_vseq_c` randomizes a signed integer between -100ns and 100ns. If the random delay is positive, it delays the frame first. If it is negative, it delays the ALU sequence instead.

```
110.    verif/vcomponents/frm/frm_vseq_lib.sv
class basic_delay_vseq_c extends uvm_sequence;
    `uvm_object_utils(frm_pkg::basic_delay_vseq_c)
    `uvm_declare_p_sequencer(vsqr_c)

    //-----
    // Group: Fields

    // frame_delay_ns
    // Delay of sending the frame, with respect to starting the ALU transactions
    // If frame_delay_ns is negative, then send the frame first, otherwise, send transactions first
    rand int frame_delay_ns;
    constraint frame_delay_cnstr { frame_delay_ns inside {[-100:100]}; }

    //-----
    // Group: Methods
    function new(string name="basic_delay_vseq");
        super.new(name);
    endfunction : new

    //////////////////////////////////////
    // func: body
    virtual task body();
        frame_c frame;
        alu_pkg::exer_seq_c alu_exer_seq;

        // create and randomize to see how many ALU transactions to send
        `uvm_create_on(frame, p_sequencer.frm_sqr)
        frame.randomize();

        `cn_info($sprintf("Sending this frame: %s", frame.convert2string()))

        fork
            begin
                if(frame_delay_ns > 0)
                    #(frame_delay_ns * 1ns);
                `uvm_send(frame);
                get_response(rsp);
                `cn_info($sprintf("Frame completed: %s", rsp.convert2string()))
            end

            begin
                if(frame_delay_ns < 0)
                    #((-frame_delay_ns) * 1ns);
                `uvm_do_on_with(alu_exer_seq, p_sequencer.alu_sqr, { count == frame.frame_len; })
            end
        join

    endtask : body
endclass : basic_delay_vseq_c
```

The exerciser sequence is a virtual sequence because it is running on a virtual sequencer. But it only launches sequences of type `basic_delay_vseq_c`, which run on the same sequencer as it does. So, it does not need to use ``uvm_do_on`, and does not need to refer to any of the subsequencers.

```

111.    verif/vcomponents/frm/frm_vseq_lib.sv
class exer_vseq_c extends uvm_sequence;
`uvm_object_utils_begin(frm_pkg::exer_vseq_c)
`uvm_field_int(count, UVM_ALL_ON | UVM_DEC)
`uvm_object_utils_end

//-----
// Group: Fields

// var: count
// The number of basic_vseq to do
rand int count;
constraint count_cnstr { count inside {[20:100]}; }

//-----
// Group: Methods
function new(string name="exer_vseq");
    super.new(name);
endfunction : new

////////////////////////////////////
// func: body
virtual task body();
    basic_delay_vseq_c vseq;
    `cn_info($psprintf("Transmitting %0d frames.", count))
    repeat(count)
        `uvm_do(vseq)
    endtask : body
endclass : exer_vseq_c

```

Everything about this test should seem to work. There is a problem, though. While your sequence prints that it's transmitting a certain number of frames, it will probably actually send fewer than that.

The problem is that the main phase of the simulation can end during the body of your sequence, and your sequence will never complete. That's because UVM does not implicitly wait for a component's default sequence to finish before moving on to the next phase. After all, it might be a persistent sequence that never ends.

We'll discuss this in more detail in Appendix A. In the meantime, you can solve this problem by having the sequence explicitly raise and drop the phase's objection. Do this by adding the following code to your `exer_vseq_c` class:

```

112.    verif/vcomponents/frm/frm_vseq_lib.sv
virtual task body();
    basic_delay_vseq_c vseq;
    `cn_seq_raise
    `cn_info($psprintf("Transmitting %0d frames.", count))
    repeat(count)
        `uvm_do(vseq)
    `cn_seq_drop
    endtask : body

```

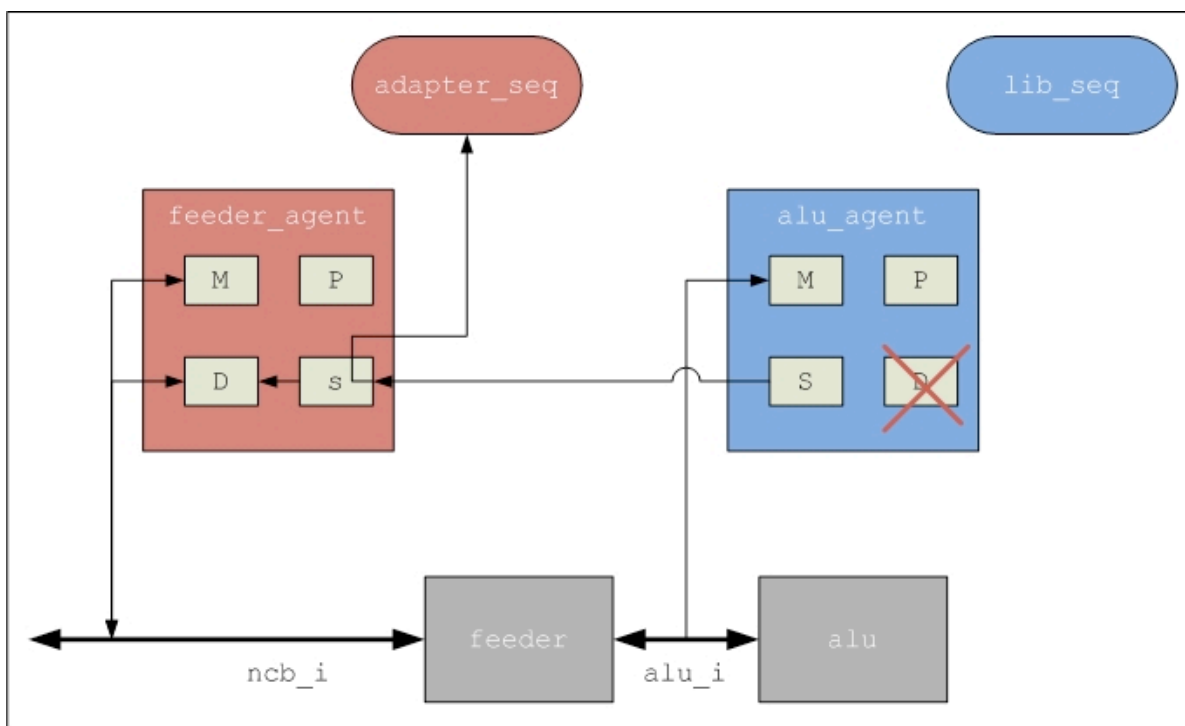
Adapter Sequences

With the example of virtual sequences, we have seen how block-level sequences can be re-used in higher-level testbenches by continuing to be run on the same sequencers for which they were designed. But what if, upon moving up to the full-chip level, you wanted to operate the same sequences on a different sequencer?

At a higher level your ALU agent may be passive, and the stimulus to the ALU block is instead fed by another RTL block. But, you've gone to a lot of trouble creating sequences that stimulate the ALU just the way you like them, why should you have to re-do them for a different interface? This is where adapter sequences can help.

One-to-One

A simple example would be a feeder block that received NCB store transactions that were then converted to ALU transactions. The objective would be to use the ALU stimulus that was already created for its block-level testbench in this higher-level testbench. To accomplish this, an ALU agent's sequencer could be directed to feed its ALU transactions stream to a port instantiated inside a feeder agent's sequencer instead of its own driver. On this sequencer is running an adapter sequence which is a persistent sequence that constantly pulls from the port just as a driver would. It then converts each transaction to the correct NCB transaction type, and drives them out its own sequencer.



This works well when there is a one-to-one correlation between the sequence items. In this case, an `alu_pkg::item_c` is directly converted to an `ncb_pkg::outb_trans_c`. And, the `alu_pkg::result_t` is converted to an `ncb_pkg::inb_trans_c`.

The basic algorithm for the adapter sequence's body would be as follows, where creation and conversion is done in an algorithm called `convert_to_ncb`:

```
113. task body();
```

```

alu_pkg::item_c alu_item;

forever begin
    p_sequencer.alu_item_port.get_next_item(alu_item);
    convert_to_ncb(outb_trans, alu_item);
    `uvm_send(outb_trans);
    alu_item.result = read_result();
    p_sequencer.alu_item_port.item_done(alu_item);
end
endtask : body

```

One-to-Many

Alternatively, what if the feeder block took several outbound NCB transactions to perform one ALU transaction? The feeder block might hypothetically operate by writing the A value to one CSR, the B value to a another CSR, and the operation type to a third CSR. An adapter sequence would also be the right choice here, as it would merely send three outbound NCB transactions for every one ALU transaction.

Many-to-One

A more complex example is one that has a many-to-one mapping of sequence items. The feeder block instead might receive NCB transactions that write a block of data into it, and then an NCB request that starts summing all of the numbers to be summed, with a polling mechanism to indicate when the transaction is complete.

We already have a sequence that sums an array. It even checks the final result for us. But, it transmits multiple transactions that need to be lumped into one transaction. The adapter sequence would have to be smart enough to know how to do this, and that may not be practical. Thought must be given to how and where complex scenario prediction should take place. In this example, an adapter sequence may not be the right solution, and putting the prediction within the sequence may have been a mistake because it would not be reusable in the higher-level testbench.

Adding Time Delays

One problem that our ALU agent suffers from is that all of the transactions are bunched together and driven as quickly as possible. As soon as one request is complete, it immediately fetches another one and drives it. While this makes for quick simulations, it probably does not represent realistic traffic conditions. There are two practical methods that can be used to solve this: adding delays in the sequence or adding delays to each transaction.

A sequence knows nothing of interfaces, signals, or clocks, because it has no virtual interfaces to attach to. If you want to create delays between transactions based on absolute time, then putting delays in a sequence before sending a sequence item works great.

If the frequency of your driver's clock is apt to change, though, the delays you incorporate into your sequences will unfortunately not vary along with it. Basing the delay on the clock frequency might be the better alternative in this case. To do so, you specify a random delay *per transaction*.

Problem 14-3

Add a field to the ALU's transaction class that indicates the number of clocks that the driver should wait before sending it. It should be constrained to a reasonable value, and optimally should be biased towards realistic conditions.

Then, implement the wait in the driver and see the results.

Solution

Adding the delay field to the `alu_pkg::item_c` class is simple. One thing you have to do, though, is ensure that when you declare the field with the macros that you add the `UVM_NOPACK` and `UVM_NOCOMPARE` flags:

```
114.    verif/vcomponents/alu/alu_item.sv
      `uvm_field_int (delay_in_clocks, UVM_ALL_ON | UVM_NOPACK | UVM_NOCOMPARE | UVM_DEC)
```

It is also worthwhile specifying in the name and comments what the timescale is for this delay. By naming it `delay_in_clocks`, it's clear how this delay will be used.

Then, declare the field and a typical constraint for it:

```
115.    verif/vcomponents/alu/alu_item.sv
      // var: delay_in_clocks
      // The number of clocks that the driver should delay before transmitting this transaction
      rand int unsigned delay_in_clocks;
      constraint delay_cnstr {
          delay_in_clocks dist { 0 :/ 40,
                                [1:10] :/20,
                                [11:30] :/20,
                                [31:100] :/ 10
                              };
      }
```

Finally, the ALU driver should obey the transaction delay as soon as it's fetched from the `seq_item_port`:

```
116.    verif/vcomponents/alu/alu_drv.sv
      seq_item_port.get_next_item(req);

      // wait the correct number of clocks before sending
      repeat(req.delay_in_clocks)
          @(drv_vi.drv_cb);
```

After doing this, you will notice that your transactions will now have gaps between them and the transaction stream is more realistic. You can then create derived transactions and use the factory to modify the distribution constraints to suit your needs. Better yet, you can add knobs to your vcomponent's CFG class as discussed in Lesson 11.

Conclusion

In this lesson, we learned the important concept of virtual sequences and sequencers. We also saw how adapter sequences can be re-used to take a sequence that resides on one interface and drive it on the next. Finally, we discussed how time delays between transactions can be achieved.

Appendix A: Phases and Heartbeats

Cavium's global environment contains both a watchdog timer and a heartbeat monitor. Together with UVM's phase objections, these components work together to ensure that simulation jobs are not running away needlessly.

Watchdog Timer

The watchdog timer will issue a fatal error when the simulation time exceeds the specified watchdog time, which can be set on a cmake command-line:

```
% cmake sim TEST=basic SIMU=VCS WDOG=300000
```

Watchdog times are always specified in nanoseconds.

Phase Objections

Phase objections exist to ensure that your various independent components do not get ahead of one another. It would be inappropriate to start sending packets into a DUT before reset was complete, for example.


Each of UVM's run-time phases ends immediately when all phase objections have been dropped. **Any task that was spawned during the phase but is still active will immediately be killed.**

The heartbeat monitor works in concert with a well-constructed environment which raises and drops objections in the correct manner. Where and when your objections take place depend on the current phase.

Reset Phase

The `cn` package contains a standard reset driver component, `cn_pkg::rst_drv_c`. This driver may be instantiated and configured to provide a simple reset on a single reset line. When this standard driver is used, it will manage the raising and dropping of the reset phase's objections for you.

If you decide to create your own reset driver, this phase's objection must be raised at the beginning of the phase, and be dropped once the reset is complete.



Configuration Phase

As discussed in Lesson 10, the configuration phase may be as simple as calling `register_block.update()`, or as complicated as you wish it to be. Like the reset phase, objections should be raised at the beginning and dropped at the end.

Main Phase

The main phase is substantially different from both the reset and configuration phases. The purpose of the main phase is to launch sequences that inject sequence items into the DUT. The main phase does not necessarily wait for the completion of these items, it merely lasts as long as it takes to send these items to the drivers.

Therefore, the main phase should end once all sequence items have been launched. Where these objections are raised and dropped depends entirely on how your sequences are launched.

If your main phase is creating and starting your sequences directly, then you must raise and drop the objections once they are complete. In Lesson 7, your first sequence was started in this manner:

```
117.
virtual task main_phase(uvm_phase phase);
    alu_pkg::alu_seq_c alu_seq = new("basic_seq");

    phase.raise_objection(this);
    `cn_info("Starting alu_seq.")
    alu_seq.start(alutb_env.alu_agent.sqr);
    phase.drop_objection(this);
endtask : main_phase
```

Alternatively, you may launch your sequences as the default sequence of the `main_phase`. In Lesson 12, this is how you launched the `sum_array_seq_c`:

```
118.
uvm_config_db#(uvm_object_wrapper)::set(this, "alutb_env.alu_agent.sqr.main_phase",
                                         "default_sequence", alu_pkg::sum_array_seq_c::type_id::get());
```

When a sequence is launched in this manner, a field called `starting_phase` is set to the main phase. Otherwise, this field is `null`. In either the sequence's body, or in pre- and post-body tasks, you can use this field to raise and drop the objection:

```
119.
virtual task pre_body();
    if(starting_phase)
        starting_phase.raise_objection();
endtask : pre_body

virtual task post_body();
    if(starting_phase)
        starting_phase.drop_objection();
endtask : post_body
```

Since you never know when a sequence is going to be configured to be a default sequence, you would have to add these tasks to every sequence you write. A recommended procedure is to create a base sequence that contains this code in your library sequence, and then derive all of your other sequences from it.

Alternatively, you can use the macros ``cn_seq_raise` and ``cn_seq_drop` as shown in Lesson 14.

Shutdown Phase

The shutdown phase is the DUT's opportunity to finish responding to all of the stimulus that was previously driven into it. This phase should end once all of the stimulus has *completed*. This is usually decided by a predictor whose scoreboards are empty, a monitor or driver awaiting responses, etc. Only components that can determine this information should be raising and dropping the shutdown phase's objection.

The following is a typical pattern for a predictor's shutdown phase:

```
120. virtual task shutdown_phase(uvm_phase phase);  
    if(scoreboard.size()) begin  
        phase.raise_objection(this, $psprintf("Waiting for %0d responses.", scoreboard.size()));  
        while(scoreboard.size())  
            @(response_received);  
        phase.drop_objection(this, "All responses complete.");  
    end  
endtask : shutdown_phase
```

The phase only raises an objection if there are outstanding packets, otherwise it does nothing. As shown here, an event must have been defined which is triggered whenever a response is received.

Once the shutdown phase is complete, your simulation will move to the extract and check phases.

Extract and Check Phases

The extract phase precedes the check phase and provides higher-level components with the opportunity to pull information from lower-level components in order to complete its prediction process.

The check phase is your environment's last chance to say whether or not the simulation passed or failed. This is where scoreboards must be checked to be empty, credits must be completed, etc.

These phases are not run-time phases. They are zero-time functions whose objections need not be raised or dropped. They are mentioned here because as you will see in the

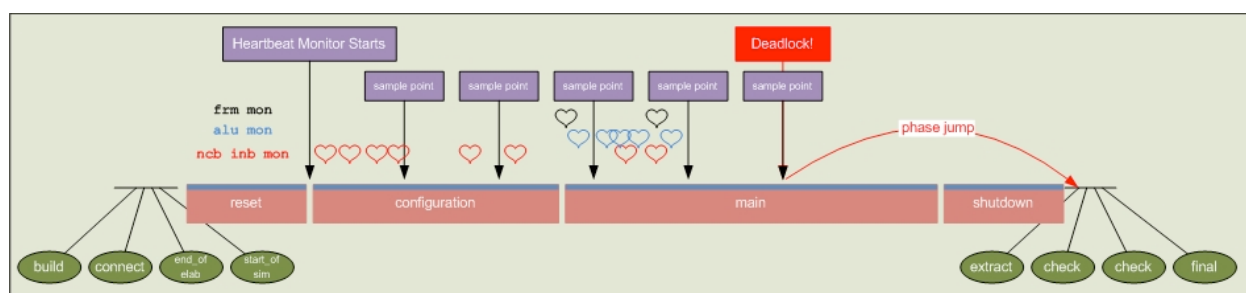
next section, it is important that your predictive components report errors during the check phase for any final conditions.

Deadlock Checking

Deadlock is when stimulus has gone into the DUT, but responses are not coming out. How long to wait before deciding that it will never come out is a function of your DUT and your environment. When a simulation is deadlocked, it should immediately exit and report a failure. It is also helpful to report what the environment is still waiting for.

The global environment's heartbeat monitor, `global_pkg::heartbeat_mon_c`, is the configurable component that is responsible for detecting a deadlock scenario. During the `end_of_elaboration` phase, components that are responsible for tracking responses from the DUT (i.e. monitors, predictors, subscribers, etc.) register themselves with the heartbeat monitor. Then, after reset has been completed, the heartbeat monitor periodically checks to ensure that *at least one* registered component has seen some form of activity from the DUT (the heartbeat). The length of time between these checks is called the *sample time*. The monitor will obey the *longest* sample time required by all registered components.

If no activity was seen during the last sample time, the simulation is considered deadlocked. In that scenario, all run-time phases are immediately halted, and the simulation proceeds to the extract, check, and final phases, where your environment has an opportunity to report what activity is still expected of the DUT. This is called a *phase jump*.



The art of getting the heartbeat monitor to work properly for your environment lies in determining which components must be registered, and how long the sample time should be. Your registered monitors will probably have different needs from one another. For example, RSL responses might come out just a few clocks after the request, whereas a complex algorithmic unit such as ZIP might take 50,000ns before a response comes out.

Each component registers itself and its required sample time during the `end_of_elaboration` phase using a standard macro:

121.

```
virtual function void end_of_elaboration_phase(uvm_phase phase);
  `global add_to_heartbeat_mon(7000ns);
endfunction : end_of_elaboration_phase
```

When a response to *stimulus* is seen from the DUT, then your component indicates a heartbeat with another macro. It's important to distinguish between responses to stimulus and responses to random traffic, such as random CSR reads.

The ALU monitor's `monitor_result` task would add this functionality:

122.

```
virtual task monitor_result();
  forever begin
    @(posedge mon_vi.mon_cb.ready);
    if(waiting_for_result == null)
      `cn_err("Monitored Result does not match any outstanding transaction.");

    `cn_info($psprintf("Monitored Result: %08X", mon_vi.mon_cb.result));
    monitored_result_port.write(mon_vi.mon_cb.result);
    `global_heartbeat("Result seen");
    waiting_for_result = null;
  end
endtask : monitor_result
```

The ALU monitor's check phase should report an error if it is still waiting for a response:

123.

```
virtual function void check_phase(uvm_phase phase);
  super.check_phase(phase);
  if(waiting_for_result != null)
    `cn_err($psprintf("Monitor is currently waiting for a result to transaction: %s",
      waiting_for_result.convert2string()));
endfunction : check_phase
```

Both of these methods use the variable `waiting_for_response`, which is a reference to the outstanding request and would be set by the `monitor_item` task.

By default, the heartbeat monitor starts in the `pre_configuration` phase, but it can be configured to start during `pre_reset` or `pre_main` instead. It stops in the `post_shutdown` phase.

Appendix B: Common Recipes

Implementing a Watchdog

Often, you will want to send something into the DUT and ensure that it completes within a “reasonable” amount of time. Otherwise, you want to report an error. Depending on your scenario, this may be measured in nanoseconds, clocks, or some other triggering event.

A simple method to do this is as follows:

```
124.
fork
    do_something();
begin
    #(5000ns);
    `cn_error("Watchdog timeout!")
end
join_any
disable fork;
```

Either the `do_something()` task will complete its mission, or the error will occur in 5,000ns.

Forking Multiple Instances of a Method

There will be myriad situations where you will want to fork off multiple instances of the same method call with different arguments. Doing this simple approach will not give the desired results:

```
for(int num=0; num < 5; num++)
    fork
        do_stuff(num);
join_none
```

Nor will this:

```
fork
    for(int num=0; num < 5; num++)
        do_stuff(num);
join
```

The problem with both of these examples is that while four tasks will be launched, they will all be launched with an argument of 4. This is due to the fact that the loop will add 4 tasks to the scheduler, and when they later launch the `num` variable will have a final value of 4.

To fix this, you need an automatic variable to accept the loop value:

```
125.
for(int num=0; num < 5; num++) begin
```



```

fork
    automatic int _num = num;
    do_stuff(_num);
join_none
end

```

Phase-Boundary Crossing Tasks

Any task that starts in a run-time phase (other than the `run_phase`), will be killed at the end of that phase. Raising and holding an objection on that phase may not be desirable if you want that task to continue into a later phase.

A way around that is to wait for a start event, launch your task during the `run_phase`, and then kill it upon receiving a finish event. In the following example, the task `do_stuff()` launches at the beginning of the main phase and continues through to the end of the shutdown phase.

```

126.
event start_event, finish_event;

////////////////////////////////////
virtual task run_phase(uvm_phase phase);
    @(start_event);
    fork
        do_stuff();
        @(finish_event);
    join_any
    disable fork;
endtask : run_phase

////////////////////////////////////
virtual task main_phase(uvm_phase phase);
    -> start_event;
endtask : main_phase

////////////////////////////////////
virtual task post_shutdown_phase(uvm_phase phase);
    -> finish_event;
endtask : post_shutdown_phase

```

Handling a Reset

Drivers and monitors should be made aware of the reset signals on the interfaces. If a test chooses to suddenly perform a reset, these components must handle the event cleanly.

A typical method for doing this is shown below in an example driver:

```

127.
virtual task run_phase(uvm_phase phase);
    forever begin
        @(posedge my_vi.rst_n);

        fork
            drive_items();
            @(negedge my_vi.rst_n);
        join_any
        `cn_info("Saw a reset. Time to clean up.")
        disable fork;
        cleanup();
    end
end

```

```
endtask : run_phase
```

The phase first waits for the positive edge of the `rst_n` signal to indicate that reset is complete. The `drive_items` task is expected to run forever, and so only the falling edge of the `rst_n` signal will cause the `join_any` to fall through. The `disable fork` statement will kill the `drive_items()` task, and the cleanup function will reset any of the driver's fields which track state. The whole task is wrapped in a `forever` block so that it loops back and is able to drive more items once the reset event is finished.